

QUESTION: 1

You have a Microsoft Foundry project that contains an agent.

The agent uses a knowledge source built from documents stored in Azure Blob Storage. The documents include digitally scanned PDFs that contain multipage tables. You have an ingestion job that extracts only plain text, causing loss of table structure, headings, and page- number metadata.

Users frequently ask questions that require the retrieval of specific table rows across the pages.

You need to configure an ingestion job for a Retrieval Augmented Generation (RAG) pipeline that performs optical character recognition (OCR) on scanned PDFs, preserves tables and headings as structure-aware chunks, and stores page-number metadata with each chunk.

How should you configure the ingestion job?

- A.** Use basic parsing and fixed-size chunking.
- B.** Use advanced data parsing to reingest the documents.
- C.** Use OCR and page-level chunking.
- D.** Use page-level OCR extraction and store each page as a single chunk.

Answer: B

The correct configuration is advanced data parsing because the issue is not merely OCR; the ingestion job must preserve document structure for reliable RAG retrieval. Microsoft guidance for advanced parsing states that it automatically detects tables across all pages, including tables in scanned documents, merges tables that span multiple pages, restores column headers, and creates table chunks with metadata such as table index, shape, page numbers, section headings, and table previews. This directly satisfies the requirement to retrieve specific rows from multipage tables while retaining source-page context.

Basic parsing with fixed-size chunking would flatten the document into arbitrary text fragments, which is the current failure mode. OCR with page-level chunking improves text extraction from scanned PDFs, but it does not provide structure-aware chunks that preserve headings and table relationships across pages. Storing each page as a single chunk is too coarse for row-level retrieval and can bury relevant table rows in excessive context. Advanced data parsing is purpose-built for RAG ingestion because it produces semantically meaningful, retrievable chunks and enriches them with metadata needed for citations and grounding.

Reference topics: RAG ingestion, advanced parsing, OCR, table extraction, structure-aware chunking, page metadata, and Azure Blob Storage document ingestion.

QUESTION: 2

Your company is piloting a customer support agent in a Microsoft Foundry project name Project1. Project1 is connected to an existing Application Insights resource, and the company 's support team reviews runs in the Traces tab.

The Foundry Agent Service is configured to perform the following actions:

- * Retrieve the Application Insights connection string by calling `project_client.telemetry.get_application_insights_connection_string()`.
- * Call `configure_azure_monitor(connection_string=...)` to enable telemetry.

A separate LangChain service configured to use OpenTelemetry and has the following configurations:

- * Uses `AzureAIOpenTelemetryTracer(connection_string=..., enable_content_recording=False)`
- * Passes the tracer by using `config={"callbacks": [azure_tracer]}`

Company policy has the following requirements:

- * Telemetry from LangChain and OpenTelemetry must be distinguishable within the same Application Insights resource.
- * Secrets and credentials must NOT be stored in prompts, tool arguments, or span attributes.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area		
Statements	Yes	No
The LangChain service will appear in Traces without configuring a tracer.	☐	☐
Setting different OTEL_SERVICE_NAME values separates the services in Application Insights.	☐	☐
When using <code>enable_content_recording=False</code> , prompts and tool data will be captured in the telemetry.	☐	☐

Answer:

Answer Area

Statements

The LangChain service will appear in Traces without configuring a tracer,

Yes

☐

No

☒

Setting different OTEL_SERVICE_NAME values separates the services in Application Insights.

☒☐

When using enable_content_recording=False, prompts and tool data will be captured in the telemetry.

☐☒

Explanation:

The LangChain service will appear in Traces without configuring a tracer: No
Setting different OTEL_SERVICE_NAME values separates the services in Application Insights: Yes
When using enable_content_recording=False, prompts and tool data will be captured in the telemetry: No
The first statement is No because a separate LangChain or LangGraph application must emit telemetry through the configured tracing integration. Microsoft's LangChain tracing guidance states that you configure AzureAIOpenTelemetryTracer, attach it to the runnable or agent through callbacks, and then inspect the emitted traces in Azure Monitor. The troubleshooting guidance also states that missing LangChain or LangGraph spans are caused by tracing callbacks not being attached to the run.

The second statement is Yes . In OpenTelemetry, OTEL_SERVICE_NAME maps to the service.name resource attribute. Azure Monitor Application Insights uses cloud role names to represent separate services, and Microsoft states that when multiple services emit to the same Application Insights resource, cloud role names must be set so services are represented properly.

The third statement is No . enable_content_recording=False is specifically used to redact message content and tool call arguments from traces. Microsoft also advises disabling content recording in production and not storing secrets, credentials, or tokens in prompts or tool arguments. Reference topics: Microsoft Foundry tracing, LangChain tracing, OpenTelemetry service naming, Application Insights, and secure telemetry configuration.

QUESTION: 3

You need to recommend a plan to create a customer support agent by using the Microsoft Foundry Agent Service. The agent must meet the following requirements:

- * Retain user preferences across multiple conversations.**
- * Enable users to provide contextual grounding by directly uploading documents during a chat.**

Which Foundry capability should you recommend for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

To retain user preferences across conversations, use:

- Agent memory that uses persistent storage
- Conversation history
- Orchestration-managed session context

To enable users to provide contextual grounding during chats, use the:

- Azure AI Search tool
- Code interpreter tool
- File search tool

Answer:

To retain user preferences across conversations, use:

- Agent memory that uses persistent storage
- Conversation history
- Orchestration-managed session context

To enable users to provide contextual grounding during chats, use the:

- Azure AI Search tool
- Code interpreter tool
- File search tool

Explanation:

To retain user preferences across conversations, use: Agent memory that uses persistent storage To enable users to provide contextual grounding during chats, use the: file search tool The correct capability for retaining user preferences is agent memory that uses persistent storage .

Microsoft Foundry Agent Service memory is a managed long-term memory capability that enables continuity across sessions, devices, and workflows. It is specifically intended to let agents retain user preferences, maintain relevant historical context, and personalize responses across separate conversations. Memory stores provide the persistent storage layer, and scope can be used to segment memories for secure user-specific experiences. The correct capability for contextual grounding from user-uploaded documents is the file search tool .

Microsoft describes file search as the tool that enables Foundry agents to search through documents and retrieve relevant information from outside the base model, including proprietary product information and user- provided documents. The file search workflow supports uploading files, creating a vector store, enabling the tool on the agent, and querying those documents through the agent.

Conversation history alone supports continuity within a conversation, but it is not durable

preference memory across multiple conversations. An Azure AI Search tool is better for preconfigured enterprise indexes, while file search is the direct document-upload grounding capability. Reference topics: Foundry Agent Service memory, memory stores, File Search tool, vector stores, and grounded agent responses.

QUESTION: 4

You need to configure Agent1 to answer customer questions about only the Contoso products. The solution must meet the business requirements.

What should you do?

- A.** Apply top-p sampling.
- B.** Modify the system message instructions.
- C.** Add few-shot examples.
- D.** Increase the value of the temperature parameter.

Answer: B

The correct answer is B. Modify the system message instructions . The case study states that Agent1 answers general questions about Contoso products and that a business requirement is for Agent1 to answer questions only about products sold by Contoso. This requirement defines the agent's allowed domain and refusal boundary, so it must be expressed in the agent's system-level instructions. Microsoft Foundry guidance states that system messages steer Azure OpenAI chat model behavior and are used to define the assistant's role, boundaries, output format, and safety or quality constraints.

The system message should instruct Agent1 to answer only Contoso-product questions, use Contoso product documentation when available, and decline questions about non-Contoso products. This directly enforces the intended business scope at the highest instruction level. Few-shot examples can reinforce desired behavior but are not the primary control for defining mandatory operating boundaries. Top-p sampling and temperature are decoding controls; they influence randomness and diversity, not whether the agent restricts answers to a specific product domain. Increasing temperature would likely reduce consistency. Reference topics: Microsoft Foundry agent instructions, system message design, prompt engineering, response boundaries, and grounded generative AI behavior.

QUESTION: 5

You have a Microsoft Foundry project that contains an agent. The agent has a Model Context Protocol (MCP) tool that queries a knowledge base stored in Azure AI Search. Some agent runs return answers from the base model without invoking the knowledge

base, which results in responses without grounded citations.

You are provided with the following code snippet that runs the agent.

```
run = project_client.agents.runs.create_and_process(  
    thread_id=thread.id,  
    agent_id=agent.id,  
)
```

You need to add the correct `tool_choice` parameter to the code to deterministically force the agent to invoke the MCP tool on each run.

What should you add?

- A. `tool_choice = { " type " : " mcp " }`
- B. `tool_choice = { " auto " }`
- C. `tool_choice = { " type " : " knowledge_base " }`
- D. `tool_choice = { " required " }`

Answer: D

The correct selection is D . In Microsoft Foundry Agent Service, `tool_choice` is the runtime control used to influence whether the model may answer directly or must invoke a tool. Microsoft's tool best-practice guidance states that `auto` lets the model decide whether to call tools, `none` prevents tool calls, and `required` means the model must call one or more tools. This directly addresses the issue where some runs answer from the base model and skip the knowledge base.

For an agentic retrieval solution backed by Azure AI Search through an MCP tool, Microsoft's tutorial states that setting `tool_choice = " required "` ensures the agent always uses the knowledge base tool when processing queries. This produces grounded answers because the run is forced into tool invocation before responding.

`auto` is incorrect because it preserves the nondeterministic behavior already causing missing citations. `{ " type`

`": " knowledge_base " }` is not a valid Foundry tool-choice type. `{ " type " : " mcp " }`

describes an MCP tool type in some Responses API schemas, but the deterministic guarantee for this agent run scenario is the required tool-call mode. Reference topics:

Microsoft Foundry Agent Service, MCP tools, Azure AI Search agentic retrieval, `tool_choice`, and grounded citations.

QUESTION: 6

You have a customer support agent that uses the Microsoft Foundry Agent Service.

Sometimes, customers return to a session days later to continue the same support case, and the agent must resume with the full historical context. The agent must provide the following:

- * Multi-turn continuity within the session**

- * Cross-session continuity for the same case**

- * Access to the full interaction history, including user messages, agent messages, tool calls, and tool outputs** You need to ensure that the agent automatically reloads the complete history on each new turn.

What should you do?

A. Persist only the final model response stored in the client application and prepend the response to future prompts.

B. Enable memory summarization on the agent definition to persist the context automatically.

C. Create and reuse a conversation by storing the conversation's ID and supplying the ID on subsequent requests.

Answer: C

The correct approach is to create and reuse a conversation by storing the conversation's ID and supplying that ID on subsequent requests . In Microsoft Foundry Agent Service, conversations are durable objects with unique identifiers that can be reused across sessions. The official runtime guidance states that conversations store items, including messages, tool calls, tool outputs, and other data, and are intended for multi-turn continuity, cross-session continuity, and inspection of what happened over time. This directly satisfies the requirement to resume the same support case days later with the full historical context.

Persisting only the final model response is insufficient because it loses the full interaction chain, especially tool calls and tool outputs that may be essential to case state. Memory summarization is also not the best fit because the requirement asks for the complete history, not a compressed representation that may omit details.

Reusing the conversation ID allows Foundry to maintain the conversation server-side so the next turn can reuse prior context without the client manually rebuilding prompts.

Reference topics: Foundry Agent Service runtime components, conversations, conversation items, multi-turn continuity, cross-session continuity, and tool output history.

QUESTION: 7

You are building a speech processing solution in Microsoft Foundry for a customer

support platform.

The platform will transcribe live phone calls, so that supervisors at your company can view call transcripts and detect issues while the calls are in progress. The call audio will arrive as a continuous stream from the telephony system.

You need to ensure that the call transcripts appear within only a few seconds of the audio stream.

What should you do?

- A.** Run a batch transcription job on recorded audio files.
- B.** Use real-time speech to text to process streaming audio input.
- C.** Use speech translation to generate the transcripts into multiple languages.
- D.** Use text to speech by using a custom neural voice.

Answer: B

The correct answer is B. Use real-time speech to text to process streaming audio input .

The scenario requires live transcription from a continuous telephony stream, with transcript text appearing within a few seconds while the call is still in progress. Azure Speech in Foundry Tools real-time speech recognition is specifically intended for immediate transcription scenarios such as call center assistance, dictation, and live meeting captioning. Microsoft's Speech guidance describes real-time speech to text as processing audio input and returning transcriptions in real time, which matches the supervisor monitoring requirement.

Batch transcription is inappropriate because it processes stored audio files after recording, not an active live stream. Speech translation is used when the primary goal is translating speech into another language, not simply producing live same-language call transcripts. Text to speech performs the reverse operation by generating spoken audio from text and does not transcribe inbound calls. Real-time speech to text provides the low-latency streaming recognition path required for live operational monitoring. Reference topics: Azure Speech in Foundry Tools, real-time speech recognition, streaming audio input, call center transcription, and live captions.

QUESTION: 8

You have a Microsoft Foundry project that ingests scanned PDF invoices stored in Azure Blob Storage. Each invoice contains printed line items and has a table-based layout.

Extracted results are stored as structured JSON and used as grounding data for an

agent in a Retrieval Augmented Generation (RAG) solution.

You need to create a single analyzer that meets the following requirements:

*** Extracts the invoice number, invoice date, vendor name, and total amount across varying templates**

*** Returns confidence scores so that results with confidence below 0.80 can be routed for supervisor review What should you use?**

A. the Azure Content Understanding in Foundry Tools prebuilt-layout analyzer

B. a Foundry agent that has groundedness guardrails enabled to extract invoice fields and confidence scores

C. a custom Azure Content Understanding in Foundry Tools analyzer that defines the required fields as the extracted fields and the returned confidence scores for routing

D. the Azure Content Understanding in Foundry Tools prebuilt-documentSearch analyzer and search.score from the Azure AI Search results for routing

Answer: C

The correct answer is C because the requirement is structured field extraction from invoices across varying templates, not only OCR or layout preservation. Azure Content Understanding analyzers are reusable configurations that combine content extraction, AI-powered analysis, and structured data output, and Microsoft states that custom analyzers can be created for specific extraction needs. In this case, the analyzer schema should define fields such as invoice number, invoice date, vendor name, and total amount so the output can be returned as structured JSON for downstream RAG grounding.

The confidence-routing requirement also points to Content Understanding field confidence scores. Microsoft documentation states that every field can include a confidence score from 0 to 1, and that confidence scores can be used to automate high-confidence results while routing low-confidence results for human review. A threshold such as 0.80 is therefore an application routing rule based on the returned field confidence. The prebuilt-layout analyzer preserves layout but does not define invoice-specific business fields. Groundedness guardrails evaluate generated answers, not invoice field extraction. Azure AI Search search.score measures retrieval relevance, not extraction confidence. Reference topics: Content Understanding custom analyzers, document field extraction, structured JSON output, confidence scoring, and RAG grounding.

QUESTION: 9

You have a Microsoft Foundry project that serves a high-volume chat app.

**Most requests are simple FAQs, but some require advanced reasoning.
You need to reduce costs and latency for common queries, without degrading the quality of the responses to complex questions.**

What should you do?

- A.** Increase the value of the max_tokens parameter for all the requests.
- B.** Route all the requests to a smaller model.
- C.** Route all the requests to the most capable model.
- D.** Use a model cascade that routes the requests to different models.

Answer: D

The correct choice is to use a model cascade that routes the requests to different models . In Microsoft Foundry, this pattern aligns with model routing: simple, low-risk prompts can be handled by smaller, faster, lower-cost models, while complex prompts can be escalated to more capable or reasoning models. Microsoft's Foundry model router guidance states that the router optimizes cost and latency while maintaining comparable quality by using smaller, cheaper models when they are sufficient and larger or reasoning models when the task requires more advanced capability.

This directly matches the scenario: most traffic consists of simple FAQs, so routing those requests to efficient models reduces average latency and token-processing cost.

Advanced reasoning requests still receive high- quality responses because they are routed to models with stronger reasoning capability. Microsoft's model router documentation also explains that routing decisions consider prompt difficulty, cost, quality, latency, and conversation context, making it suitable for diverse chat workloads.

Increasing max_tokens for all requests would usually increase cost and latency. Sending all requests to a smaller model risks poor quality for complex questions, while sending all requests to the most capable model wastes cost and latency on simple FAQs. Reference topics: Microsoft Foundry model routing, model selection, generative AI optimization, latency management, and cost-aware AI application design.

QUESTION: 10

You have a chat app in a Microsoft Foundry project and an Azure AI Search vectorized index.

You need to connect to the index to meet the following requirements:

- * Complex questions must retrieve information from multiple chunks.**
- * Multi-turn conversations must influence retrieval planning.**

*** Retrievals must run in parallel to reduce latency.**

Which retrieval approach should you use?

- A.** classic Retrieval Augmented Generation (RAG)
- B.** chain of thought
- C.** agentic Retrieval Augmented Generation (RAG)
- D.** iterative retrieval

Answer: C

The correct answer is agentic Retrieval Augmented Generation (RAG) because the requirements describe the agentic retrieval pipeline in Azure AI Search. Agentic retrieval is designed for chat and copilot scenarios where a user's request can be complex, conversational, and dependent on prior turns. Azure AI Search agentic retrieval uses an LLM-assisted planning stage to break a complex request into focused subqueries, allowing the system to retrieve grounding information from multiple chunks rather than relying on a single query path.

Microsoft's Azure AI Search guidance describes agentic retrieval as a multi-query pipeline for complex questions in chat and agent workflows, with subqueries that can include chat history for additional context.

This also satisfies the latency requirement because agentic retrieval runs the generated subqueries in parallel and then merges and reranks the best results for use by the generative model. Classic RAG is simpler and typically sends a single query to search, making it less suitable for multi-hop or conversational retrieval planning. Chain of thought is a reasoning technique, not an Azure AI Search retrieval approach, and iterative retrieval does not specifically provide the built-in query planning, conversation-aware retrieval, and parallel execution described here. Reference topics: Azure AI Search agentic retrieval, RAG with Azure AI Search, knowledge bases, query planning, and generative AI grounding.

QUESTION: 11

You have a Microsoft Foundry project that contains three agents as shown in the following table.

Name	Description
TriageAgent	Classifies incoming customer requests
PolicyAgent	Answers policy questions by searching internal content
ActionAgent	Creates or updates tickets by calling an HTTP API

You need to orchestrate the agents to ensure that the customer requests meet the

following requirements:

*** Support a deterministic, step-based process that uses conditional branching and shared state across the agents.**

*** Optionally trigger a ticket action based on the triage result.**

The solution must minimize development effort.

What should you include in the solution?

A. separate agent runs coordinated in the application code

B. a multi-agent group chat session

C. a workflow

D. threads and runs without a workflow

Answer: C

The correct answer is a workflow . Microsoft Foundry workflows are designed to orchestrate agents and business logic as declarative, predefined sequences of actions. The official workflow guidance states that workflows are ideal when you need to orchestrate multiple agents in a repeatable process, add branching logic such as if/else, and handle variables without writing application orchestration code. This directly matches the requirement for a deterministic, step-based process with conditional branching and shared state.

In this scenario, TriageAgent can classify the request first, the workflow can store the triage result, and conditional logic can determine whether to invoke PolicyAgent, ActionAgent, or both. The ticket action is optional, so it should be triggered through a workflow condition based on the triage output. This minimizes development effort because the branching, sequencing, and variable handling are managed in the Foundry workflow rather than being manually implemented across separate runs in application code.

A group chat session is better for dynamic agent handoff, not a strict deterministic process. Threads and runs or separate app-coordinated calls require more custom orchestration.

Reference topics: Microsoft Foundry workflows, multi-agent orchestration, conditional branching, variable handling, and agent-driven workflows.

QUESTION: 12

You are building a web app named App1 that generates responses by using a model deployed to a Microsoft Foundry project named Project1.

Before sending the prompts to the model, App1 must retrieve documents by using Azure AI Search.

You need to integrate Project1 and App1. The solution must meet the following

requirements:

- * Multiple client applications must use the same search configuration.**
- * A security policy must prevent key-based authentication.**
- * Administrative effort must be minimized.**

What should you do?

- A.** Enable a managed identity for each application and call Azure AI Search directly.
- B.** Create a custom HTTP connection in Foundry and manually configure Azure AI Search endpoints per application.
- C.** Call Azure AI Search directly from each application by using Microsoft Entra authentication.
- D.** Configure an Azure AI Search connection in Project1 and reference the connection in each application.

Answer: D

The correct solution is to configure an Azure AI Search connection in Project1 and reference that connection from each application. Microsoft Foundry project connections are intended to centralize external resource configuration for a project. The official Foundry guidance states that when adding a connection, you select the external service, such as Azure AI Search, and choose the authentication method for that resource.

This directly satisfies the requirement for multiple client applications to use the same search configuration instead of duplicating endpoints, indexes, and credentials in each app.

The security requirement is met by configuring the connection with keyless authentication.

The Azure AI Search integration guidance specifies that a Foundry project connection requires the search endpoint and either key-based authentication or keyless authentication with Microsoft Entra ID. For keyless authentication, RBAC roles are assigned to the project's managed identity, eliminating hard-coded API keys.

Options A and C use secure identity-based access but force every client application to maintain its own Azure AI Search configuration, increasing administration and violating the shared-configuration requirement. Option B also increases administrative work and does not use the native Azure AI Search connection type. Reference topics: Microsoft Foundry project connections, Azure AI Search grounding, Microsoft Entra ID authentication, managed identities, and RBAC.

QUESTION: 13

You have an application named App1 that uses Azure Speech in Foundry Tools to transcribe live calls.

Transcript segments often contain both English and Spanish. App1 sends each segment to Azure Translator in Foundry Tools to translate to another language. Sometimes, mixed-language segments result in incomplete or incorrect translations. You need to reduce translation errors. The solution must ensure that the entire transcript is translated successfully.

What should you do before sending the segments to Translator?

- A. Use document translation to translate the entire transcript as a single document.
- B. Split the mixed-language segments into single-language segments and translate each segment separately.
- C. Enable automatic language detection for the translation request.
- D. Specify English as the source language in the translation request for all the segments.

Answer: B

QUESTION: 14

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

You need to improve response completeness. The solution must be implemented in the logic of the application code before responses are returned.

What should you do?

- A. Decrease the value of the temperature parameter.
- B. Add a reflection pass before the responses are returned.
- C. Switch to Retrieval Augmented Generation (RAG).
- D. Enable response streaming.

Answer: B

The correct answer is B. Add a reflection pass before the responses are returned . A reflection pass is an application-orchestration step in which the generated summary is reviewed before final delivery, typically by asking the model or an evaluator step to check

whether the answer covers the retrieved policy evidence and to revise the response when important details are missing. This directly addresses response completeness in application logic before the response is returned. The Microsoft Learn study guide explicitly includes Implement model reflection and Apply prompt engineering techniques to improve responses under optimization and operationalization of generative AI solutions.

This is also consistent with Microsoft Foundry agentic-loop guidance, which identifies reflection and planning cycles as patterns for multi-step reasoning in production agent systems. Completeness is a response-quality property: Azure AI evaluation defines completeness as whether a response contains all necessary and relevant information with respect to ground truth.

Option C is not correct because the scenario already says the agent generates summaries from retrieved policy documents, which is already a grounded retrieval pattern. Option A mainly reduces randomness, not missing content. Option D improves delivery experience, not answer completeness. Reference topics: model reflection, prompt engineering, agentic loops, response evaluation, and grounded generative AI solutions.

QUESTION: 15

You need to recommend an invoice review solution that resolves the issue reported by the finance department.

What should you include in the recommendation?

- A.** Azure Content Understanding in Foundry Tools
- B.** chat completions
- C.** Azure Document Intelligence in Foundry Tools
- D.** Image Analysis

Answer: A

The correct recommendation is Azure Content Understanding in Foundry Tools . The case study states that Contoso's finance department must manually review vendor invoices to verify that invoice details match vendor contract terms, and that the invoices contain tables, logos, and varied layouts that make consistent processing difficult. It also states that the planned solution must evaluate both the visual layout and textual content of the invoices.

Azure Content Understanding is designed for this type of multimodal document-processing workload.

Microsoft describes Content Understanding as a Foundry Tool that processes unstructured

and multimodal content, including documents and images, and transforms it into structured output for AI applications. It can use document analyzers to extract text, layout, tables, fields, and relationships from diverse document types.

Chat completions alone would not reliably extract structured invoice fields from complex layouts. Azure Document Intelligence can extract OCR, layout, and tables, but Content Understanding is the better end-to-end Foundry capability for combining visual and textual understanding with structured extraction for downstream verification. Image Analysis focuses on image-level visual features and is insufficient for invoice field and table review. Reference topics: Content Understanding, document analyzers, multimodal extraction, invoice processing, tables, layout, and structured JSON output.

QUESTION: 16

You have a Microsoft Foundry project that contains an agent.

You need to enable long-term memory to ensure that the agent can recall user preferences across separate conversations. Stored memories must be isolated per authenticated user without the client application manually generating user IDs.

How should you complete the Python code? To answer, drag the appropriate values to the correct targets.

Each value may be used once, more than once, or not at all.

NOTE: Each correct selection is worth one point.

Values	Answer Area
<pre>:: "session"</pre>	<pre>from azure.ai.projects.models import MemorySearchTool, PromptAgentDefinition</pre>
<pre>:: "{{\${conversationId}}}"</pre>	<pre>mem_store_name = "agent_mem_store"</pre>
<pre>:: "{{\${userId}}}"</pre>	<pre>memory_tool = MemorySearchTool(</pre>
<pre>:: [mem_store_name]</pre>	<pre> memory_store_name=mem_store_name,</pre>
<pre>:: [memory_tool]</pre>	<pre> scope= " ", Value ,</pre>
<pre>:: MemorySearchTool("support_mem_store")</pre>	<pre>)</pre>
	<pre>agent_def = PromptAgentDefinition(</pre>
	<pre> model="gpt-5.2",</pre>
	<pre> instructions="You are a customer support assistant.",</pre>
	<pre> tools= Value</pre>
	<pre>)</pre>

Answer:

Values

```
["session"]
["{{$conversationId}}"]
["{{$userId}}"]
[mem_store_name]
[memory_tool]
MemorySearchTool("support_mem_store")
```

Answer Area

```
from azure.ai.projects.models import MemorySearchTool, PromptAgentDefinition

mem_store_name = "agent_mem_store"

memory_tool = MemorySearchTool(
    memory_store_name=mem_store_name,
    scope="{{$userId}}",
)

agent_def = PromptAgentDefinition(
    model="gpt-5.2",
    instructions="You are a customer support assistant.",
    tools=[memory_tool]
)
```

Explanation:

scope = "{{\$userId}}"

tools = [memory_tool]

The correct scope value is {{\$userId}} because the requirement is per-authenticated-user memory isolation without the client application manually generating user identifiers. In Microsoft Foundry Agent Service memory, the scope parameter partitions memory items inside the memory store. The official guidance states that when the memory search tool is attached to an agent, setting scope to the user identity template enables per-user memory isolation; the service resolves the end-user identity from the request header when provided, or falls back to the Microsoft Entra tenant ID and object ID of the caller. This matches the requirement to isolate stored preferences by authenticated user automatically.

The tools property must be [memory_tool] because the MemorySearchTool instance is created earlier and must be attached to the PromptAgentDefinition. Foundry guidance shows the memory search tool being passed in the agent definition as tools=[tool], allowing the agent to read from and write to the configured memory store during conversations.

"session" and {{\$conversationId}} would limit continuity to a session or conversation instead of enabling long-term recall across separate conversations. [mem_store_name] is a list containing the store name, not a tool definition. Reference topics: Foundry Agent Service memory, memory stores, memory search tools, scope, and per-user isolation.

QUESTION: 17

You have a Microsoft Foundry project named Project1.

Project1 contains an application that processes PDF vendor invoices.

You need to configure Azure Document Intelligence in Foundry Tools to generate a Markdown output that preserves the sections and table structure of the PDFs. The

solution must minimize development effort.

What should you do?

- A.** Increase the confidence threshold.
- B.** Configure output=figures when you analyze the PDF.
- C.** Configure content=markdown when you analyze the document.
- D.** Set the output_content_format=ContentFormat. MARKDOWN value.

Answer: D

The correct answer is D. Set the output_content_format=ContentFormat.MARKDOWN value . Azure Document Intelligence Layout API can return extracted document content in Markdown format, preserving semantic structure such as headings, paragraphs, sections, tables, and other layout elements. Microsoft's Document Intelligence layout guidance shows the Python SDK pattern for analyzing a document with the prebuilt-layout model and setting output_content_format=ContentFormat.MARKDOWN in the begin_analyze_document call. The Markdown output is returned in the top-level content section of the analysis result.

This minimizes development effort because the service produces structure-preserving Markdown directly, rather than requiring custom post-processing to reconstruct sections and table formatting from raw OCR spans. Microsoft's Markdown output documentation states that specifying Markdown output produces semantically structured content that maintains paragraphs, headings, tables, and other document elements in their proper hierarchy.

Option A only changes validation behavior and does not generate Markdown. Option B requests figures, not structured Markdown. Option C uses an incorrect parameter name; the documented SDK setting is output_content_format, not content. Reference topics: Azure Document Intelligence Layout API, Markdown output, PDF analysis, table extraction, and Foundry Tools document processing.

QUESTION: 18

You are planning a Microsoft Foundry project named Project1 that will contain multiple agents. Each agent will access the same Azure AI Search resource. You need to recommend a solution to centrally manage the Azure AI Search credentials within Project1. The solution must be implemented across all the agents. What should you recommend?

- A.** Enable role-based access control (RBAC) for the Azure AI Search resource.

- B.** Add a connection to the Azure AI Search resource.
- C.** Disable key-based access control on the Azure AI Search resource.
- D.** Create a managed private endpoint that connects to the Azure AI Search resource.

Answer: B

The correct recommendation is B. Add a connection to the Azure AI Search resource . Microsoft Foundry project connections are used to centrally define access from a project to external resources, including Azure AI Search. The official connection guidance states that you can add a connection by selecting an external service such as Azure AI Search , choosing the resource, and selecting the authentication method for that resource. This creates a reusable project-level configuration rather than requiring each agent to store or duplicate search credentials independently.

For Foundry agents that use Azure AI Search, the Azure AI Search tool requires a `project_connection_id`, which is the resource ID of the project connection to Azure AI Search. This allows multiple agents to reference the same managed connection while using the configured endpoint and authentication settings consistently. RBAC may be part of a keyless authentication design, but it does not by itself create a centrally managed Foundry credential configuration. Disabling key-based access improves security posture but does not connect the agents. A managed private endpoint addresses network isolation, not credential centralization.

Reference topics: Microsoft Foundry project connections, Azure AI Search tool, project connection IDs, authentication, and agent grounding.

QUESTION: 19

You have a Microsoft Foundry project that uses Azure AI Search to ground an agent in internal documentation.

After a recent content update, users report that the agent ' s answers have become less accurate.

You need to identify whether the retrieved content is negatively influencing the model ' s generated responses.

Which observability signal should you review?

- A.** prediction drift metrics
- B.** groundedness evaluation metrics
- C.** latency breakdown traces

D. indexer status and failure history

Answer: B

The correct observability signal is B. groundedness evaluation metrics . In a RAG solution, the key diagnostic question is whether the generated answer is supported by the retrieved context. Microsoft Foundry' s built-in evaluator reference defines Groundedness as the metric that measures how grounded the response is in the retrieved context, with scoring that indicates whether the model's claims are supported by the provided source material. This matches the issue after a content update. If retrieved chunks are stale, misleading, incomplete, or poorly aligned with the user query, groundedness results can show that generated responses are not reliably supported by the retrieved documentation. The RAG evaluator guidance explains that groundedness focuses on whether the response avoids content outside the grounding context, while other process metrics such as retrieval evaluate how relevant the retrieved chunks are. Latency traces are useful for performance troubleshooting, not response accuracy. Indexer status can reveal ingestion failures, but it does not show whether retrieved content is influencing generated answers negatively. Prediction drift is a model monitoring concept and is not the primary signal for RAG grounding quality. Reference topics: Microsoft Foundry observability, RAG evaluators, groundedness, retrieved context, and response quality evaluation.

QUESTION: 20

You have an Azure Speech in Foundry Tools resource that hosts a custom speech to text model deployed to a custom endpoint. An agent uses the endpoint to perform real-time speech recognition.

You are approaching the expiration date of the custom speech to text model.

What is the expected behavior when the model expires?

- A.** Speech recognition requests will fall back to the most recent base model for the same locale.
- B.** Speech recognition requests will continue to use the expired custom model until the model is removed manually.
- C.** Speech recognition requests will return a 4xx error until a new custom model is deployed.
- D.** The custom model will be deleted automatically when the model expires.

Answer: A

The correct answer is A. Speech recognition requests will fall back to the most recent base model for the same locale . Microsoft's custom speech model lifecycle guidance states that when a model expires, it is no longer available for transcription. For the custom endpoint route, speech recognition requests fall back to the most recent base model for the same locale. The documentation also warns that recognition results might still be returned, but the transcription may no longer reflect the domain-specific adaptation of the custom model.

This distinguishes real-time custom endpoint behavior from batch transcription behavior. Batch transcription requests that specify an expired model fail with a 4xx error, but that is not the route described in this question.

The agent is using a custom endpoint for real-time recognition, so fallback to the latest base model is the expected behavior. The model is not automatically deleted merely because it expires, and it does not continue to use the expired custom model indefinitely. The operational recommendation is to update the endpoint's model before expiration by redeploying the endpoint with a newer custom model. Reference topics: Azure Speech custom speech model lifecycle, custom endpoints, model expiration, real-time speech recognition, and endpoint redeployment.

QUESTION: 21

You need to configure Agent1 to answer customer questions about only the Contoso products. The solution must meet the business requirements.

What should you do?

- A.** Apply top-p sampling.
- B.** Modify the system message instructions.
- C.** Add few-shot examples.
- D.** Increase the value of the temperature parameter.

Answer: B

The correct answer is B. Modify the system message instructions . The case study states that Agent1 answers general questions about Contoso products and that the business requirement is for Agent1 to answer questions only about the products sold by Contoso . This is a behavioral boundary for the agent, so it should be implemented in the highest-priority instructions that define the agent's role, allowed scope, and refusal behavior. Microsoft Foundry guidance states that a system message is used to steer model behavior,

define the assistant's role and boundaries, and add safety or quality constraints for the scenario. The system message should instruct Agent1 to answer only when the question concerns Contoso products, use the configured Contoso product documentation as grounding, and politely refuse or redirect questions about non-Contoso products. Top-p sampling and temperature control randomness, not business-domain scope. Increasing temperature would make responses less deterministic. Few-shot examples can support desired behavior, but examples alone are weaker than explicit system-level instructions for defining operating boundaries. Reference topics: system message design, prompt engineering, agent instructions, response constraints, and grounded generative AI behavior.

QUESTION: 22

You have a Microsoft Foundry project that contains an internal Q & A agent.

Users report the following issues when they ask the agent questions:

- * An increase in the following response: "No relevant information found"**
- * Periodic HTTP 429 rate limit exceeded errors during peak hours**

You need to identify whether each issue is caused by model unavailability, resource limits, or inference failures.

What should you do? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Metrics to enable:

- Model Availability Rate and Provisioned Utilization
- Only Tokens Cache Match Rate
- Only Total Requests filtered to status code 200
- Time To Response and Total Tokens

Diagnostic log to collect:

- AllMetrics
- audit
- RequestResponse
- trace

Answer:

Metrics to enable:

Diagnostic log to collect:

Explanation:

Metrics to enable: Model Availability Rate and Provisioned Utilization

Diagnostic log to collect: trace

The correct metrics are Model Availability Rate and Provisioned Utilization . Model Availability Rate identifies whether failures are caused by service-side model unavailability because it is calculated from total calls minus server errors. Provisioned Utilization identifies resource-limit conditions because Microsoft states that when utilization reaches or exceeds 100%, calls are throttled and HTTP 429 errors are returned. This directly maps to the reported peak-hour rate-limit errors.

The correct diagnostic log is trace . For an internal Q & A scenario, trace logging is required to analyze question-answering behavior, including unanswered questions. Microsoft's custom question answering analytics guidance states that diagnostic logging stores telemetry and chat logs and instructs enabling Trace in addition to Audit, RequestResponse, and AllMetrics. The sample Kusto queries for unanswered questions inspect the answer, question, score, and knowledge base ID, and filter unanswered results by a score of zero.

RequestResponse is useful for request status and latency, and Audit is for administrative operations, but neither is the best diagnostic category for analyzing Q & A inference behavior such as "No relevant information found." Reference topics: Foundry monitoring, model availability, provisioned utilization, diagnostic logging, custom question answering analytics, and trace logs.

QUESTION: 23

You have an application that processes scanned PDF invoices. The invoices have

varied layouts and include multipage tables.

You have a pipeline that uses optical character recognition (OCR) and extracts totals and invoice numbers.

The results are often incorrect because the document structure is ignored.

You need to implement a solution that provides OCR, layout analysis, and template-generalizing field extraction. The solution must NOT require training a custom model.

The solution must minimize administrative effort.

What should you include in the solution?

- A. an Azure Machine Learning model**
- B. Azure Content Understanding in Foundry Tools**
- C. Azure Language in Foundry Tools**

Answer: B

The correct answer is Azure Content Understanding in Foundry Tools . The scenario requires more than basic OCR because scanned invoices have varied layouts and multipage tables. Content Understanding is designed for intelligent document processing and provides OCR, layout detection, table extraction, field extraction, confidence scores, and grounding in a managed service. Microsoft describes Content Understanding as a service that transforms unstructured content into structured outputs and supports invoice processing by extracting and validating fields from complex documents.

This also meets the requirement to avoid training a custom model. Content Understanding includes prebuilt and domain-specific analyzers, including invoice and procurement-style document processing, and Microsoft states that these analyzers provide structured extraction without custom training. It generalizes across visual template variations by using semantic document categories rather than requiring separate models per invoice layout. Azure Machine Learning would increase administrative effort because it requires model development, training, deployment, and monitoring. Azure Language is optimized for text analytics tasks such as classification and entity extraction after text is available, but it does not provide document layout analysis or multipage table structure extraction. Reference topics: Content Understanding, intelligent document processing, OCR, layout analysis, analyzers, field schemas, and structured extraction.

QUESTION: 24

You have a Microsoft Foundry project that contains an agent and an image generation model deployment.

The agent generates original images from user-supplied product photos.

You need to ensure that the generated images maintain the product identity and visual characteristics of the provided photo.

What should you do?

- A.** Set the input_fidelity parameter to high.
- B.** Apply a groundedness detection filter.
- C.** Include a prompt and input image in the request.
- D.** Decrease the value of the temperature parameter.

Answer: A

The correct answer is A. Set the input_fidelity parameter to high . The scenario requires the generated image to preserve the identity and visual characteristics of the user-supplied product photo. In Azure OpenAI image editing and generation workflows, input_fidelity controls how strongly the model attempts to match the style and features of the input image. Microsoft's documentation states that this parameter lets you make subtle edits without changing unrelated areas, and that high input fidelity preserves input-image features more accurately than standard mode.

Including a prompt and input image is necessary for image-guided generation, but it does not by itself maximize preservation of the product's appearance. The explicit preservation control is input_fidelity, and the requirement specifically asks to maintain product identity and visual characteristics. A groundedness detection filter applies to validating generated text against source data, not preserving visual features in image generation. Lowering temperature may reduce randomness in text generation, but it is not the image-control parameter used to retain product-specific visual details. Reference topics: Azure OpenAI image generation, image edit API, input images, input_fidelity, image-to-image generation, and visual identity preservation.

QUESTION: 25

You have a Microsoft Foundry project that contains a prompt agent used by a customer support web app.

The agent is invoked from a Python service that does NOT run in the Foundry portal.

You need to implement end-to-end tracing to capture latency breakdowns and exceptions across agent runs.

Which two components can you use? Each correct answer presents a complete

solution.

NOTE: Each correct selection is worth one point.

- A.** a Log Analytics workspace
- B.** OpenTelemetry
- C.** Application Insights
- D.** the Azure Monitor Agent
- E.** Microsoft Sentinel

Answer: B,C

The correct components are OpenTelemetry and Application Insights . Microsoft Foundry tracing for prompt agents is designed to capture detailed telemetry for agent execution, including latency, exceptions, prompt activity, and retrieval operations. For a Python service that invokes the agent outside the Foundry portal, OpenTelemetry is the appropriate instrumentation mechanism because the Foundry SDK tracing setup uses OpenTelemetry packages and Azure SDK tracing integration for client-side traces. This enables distributed tracing across the external Python service and the agent run. Application Insights is the telemetry backend used by Foundry tracing. The Foundry tracing setup requires an Azure Monitor Application Insights resource to store traces, and traces can then be viewed in Foundry or directly in Azure Monitor Application Insights. Application Insights also provides performance and failure investigation experiences for response times, slow transactions, errors, and exceptions.

A Log Analytics workspace may underlie Application Insights data storage, but it is not the agent tracing component itself. The Azure Monitor Agent collects machine and guest telemetry, not application-level agent traces from a Python SDK invocation. Microsoft Sentinel is a security information and event management solution, not an end-to-end agent tracing mechanism. Reference topics: Microsoft Foundry tracing, Azure Monitor Application Insights, OpenTelemetry instrumentation, prompt agent observability, and Python SDK telemetry.

QUESTION: 26

You have a Microsoft Foundry project that contains a customer support agent grounded in internal documentation.

After a recent update, users report the following issues:

- * Some answers are unsupported by retrieved documents.**
- * A small number of responses are flagged for policy violations.**

You need to evaluate each issue.

Which observability signals should you use for each issue? To answer, drag the appropriate observability signals to the correct issues. Each observability signal may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Observability signals

- Groundedness evaluation metrics
- Latency breakdown traces
- Risk and safety metrics
- Token usage analytics

Answer Area

Unsupported responses: Observability signal

Policy violations: Observability signal

Answer:

Observability signals

- Groundedness evaluation metrics
- Latency breakdown traces
- Risk and safety metrics
- Token usage analytics

Answer Area

Unsupported responses: Groundedness evaluation metrics

Policy violations: Risk and safety metrics

Explanation:

Unsupported responses: Groundedness evaluation metrics

Policy violations: Risk and safety metrics

For unsupported responses, use Groundedness evaluation metrics . In a Retrieval Augmented Generation scenario, the key question is whether the generated answer is supported by the retrieved context. Microsoft Foundry built-in evaluators define Groundedness as the RAG metric that measures how grounded a response is in retrieved context and returns a model-based score; Groundedness Pro evaluates whether the response is grounded in retrieved context by using Azure AI Content Safety. This directly matches answers that are unsupported by internal documentation.

For policy violations, use Risk and safety metrics . Microsoft Foundry risk and safety evaluators assess generated responses for safety risks such as hate and unfairness, sexual content, violence, self-harm, protected material, indirect attacks, code vulnerability,

ungrounded attributes, prohibited actions, and sensitive data leakage. The guidance states that these evaluators assign risk and safety severity or pass/fail outcomes for AI responses and agent behavior.

Latency breakdown traces diagnose performance, not correctness or policy compliance.

Token usage analytics diagnose cost and prompt/response size, not unsupported claims or safety violations. Reference topics:

Microsoft Foundry observability, RAG evaluators, groundedness, risk and safety evaluators, and agent quality evaluation.

QUESTION: 27

You have a Microsoft Foundry project that contains a deployed chat model.

You have a Python service that sends API requests to the model. The service is integrated with an automated validation system that compares generated outputs against approved response patterns.

Stakeholders report that small wording differences are causing validation mismatches.

You need to update the request parameters to improve output stability. The solution must maximize reasoning quality.

How should you complete the Python code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

```
message = client.messages.create(  
    model="deployment-name",  
    messages=[  
        {"role": "user", "content": "Summarize the release notes in 3 bullet points."}  
    ],  
    max_tokens=800,  
    temperature=  
        0  
        1  
        2  
    ,  
    thinking={"type": "enabled"},  
    output_config={"effort":  
        "high"  
        "low"  
        "medium"  
    }  
)
```

Answer:

```

message = client.messages.create(
    model="deployment-name",
    messages=[
        {"role": "user", "content": "Summarize the release notes in 3 bullet points."}
    ],
    max_tokens=800,
    temperature=
        0
        1
        2
    thinking={"type": "enabled"},
    output_config={"effort":
        "high"
        "low"
        "medium"
    }
)

```

Explanation:

temperature = 0

output_config = { " effort " : " high " }

The correct configuration is temperature = 0 and output_config = { " effort " : " high " }. The requirement is to reduce small wording variations that are breaking automated validation.

In chat completion requests, temperature controls sampling randomness. Microsoft's Azure OpenAI reference states that temperature ranges between 0 and 2, and that higher values make output more random while lower values make output more focused and deterministic. Therefore, the most stable setting from the available choices is 0, because it minimizes randomness and improves repeatability for validation-sensitive response patterns.

The solution must also maximize reasoning quality. The code already enables thinking with thinking={ " type

" : " enabled " }, so the remaining reasoning-quality control is the effort setting. Microsoft Foundry model guidance states that the effort parameter controls the quality/cost tradeoff and supports low, medium, and high effort levels. Selecting " high " maximizes reasoning quality among the available options.

Using temperature values of 1 or 2 would increase variability and make validation mismatches more likely.

Selecting low or medium effort would not meet the requirement to maximize reasoning quality. Reference topics: Microsoft Foundry model inference, chat model parameters, temperature, thinking, effort, and output stability.

QUESTION: 28

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You run an evaluation flow that scores responses for completeness and blocks responses that fall below a defined threshold.

Does this meet the goal?

A. Yes

B. No

Answer: B

The solution does not meet the goal. A completeness evaluation flow is useful for detecting incomplete responses, but detection and blocking do not improve the response itself.

Microsoft Foundry RAG evaluators define Response Completeness as a metric that measures whether a response covers all critical information from the expected response or ground truth. It is a system evaluation signal used to assess response quality and produce pass/fail or scored results.

In this scenario, the issue is that the agent omits required regulatory clauses even though the clauses are present in retrieved content. Blocking low-scoring responses would prevent incomplete answers from being returned, but it would not revise the summary, add the missing clauses, or improve the generation process.

The appropriate improvement is to add a response-generation control such as a reflection or verification pass that checks the draft summary against the retrieved policy content and

regenerates or amends the answer before returning it. Evaluation can support the quality gate, but by itself it is an assessment mechanism, not a completeness-enhancement mechanism. Reference topics: Microsoft Foundry RAG evaluators, response completeness, grounded generation, reflection, and response quality optimization.

QUESTION: 29

You need to configure an indexing pipeline for Agent1 to retrieve the relevant product information in storage1. The solution must meet the technical requirement.

Which two built-in skills should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Language Detection
- B. Entity Recognition
- C. Merge
- D. Azure OpenAI Embedding
- E. Text Split
- F. key phrase extraction

Answer: D,E

The correct built-in skills are Azure OpenAI Embedding and Text Split. The case study requires an indexing pipeline that enables semantic and vector search over the product sheets stored in Azure Blob Storage, so Agent1 can retrieve relevant product information for natural language customer questions. For a RAG pipeline, long PDF content must first be broken into retrievable chunks, and each chunk must then be vectorized for semantic similarity retrieval.

Microsoft's Azure AI Search integrated vectorization guidance states that you create a skillset that calls the Text Split skill for chunking and the Azure OpenAI Embedding skill to vectorize the chunks. The Text Split skill breaks text into chunks and provides positional metadata, making it suitable when downstream embedding skills have input-length limits. The Azure OpenAI Embedding skill connects to an embedding model deployed in Azure OpenAI or a Microsoft Foundry project and generates embeddings during indexing. Language Detection, Entity Recognition, and key phrase extraction can enrich text, but they do not create vector embeddings. Merge is useful for combining OCR text with document text, but it does not satisfy the core vector-search requirement. Reference topics: Azure AI

Search skillsets, Text Split skill, Azure OpenAI Embedding skill, integrated vectorization, and RAG indexing.

QUESTION: 30

You have an app named App1 that uses a Microsoft Foundry multimodal model deployment.

App1 runs optical character recognition (OCR) on uploaded images and appends the OCR output to the prompt as additional context.

Some uploaded images contain embedded text.

You need to prevent potentially malicious instructions from being processed by the model.

What should you use?

- A.** protected material text
- B.** prompt shields for user prompts
- C.** image moderation
- D.** prompt shields for documents

Answer: D

The correct answer is D. prompt shields for documents . The OCR text is extracted from uploaded images and appended as additional context, so it is third-party content rather than a direct trusted user instruction.

Microsoft Foundry Prompt Shields distinguish between user prompt attacks , which are malicious instructions directly supplied as user prompts, and document attacks , which are hidden or embedded instructions in third-party content such as documents, emails, webpages, or grounded data. Microsoft also notes that Prompt Shields analyze indirect attacks embedded in input documents or images.

This scenario is an indirect prompt injection pattern: an uploaded screenshot may contain text that says to ignore the system prompt or perform unauthorized actions. Because the OCR output is being passed to the multimodal model as contextual content, the appropriate protection is Prompt Shields for documents, which scans externally sourced context for malicious embedded instructions before generation. Protected material text detects known copyrighted text in model outputs, not prompt injection. Image moderation detects harmful visual content categories, not malicious instructions in extracted context. Prompt Shields for user prompts are less precise here because the risk comes from document-like content derived from the uploaded image.

Reference topics: Prompt Shields, document attacks, indirect prompt injection, multimodal safety, OCR- derived context, and Foundry guardrails.

QUESTION: 31

You have a Microsoft Foundry project that contains a customer support agent. The agent calls an internal knowledge API tool before generating responses.

Users report the following issues:

- * Some requests take more than 15 seconds to complete.**
- * Some responses are incorrect, even when the knowledge API returns the expected data.**

You need to inspect individual agent runs to view the ordered sequence of large language model (LLM) calls, tool invocations, and timing information.

Which observability capability should you use?

- A. token usage**
- B. safety metrics**
- C. tracing**
- D. monitoring**

Answer: C

The correct capability is tracing because the requirement is to inspect the execution path of an individual agent run. Microsoft Foundry tracing captures detailed telemetry for agent behavior, including LLM calls, tool invocations, agent decision flows, inputs, outputs, tool results, token consumption, duration, and latency.

This is the appropriate observability mechanism when you need to determine which step introduced a delay, whether the agent called the internal knowledge API, what data the tool returned, and how the model used that data before producing the final response.

Microsoft's Foundry observability guidance describes distributed tracing as the mechanism that provides visibility into LLM calls, tool invocations, agent decisions, and inter-service dependencies.

Token usage is useful for cost analysis and prompt optimization, but it does not show ordered run steps or tool-call sequencing. Safety metrics evaluate risk-related output behavior, not latency or tool execution.

General monitoring provides aggregate health, latency, success-rate, and dashboard views, but the question asks for per-run sequence inspection and timing breakdowns. Foundry agent tracing specifically supports debugging unexpected behavior and monitoring latency

across requests. Reference topics: Microsoft Foundry observability, agent tracing, OpenTelemetry-based traces, tool invocations, LLM call inspection, and latency diagnostics.

QUESTION: 32

You have a Microsoft Foundry project that contains a high-traffic agent.

After a recent update, operational costs increase significantly.

Monitoring confirms that the volume of user traffic to the agent remains unchanged.

You suspect that changes to the request or response characteristics are causing the increase.

You need to identify whether the additional costs are driven by the model input size, the model output size, or expanded tool usage.

Which observability capability should you use?

- A.** evaluation metrics
- B.** token usage
- C.** latency
- D.** run success rate

Answer: B

The correct capability is token usage . In Microsoft Foundry observability, token consumption is the primary signal for diagnosing model-cost changes when request volume is unchanged. Token usage lets you distinguish whether costs increased because prompts became larger, retrieved or tool-provided context expanded, responses became longer, or agent execution added more model calls. Microsoft Foundry monitoring dashboards track operational metrics such as token consumption, latency, error rates, and quality scores, and the agent monitoring dashboard is specifically intended to help analyze token usage, latency, success rates, and evaluation outcomes for production traffic.

This directly matches the scenario because the issue is not more traffic, but changed request or response characteristics. Input tokens reveal whether the prompt, chat history, grounding data, or tool outputs being sent to the model increased. Output tokens reveal whether the model is generating longer completions.

Expanded tool usage can also increase cost indirectly by adding more tool results, intermediate calls, and context into subsequent model requests; Foundry tracing and observability capture tool usage and token consumption for agent runs.

Evaluation metrics assess response quality and safety, not cost drivers. Latency identifies

performance delays, and run success rate measures reliability. Reference topics: Microsoft Foundry observability, agent monitoring dashboard, token consumption, cost analysis, tool usage, and production monitoring.

QUESTION: 33

You have a Microsoft Foundry project that contains an agent.

The agent uses tools to retrieve internal content and call external APIs. The agent is configured to let the model decide when to call the tools.

You need to publish the agent for a compliance workflow. The solution must meet the following requirements:


- * Each workflow run must include a retrieval step before generating a response.**
- * Tool calls must authenticate by using the published agent's own identity.**
- * Tool access must use an identity isolated from other project resources.**
- * Tool access must support audit tracing.**

What should you do? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

The screenshot shows a configuration interface for a Microsoft Foundry agent. It features two dropdown menus. The first dropdown, labeled 'Set tool_choice to:', has three options: 'auto' (selected), 'none', and 'required'. The second dropdown, labeled 'Configure the tool to authenticate by:', has three options: 'Storing API keys in prompts', 'Using the shared project agent identity', and 'Using a distinct agent identity bound to the client application'.

Answer:



Explanation:

Set tool_choice to: required

Configure the tool to authenticate by: Using a distinct agent identity bound to the client application Set tool_choice to required because the compliance workflow must deterministically include a tool-based retrieval step before the agent generates a response.

Microsoft Foundry Agent Service guidance states that tool_choice provides the most deterministic control over tool use: auto lets the model decide, none prevents tool calls, and required forces the model to call one or more tools. This directly corrects the current nondeterministic behavior where the model decides whether to call tools.

For authentication, use a distinct agent identity bound to the client application . Microsoft Foundry creates a shared identity for unpublished or in-development agents, but publishing an agent automatically creates a dedicated agent identity blueprint and agent identity associated with the agent application resource. Published agents authenticate tool calls by using that unique agent identity, and RBAC permissions must be assigned to the new identity. This provides isolation from the broader shared project identity and supports independent audit trails for compliance workflows.

Storing API keys in prompts violates security guidance and prevents robust audit attribution. The shared project agent identity is easier for development, but it has a broader blast radius and does not meet the isolation requirement. Reference topics: Foundry Agent Service tool choice, tool authentication, published agent identities, RBAC, and auditability.

QUESTION: 34

You have a Microsoft Foundry project that contains an agent.

You need to process mixed-format documents that contain scanned text, tables, and multicolumn layouts. The extracted content must preserve the document structure and be converted into the Markdown format for downstream reasoning.

What should you configure first?

- A. an Azure Content Understanding in Foundry Tools analyzer
- B. a generative chat completion request
- C. an Azure Language in Foundry Tools text analysis model deployment
- D. an Azure OpenAI Responses API call that uses a multimodal model

Answer: A

The correct answer is A. an Azure Content Understanding in Foundry Tools analyzer . In Azure Content Understanding, an analyzer is the reusable configuration that defines what content type is processed, which elements are extracted, and how the output is structured. Microsoft's analyzer guidance states that analyzers define extraction for text, layout, tables, fields, and output formats such as Markdown and JSON. This makes the analyzer the first required configuration step before downstream reasoning or agent orchestration can use the extracted content.

This matches the requirement because the documents contain scanned text, tables, and multicolumn layouts.

Content Understanding document analysis is designed to transform unstructured documents into structured, machine-readable output while preserving document structures and relationships. Its Markdown representation converts unstructured documents into GitHub Flavored Markdown while maintaining content and layout for downstream use.

A generative chat completion request or Azure OpenAI Responses API call could reason over extracted content, but it is not the correct first step for OCR, layout preservation, and Markdown conversion. Azure Language focuses on text analysis after content has already been extracted. Reference topics: Content Understanding analyzers, document extraction, OCR, layout analysis, Markdown output, and downstream reasoning.

QUESTION: 35

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You increase the value of the max_tokens parameter.

Does this meet the goal?

A. Yes

B. No

Answer: B

The solution does not meet the goal. Increasing max_tokens only raises the maximum number of tokens the model is allowed to generate. Microsoft's Azure OpenAI reference defines max_tokens as the maximum number of tokens allowed for the generated answer, and the quota guidance notes that increasing it can help when responses are being truncated.

In this scenario, the problem is not described as output truncation. The required regulatory clauses are already present in the retrieved policy documents, but the agent omits them during summarization. That is a response completeness issue: Microsoft Foundry RAG evaluator guidance defines response completeness as the recall aspect of the response, meaning the response should not miss critical information compared with expected content or ground truth.

A larger token budget might permit a longer answer, but it does not force the model to identify, verify, or include each mandatory clause. It can also increase cost and latency. The appropriate control is a reflection or completeness verification pass that checks the draft against the retrieved policy clauses and regenerates or revises the response when required content is missing. Reference topics: RAG response completeness, model output limits, max_tokens, reflection, and response validation.

QUESTION: 36

You are deploying a support agent that enables users to upload photos.

You need to automatically classify uploaded images for harmful content. The solution must block content based on severity levels.

What should you do?

A. Implement image moderation.

B. Enable prompt shields.

- C. Apply keyword scanning to optical character recognition (OCR) output by using Azure Vision in Foundry Tools.
- D. Use blocklists.

Answer: A

The correct answer is A. Implement image moderation . Azure AI Content Safety provides image analysis that classifies uploaded images for harmful content, including harm categories such as hate, sexual content, violence, and self-harm. Microsoft's Content Safety overview states that the Analyze Image API scans images for harmful content with multi-severity levels, which directly matches the requirement to automatically classify uploaded photos and block content based on configured severity thresholds.

Prompt Shields are intended to detect prompt injection and jailbreak-style attacks against generative models, not to classify image harm categories. Keyword scanning OCR output would only detect visible text extracted from the image and would miss visual harm in the image itself. Blocklists can help match known words or custom patterns, but they are not a complete image safety classifier and do not provide the built-in severity- based image harm classification required here. Image moderation is therefore the correct control for user-uploaded photos. Reference topics: Azure AI Content Safety, image moderation, harm categories, severity levels, Foundry guardrails, and responsible AI controls.

QUESTION: 37

You have a Microsoft Foundry project that contains an agent.

The agent ingests scanned PDF vendor invoices that contain tables and embedded QR codes.

The agent must preserve the PDF layout in the extracted output to ensure that downstream processing can reference sections and tables.

You plan to call Azure Content Understanding in Foundry Tools.

You need to extract content and layout elements and detect QR codes without requiring a language model deployment.

Which built-in analyzer should you use?

- A. prebuilt-layout
- B. prebuilt-documentFieldSchema
- C. prebuilt-read
- D. prebuilt-documentSearch

Answer: A

The correct built-in analyzer is prebuilt-layout because the requirement is to preserve document layout while extracting content from scanned PDFs. Microsoft's Content Understanding prebuilt analyzer guidance states that prebuilt-layout extracts content and layout elements such as words, figures, paragraphs, and tables, identifies document structure including sections and formatting, and provides detailed layout information beyond basic text extraction. It also states that prebuilt-layout does not require a language model or embedding model, which directly satisfies the no language model deployment requirement.

QR codes are handled through barcode extraction. The analyzer configuration reference states that enableBarcode detects and extracts barcodes and QR codes, returns decoded values, and supports QR Code and Micro QR Code among other barcode types. This option is supported by document-based analyzers, making it compatible with layout-oriented document processing.

prebuilt-read is insufficient because it provides OCR and barcode extraction but foundational text extraction without layout analysis. prebuilt-documentSearch is optimized for RAG ingestion and semantic analysis, which is broader than required. prebuilt-documentFieldSchema proposes extraction schemas rather than extracting full document layout. Reference topics: Content Understanding prebuilt analyzers, layout analysis, OCR, barcode detection, QR code extraction, and document-based analyzers.

QUESTION: 38

You need to recommend a solution to support the planned changes and technical requirements for Agent1 to use the product information stored in storage1.

What should you include in the recommendation?

- A.** Azure AI Search
- B.** Azure Translator in Foundry Tools
- C.** Azure Document Intelligence in Foundry Tools
- D.** Grounding with Bing Search

Answer: A

The correct recommendation is Azure AI Search. The case study states that the product detail sheets are stored as PDFs in storage1, and that Agent1 must be enabled to retrieve

and use detailed product information from those sheets. It also specifies that the indexing pipeline must enable semantic and vector search, and that Agent1 must answer natural language questions about product details by using the product sheet information.

Azure AI Search is the Azure service designed to ingest content from sources such as Azure Blob Storage, create searchable indexes, and support keyword, semantic, hybrid, and vector retrieval for Retrieval Augmented Generation (RAG) solutions.

Microsoft's Azure AI Search guidance states that integrated vectorization can chunk content and generate embeddings during indexing, enabling vector search over source documents. It also states that Azure AI Search supports text and vector queries and can improve raw content for search-related scenarios through enrichment pipelines. Azure Translator is unrelated to retrieval. Document Intelligence can extract document structure, but it is not the retrieval index for Agent1. Grounding with Bing Search retrieves public web content, not Contoso's private PDFs in storage1. Reference topics: Azure AI Search, RAG, semantic search, vector search, Azure Blob Storage indexing, and agent grounding.

QUESTION: 39

You have a Microsoft Foundry project that processes procurement documents submitted by suppliers.

You need to implement two pipelines by using Azure Content Understanding in Foundry Tools. The solution must meet the following requirements:

*** Include a pipeline named Pipeline1 that supports cost-effective, high-volume processing of standalone PDF invoices.**


*** Include a pipeline named Pipeline2 that supports cross-document validation by using multi-step reasoning and reference data.**

How should you configure each pipeline? To answer, drag the appropriate configurations to the correct pipelines. Each configuration may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Configurations	Answer Area
⋮ Multi-file task in pro mode	Pipeline1: Configuration
⋮ Multi-file task in standard mode	Pipeline2: Configuration
⋮ Single-file task in pro mode	
⋮ Single-file task in standard mode	

Answer:



Explanation:

Pipeline1: Single-file task in standard mode

Pipeline2: Multi-file task in pro mode

Pipeline1 should use a single-file task in standard mode because the workload is high-volume processing of standalone PDF invoices. Azure Content Understanding standard mode is intended for individual files that require straightforward structured extraction, and Microsoft describes it as minimizing cost and latency for broad, data-centric processing scenarios. This makes it the best fit for cost-effective invoice extraction where each PDF can be processed independently.

Pipeline2 should use a multi-file task in pro mode because the requirement includes cross-document validation, multi-step reasoning, and reference data. Microsoft guidance states that pro mode is designed for advanced scenarios requiring multi-step reasoning and cross-file analysis, including processing multiple input files in a single request, validating or enriching data across documents, and using reference data to guide extraction and validation.

Single-file pro mode would add unnecessary capability for Pipeline1 and would not optimize for cost-effective high-volume standalone processing. Multi-file standard mode does not meet the pro-mode requirement for reference-data-based reasoning. Reference topics: Azure Content Understanding standard mode, pro mode, single-file tasks, multi-file tasks, field extraction, and procurement document validation.

QUESTION: 40

You have a Microsoft Foundry project named Project1 that contains the following:

*** An OpenAPI tool that calls an external API**

*** A project connection named Connection1 that stores the API key of the external API**

When an agent calls the OpenAPI tool, the API returns a 401 unauthorized error, and traces show that the API key header is NOT being sent.

You need to ensure that the OpenAPI tool automatically includes the API key from

Connection1 on all requests.

What should you do?

- A.** Configure the tool to use the default connection of Project1.
- B.** Connect the tool to Connection1.
- C.** Enable identity passthrough so that the tool uses the Microsoft Entra token of the caller.
- D.** Add the API key header manually to the OpenAPI specification.

Answer: B

The correct action is to connect the OpenAPI tool to Connection1 . In Microsoft Foundry Agent Service, an OpenAPI tool does not automatically use every project connection in the project. For API key authentication, the tool must be explicitly configured to use the project connection that contains the required secret.

Microsoft's OpenAPI tool guidance states that API key or token authentication requires a project connection configured with the API key or token, and the tool is then created or configured to use that connection for authentication. The documentation also notes that the agent calls the external API by using the stored API key when the OpenAPI tool is configured with the project connection.

Option A is incorrect because a project's default connection is not automatically the authentication source for a specific OpenAPI tool. Option C is incorrect because identity passthrough or managed identity only applies when the target API accepts Microsoft Entra ID tokens; this scenario uses an external API key. Option D is also incorrect as a credential value should not be manually embedded in the OpenAPI specification. The specification defines the security scheme, while the secret value is stored in the Foundry connection.

Reference topics: OpenAPI tools, project connections, API key authentication, tool authentication configuration, and agent tracing.

QUESTION: 41

You have a Microsoft Foundry project that contains an agent used by the financial analysts at your company.

You need to optimize the agent workflow by providing additional data access and processing capabilities. The solution must meet the following requirements:

- * Ensure that the agent can perform calculations during conversations**
- * Ensure that the agent can access up-to-date information from public websites.**
- * Ensure that the agent can retrieve information from documents uploaded directly to the agent.**

What should you use for each requirement? To answer, drag the appropriate tools to the correct requirements.

Each tool may be

used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Tools	Answer Area
☐ Code interpreter	Access up-to-date information from public websites: Tool
☐ Computer use	Perform calculations during conversations: Tool
☐ File search	Retrieve information from documents uploaded directly to the agent: Tool
☐ Grounding with Bing Search	
☐ Microsoft Fabric	

Answer:

Tools	Answer Area
☐ Code interpreter	Access up-to-date information from public websites: Grounding with Bing Search
☐ Computer use	Perform calculations during conversations: Code interpreter
☐ File search	Retrieve information from documents uploaded directly to the agent: File search
☐ Grounding with Bing Search	
☐ Microsoft Fabric	

Explanation:

Access up-to-date information from public websites: Grounding with Bing Search
Perform calculations during conversations: Code interpreter
Retrieve information from documents uploaded directly to the agent: File search
The correct tool for public, current web information is Grounding with Bing Search . Microsoft Foundry Agent Service identifies Grounding with Bing Search as the built-in tool that enables an agent to access and return information from the internet, which fits the requirement for up-to-date public website data. (learn.

microsoft.com)

For calculations during conversations, use Code interpreter . Microsoft's Foundry guidance states that Code Interpreter enables an agent to run Python code in a sandboxed execution environment and solve data analysis and math tasks iteratively. This is the correct fit for financial analysts who need calculations, analysis, and potentially chart generation during the conversation.

For documents uploaded directly to the agent, use File search . Microsoft describes File Search as the tool that enables Foundry agents to search through documents, retrieve

relevant information, and augment model responses with knowledge from uploaded files such as PDFs, Word documents, and proprietary content.

Computer use is for interacting with graphical applications, not calculation or document retrieval. Microsoft Fabric is for enterprise data and analytics integration, not direct uploaded document retrieval. Reference topics: Foundry Agent Service tools, Code Interpreter, File Search, and Grounding with Bing Search.

QUESTION: 42

You need to configure Agent1 to meet the security and compliance requirements.

What should you use?

- A.** prompt shields
- B.** Personally Identifiable Information (PII) Detection
- C.** self-harm content filtering
- D.** violence content filtering

Answer: B

The correct answer is B. Personally Identifiable Information (PII) Detection . The case study states that Agent1 must never reveal customer information , even if a document containing customer data is added accidentally to the product sheet repository in storage1. This is a privacy and compliance control requirement, so the appropriate capability is PII Detection. Azure Language PII Detection is a Foundry Tools capability that identifies, classifies, and redacts sensitive information across text, conversations, and native documents. Microsoft states that PII Detection can be used to implement privacy controls, reduce sensitive data exposure, and support compliance requirements. In this scenario, PII Detection should be applied to retrieved product-sheet content and generated responses so customer names, contact details, identifiers, and other sensitive values are not exposed to users.

Prompt Shields are important for a separate requirement: protecting Agent1 from malicious instructions hidden in documents or embedded text. Microsoft describes Prompt Shields for documents as protection against hidden instructions embedded in external content. However, the option that directly satisfies the requirement to prevent disclosure of customer information is PII Detection. Self-harm and violence filters address harmful-content categories, not privacy leakage.

QUESTION: 43

You need to configure the model deployment for Agent1 to meet the technical requirements.

What should you configure? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Deployment type:



A screenshot of a dropdown menu for 'Deployment type'. The menu is open, showing three options: 'Standard', 'Global Standard', and 'Global Provisioned'. The 'Standard' option is highlighted with a green dashed box.


Version update policy:



A screenshot of a dropdown menu for 'Version update policy'. The menu is open, showing three options: 'Once the current version expires', 'Opt out of automatic model version upgrades', and 'Upgrade once a new default version becomes available'. The 'Once the current version expires' option is highlighted with a green dashed box.

Answer:

Deployment type:



A screenshot of a dropdown menu for 'Deployment type'. The menu is open, showing three options: 'Standard', 'Global Standard', and 'Global Provisioned'. The 'Standard' option is highlighted with a green dashed box.

Version update policy:



A screenshot of a dropdown menu for 'Version update policy'. The menu is open, showing three options: 'Once the current version expires', 'Opt out of automatic model version upgrades', and 'Upgrade once a new default version becomes available'. The 'Once the current version expires' option is highlighted with a green dashed box.

Explanation:

Deployment type: Standard

Version update policy: Once the current version expires

The correct deployment type is Standard . The case study specifies that Project1 is deployed in an EU Azure region and that model-processed data must remain within the EU. It also requires scalable, high-throughput generative AI workloads that dynamically handle variable customer support traffic without reserved throughput capacity. In Microsoft Foundry Models, Standard is a pay-per-token deployment type that processes data in a single Azure region, while Global Standard can process requests across regions and Global Provisioned uses reserved provisioned throughput. Microsoft's deployment-type guidance

identifies Standard as single-region, pay-per-token, whereas Global Provisioned is cross-region with reserved capacity.

The correct version update policy is Once the current version expires . This keeps Agent1 on the selected model version during its supported lifecycle, which supports stable and consistent responses, but still preserves continuity by automatically moving to a supported replacement when the current version is retired.

Microsoft's model versioning guidance states that this policy updates only when the current model version expires, while upgrading when a new default is available changes the deployment sooner and opting out can cause the deployment to stop working after retirement. Reference topics: deployment types, regional data processing, model versioning, throughput capacity, and stable production deployments.

QUESTION: 44

You have a Microsoft Foundry project that contains an agent. The agent uses Azure AI Search as the retriever.

You plan to ingest PDFs into an Azure AI Search index to ensure that the agent can ground responses in texts in both documents and embedded images.

Users require citations that link to the source files.

You need to ensure that during indexing, the images are extracted into a structure that can be used as input for the built-in optical character recognition (OCR) skill.

Which indexing approach should you use?

- A.** a skillset to run the OCR skill directly against the content field of the index
- B.** the outputFieldMappings parameter to write image data to a searchable field
- C.** an indexer to extract image data into a normalized_images collection
- D.** a Shaper skill to restructure the OCR input

Answer: C

The correct indexing approach is to use an indexer to extract image data into a normalized_images collection . In Azure AI Search enrichment pipelines, embedded images in PDFs are not passed directly from the text content field into OCR. Instead, the indexer must perform document cracking and image extraction by enabling the indexer image action. Microsoft's Azure AI Search documentation states that image-processing skills such as OCR and image analysis expect normalized images, and that enabling imageAction causes embedded images to be extracted and normalized for downstream skills.

The OCR skill is designed to receive image input from `/document/normalized_images/*`. Microsoft's skillset tutorial specifically states that the OCR skill assumes a `normalized_images` field exists and that this field is generated by setting the `indexer` `imageAction` configuration to `generateNormalizedImages`. The document extraction skill reference also confirms that `generateNormalizedImages` creates an array of normalized images during document cracking for OCR and image analysis.

Option A is incorrect because OCR does not run directly against the index content field.

Option B maps outputs after enrichment; it does not extract images. Option D reshapes data but does not create the required normalized image collection. Reference topics: Azure AI Search indexers, AI enrichment, OCR skill, `imageAction`, and `normalized_images`.

QUESTION: 45

You are creating an image-editing workflow in a Microsoft Foundry project.

The workflow must meet the following requirements:

- * Ensure that background objects can be removed by applying a mask-based inpainting edit.**

- * Preserve the original lighting and style of the edited images.**

- * Use the built-in image editing controls, NOT a custom model.**

You need to ensure that image edits apply exclusively inside the masked area.

How should you configure the workflow?

A. Enable `text_to_image` mode and a prompt describing the desired background removal.

B. Set generation mode to `image_variation` and provide the original image as a reference.

C. Enable `image_to_image` mode and a high-strength value to regenerate the full image based on the prompt.

D. Enable `mask_inpainting` and supply both the input image and a mask indicating which part of the image to modify.

Answer: D

The correct configuration is D. Enable `mask_inpainting` and supply both the input image and a mask indicating which part of the image to modify. The requirement is not to generate a new image, but to edit a specific region of an existing image while preserving the surrounding lighting, composition, and style. Azure OpenAI image editing in Microsoft Foundry supports modifying existing images by submitting an input image plus a prompt. For masked edits, the mask explicitly defines the part of the image the model is allowed to change; Microsoft states that the mask parameter defines the area to edit and must match

the input image dimensions.

text_to_image would create a new image from a prompt and cannot guarantee preservation of the original image. image_variation generates related variants rather than targeted removals. image_to_image with high strength can regenerate broader areas and may alter unrelated visual details. Mask-based inpainting is the built-in editing control that limits modification to the selected region. Reference topics: Azure OpenAI image editing, mask inpainting, image edit API, input image, mask parameter, and computer vision image generation workflows.

QUESTION: 46

You have a Microsoft Foundry project that contains an agent.

The knowledge source for the agent is a set of scanned PDF troubleshooting guides stored in Azure Blob Storage. The guide pages contain two-column layouts and tables.

You use Azure Content Understanding in Foundry Tools to process the PDFs.

You plan to ingest the processed content into an index for Retrieval Augmented Generation (RAG) and store extracted fields for downstream automation.

Stakeholders must be able to verify where each extracted field value came from in the original PDF and route low-reliability extractions for manual review.

You need to ensure that the Content Understanding document analyzer output includes a per-field confidence score and source grounding locations within the source document.

What should you do?

- A. Enable estimateFieldSourceAndConfidence.**
- B. Configure the analyzer to use generative extraction for all fields.**
- C. Set enableSegment to true.**
- D. Provide labeled samples.**

Answer: A

The correct answer is A. Enable estimateFieldSourceAndConfidence . Azure Content Understanding document analyzers support an opt-in confidence and grounding capability for field extraction. Microsoft documentation states that to opt in for confidence and grounding, you set estimateFieldSourceAndConfidence = true in the analyzer configuration, or configure estimateSourceAndConfidence = true for specific fields.

This enables each extracted field to include a confidence score and references back to the

original document source location.

This directly satisfies both stakeholder requirements: source grounding allows users to verify where the extracted value came from in the scanned PDF, and the confidence score supports downstream automation rules, such as sending low-confidence fields to manual review. Microsoft's analyzer improvement guidance describes confidence scoring as a value between 0 and 1 and grounding as references or citations for extracted outputs to the original document content.

Generative extraction does not guarantee per-field confidence and source grounding. `enableSegment` is used for document segmentation, not confidence scoring. Labeled samples can improve extraction quality, but they do not by themselves enable confidence and grounding output. Reference topics: Content Understanding document analyzers, field extraction, confidence scoring, grounding, and RAG ingestion.

QUESTION: 47

You have a Microsoft Foundry project named Project1 that contains an agent. The agent uses an OpenAPI 3.0 specification to call an external weather service. The weather service requires a key to be passed in an HTTP header. The key value is stored as a connection in Project1.

You need to ensure that the key value from the connection is included automatically whenever the OpenAPI tool is invoked.

What should you configure in the OpenAPI specification?

- A.** an Azure Key Vault connection
- B.** a header parameter defined for each operation
- C.** an API key security scheme
- D.** a Bearer token security scheme

Answer: C

The correct configuration is an API key security scheme . For Microsoft Foundry Agent Service OpenAPI tools, the OpenAPI specification must declare authentication through the `components.securitySchemes` section and use a scheme of type `apiKey` when the external service expects a key in a header. Microsoft's OpenAPI tool guidance states that API key authentication requires updating the OpenAPI spec security schemes with one scheme of type `apiKey`, and the tool then uses the associated project connection to supply the key value at runtime. This allows the key stored in Project1's connection to be injected automatically when the tool is invoked.

A header parameter defined separately for each operation is not the correct approach because credentials should not be modeled as ordinary operation parameters. The Foundry guidance explicitly indicates that parameters requiring the API key should be removed from the OpenAPI spec because the API key is stored and passed through a connection. A Bearer token security scheme is used for bearer-token-style authorization, not a generic weather API key passed in a custom HTTP header. Azure Key Vault is a secret store, but the scenario already stores the key in a Foundry project connection. Reference topics: Microsoft Foundry Agent Service, OpenAPI tools, project connections, API key authentication, and OpenAPI security schemes.

QUESTION: 48

You have a customer support agent built by using the Microsoft Foundry Agent Service. The agent calls an Azure OpenAI model deployment.

During load testing, calls intermittently fail and return an HTTP 429 rate limit exceeded error.

You need to handle throttling to reduce call failures and improve reliability under load. The solution must remain within the service and model limits.

What should you do?

- A.** Implement a retry policy that uses exponential backoff and jitter.
- B.** Create a new thread and retry the calls immediately.
- C.** Reduce the number of registered tools.
- D.** Split uploaded content into smaller files.

Answer: A

The correct answer is A. Implement a retry policy that uses exponential backoff and jitter . HTTP 429 indicates that the request rate or token rate has exceeded the configured service limits for the model deployment. Microsoft Foundry Agent Service limits guidance specifically recommends implementing exponential backoff with jitter in application retry logic when agents receive rate-limit 429 errors. It also recommends reviewing Azure OpenAI quotas and token-per-minute and request-per-minute limits for the deployment. This approach improves reliability while remaining within service limits because retries are delayed progressively instead of immediately adding more pressure to an already throttled deployment. Microsoft Foundry Models quota guidance also states that unsuccessful requests still count toward per-minute rate limits and that continuously resending requests without backing off makes throttling worse. It recommends retry logic with exponential

backoff and using the Retry-After header when available.

Creating a new thread and retrying immediately does not change the deployment's rate limits and can worsen throttling. Reducing registered tools may simplify orchestration but does not directly address model RPM or TPM limits. Splitting uploaded content into smaller files may help ingestion scenarios, but it is not the correct throttling control for intermittent HTTP 429 model calls. Reference topics: Foundry Agent Service limits, Azure OpenAI quota management, throttling, retry policies, and production reliability.

QUESTION: 49

You have a Microsoft Foundry project.

You plan to build a customer support solution that contains an agent. The solution must meet the following requirements:

- * Provide accurate, context-aware responses grounded in internal product documentation stored in Azure AI Search.**
- * Require deep, multi-step reasoning across long contexts.**
- * Generate detailed natural language responses.**

Which type of model should you use to power the agent?

- A.** a multimodal model
- B.** a key phrase extraction model
- C.** a small language model (SLM)
- D.** a large language model (LLM)

Answer: D

The correct model type is a large language model (LLM) . The scenario requires an agent that can reason over retrieved documentation, synthesize context-aware answers, and generate detailed natural language responses. Microsoft Foundry RAG guidance defines Retrieval Augmented Generation as a pattern that combines search with large language models so responses are grounded in organizational data, which directly matches the use of Azure AI Search for internal product documentation.

An LLM is also the correct fit for deep, multi-step reasoning across long contexts. Azure AI Search agentic retrieval guidance states that complex chat and agent scenarios use an LLM to break a user query into smaller focused subqueries for better coverage over indexed content. Microsoft also describes Azure AI Search as a way to ground agents and chatbots in proprietary enterprise data for accurate, context-aware responses.

A multimodal model is unnecessary because the requirements are text-focused, not

image, audio, or video based. A key phrase extraction model performs narrow text analytics and cannot generate detailed answers.

An SLM may reduce cost and latency, but the stated need for deep reasoning and long-context synthesis favors an LLM. Reference topics: Microsoft Foundry RAG, Azure AI Search grounding, agentic retrieval, and model selection for generative AI agents.

QUESTION: 50

You have a Microsoft Foundry project that contains an agent named PaymentAgent.

PaymentAgent includes a function tool that issues customer refunds by using an external API.

You are creating a workflow in YAML.

You need to ensure that the workflow pauses for human approval and continues with the refund step only after approval is granted.

How should you complete the workflow definition? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

steps:

- id: propose_refund

type: agent

agent: PaymentAgent

- id: approval

type:

ask_question
basic_chat
data_transformation

- id: execute_refund

type: agent

agent: PaymentAgent

condition:

approval == "approved"
propose_refund.output != null
true

Answer:

```

steps:
  - id: propose_refund
    type: agent
    agent: PaymentAgent

  - id: approval
    type:
      ask_question
      basic_chat
      data_transformation

  - id: execute_refund
    type: agent
    agent: PaymentAgent
    condition:
      approval == "approved"
      propose_refund.output != null
      true

```

Explanation:

type: ask_question

condition: approval == " approved "

The approval step must use type: ask_question because the workflow must pause and wait for a human response before the refund execution proceeds. Microsoft Foundry workflows support human-in-the-loop patterns where the workflow asks the user a question and awaits input before continuing; this pattern is explicitly intended for approval requests and clarifying questions. The workflow guidance also identifies workflows as declarative sequences that orchestrate agents and business logic, including branching logic and human-in-the-loop steps.

The refund execution step must use condition: approval == " approved " so that the second invocation of PaymentAgent runs only when the approval response matches the required approval value. Using true would always execute the refund, bypassing the approval control. Using propose_refund.output != null would only confirm that the first agent step produced output; it would not prove that a human approved the refund.

data_transformation is also incorrect for the approval node because it sets or parses

values rather than pausing for user input. Reference topics: Microsoft Foundry workflows, human-in-the-loop workflow pattern, YAML workflow editing, agent orchestration, conditional execution, and workflow approval gates.

QUESTION: 51

You have a Microsoft Foundry project that contains a deployed ticket-triage agent. You discover that sometimes the agent responds without calling any tools, even when a tool is required.

You need to ensure that the agent calls a tool during execution.

How should you complete the Python code? To answer, drag the appropriate values to the correct targets.

Each value may be used

once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values

auto

required

response_format

tool_choice

tools

type

Answer Area

```
run_payload = {  
  "assistant_id": agent_id,  
  Value : Value ,  
  "metadata": {  
    "scenario": "ticket-triage"  
  }  
}
```

Answer:

Values

auto

required

response_format

tool_choice

tools

type

Answer Area

```
run_payload = {  
  "assistant_id": agent_id,  
  "tool_choice" : "required",  
  "metadata": {  
    "scenario": "ticket-triage"  
  }  
}
```

Explanation:

"tool_choice": "required"

The correct completion is "tool_choice": "required". In Microsoft Foundry Agent Service,

tool_choice controls whether the model can answer directly or must invoke a tool during a run. The official Foundry tool guidance states that tool_choice provides deterministic control over tool calling: auto allows the model to decide whether to call tools, none prevents tool use, and required forces the model to call one or more tools.

This directly addresses the issue where the ticket-triage agent sometimes responds without invoking a required tool.

The completed payload should therefore add the tool_choice property beside assistant_id, with the value " required " . The value " auto " is incorrect because it preserves the current nondeterministic behavior. The values " tools " and " type " do not force execution-time tool invocation in this payload; they are used for tool definitions or typed objects in other contexts. response_format controls output formatting, not tool execution.

Reference topics: Microsoft Foundry Agent Service, tool calling reliability, run payload configuration, tool_choice, agent execution, and deterministic tool invocation.

QUESTION: 52

You have a Microsoft Foundry project that contains an agent. The agent uses Azure Speech in Foundry Tools.

You fine-tune a baseline speech to text model for the en-us locale and publish the model.

The agent calls the Speech to text REST API and returns an error message indicating that the project ID is invalid.

You need to set the project property to the correct ID.

To what should you set the project property?

- A.** the custom speech endpoint URL
- B.** the project URL
- C.** the project ID
- D.** the custom speech project ID

Answer: D

The correct answer is D. the custom speech project ID . For custom speech fine-tuning, the Speech to text REST API uses a project property that must refer to the Custom Speech project, not the general Microsoft Foundry project. Microsoft's Custom Speech guidance states that when using the Speech to text REST API for custom speech, you must set the project property to the ID of your custom speech project. It also explicitly notes that the custom speech project ID is not the same as the Microsoft Foundry project ID.

This distinction explains the invalid project ID error. Supplying the Foundry project ID, project URL, or endpoint URL does not identify the Custom Speech project that owns the fine-tuned speech model. The custom speech endpoint URL is used when calling a deployed custom model endpoint for recognition, but it is not the value of the REST API project property. The project URL is also not accepted because the API expects the identifier value. Reference topics: Azure Speech in Foundry Tools, Custom Speech fine-tuning, Speech to text REST API, custom speech project ID, model publication, and endpoint configuration.

QUESTION: 53


You have a Microsoft Foundry project that contains a workflow for a customer support triage process.


You have an Ask a question node that stores user responses in a local variable named Var01.

You need to create the following Power Fx expressions:

- * An if/else condition expression that ensures that Var01 contains a value**
 - * A Send message expression that returns the stored user response in uppercase**
- How should you configure the expressions? To answer, select the appropriate options in the answer area.**


NOTE: Each correct selection is worth one point.

If/else condition expression: 

Send message expression: 


Answer:

If/else condition expression:



IsBlank(Local.Var01)
IsEmpty(Local.Var01)
Not(IsBlank(Local.Var01))

Send message expression:



{Local.Var01}
{Upper(Local.Var01)}
{Upper(Var01)}

Explanation:

If/else condition expression: Not(IsBlank(Local.Var01))

Send message expression: {Upper(Local.Var01)}

The correct if/else condition is Not(IsBlank(Local.Var01)) because the workflow must continue only when the Ask a question node has captured a nonblank value. In Power Fx, IsBlank checks whether a value is blank, while Not reverses that Boolean result; therefore, Not(IsBlank(Local.Var01)) evaluates to true only when Var01 contains user input.

IsEmpty(Local.Var01) is not the correct choice because IsEmpty checks whether a table contains records, not whether a text variable has a value. The Power Fx formula reference distinguishes these functions by defining IsBlank for blank values and IsEmpty for empty tables.

The Send message expression must be {Upper(Local.Var01)}. Microsoft Foundry workflow guidance uses the same pattern for an Ask a question node saved as Var01, followed by a Send message action that outputs

{Upper(Local.Var01)}. The Local. prefix is required because the response is stored in a local workflow variable, and Upper() converts the stored text response to uppercase. Reference topics: Microsoft Foundry workflows, Power Fx expressions, local variables, if/else branching, and Send message actions.

QUESTION: 54

You need to ensure that Agent1Dev Team can access Agent1. The solution must meet the security and compliance requirements.

How should you complete the Python code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
from azure.identity import DefaultAzureCredential
from azure.ai.projects import AIProjectClient
from azure.core.credentials import AzureKeyCredential
myEndpoint = "https://contoso.services.ai.azure.com/api/projects/project1"
project_client = AIProjectClient(
    endpoint=myEndpoint,
    credential=
)
myAgent = "Agent1"
agent = project_client.agents. (agent_name=myAgent)

print(f "Retrieved agent: {agent.name}")
```

Answer:

Answer Area

```
from azure.identity import DefaultAzureCredential
from azure.ai.projects import AIProjectClient
from azure.core.credentials import AzureKeyCredential
myEndpoint = "https://contoso.services.ai.azure.com/api/projects/project1"
project_client = AIProjectClient(
    endpoint=myEndpoint,
    credential=
)
myAgent = "Agent1"
agent = project_client.agents. (agent_name=myAgent)

print(f "Retrieved agent: {agent.name}")
```

Explanation:

credential = DefaultAzureCredential()

agent = project_client.agents.get(agent_name=myAgent)

The correct authentication option is DefaultAzureCredential() because the case study states that API keys must not be used to access Foundry-deployed models and that

Contoso developers must authenticate to Microsoft Foundry resources by using Microsoft Entra authentication. It also states that access to Project1 must be assigned to Agent1Dev Team by using the security group SC_Agent1_Dev . Microsoft Foundry authentication guidance recommends Microsoft Entra ID for production workloads because it supports least- privilege RBAC, per-principal auditing, and keyless authentication.

AzureKeyCredential() would violate the no-API-key requirement, and None would not provide a valid credential.

The correct agent operation is get because the task is to access an existing agent named Agent1, not create a new version or retrieve a specific published version. Microsoft Foundry SDK examples show AIProjectClient created with DefaultAzureCredential() and then using project agent operations to create, retrieve, or interact with agents by name. To meet the compliance requirement, the group SC_Agent1_Dev must also be granted the appropriate project-scoped Foundry role, such as Foundry User, for Project1. Reference topics: Microsoft Entra authentication, Foundry RBAC, AIProjectClient, and project agent access.

QUESTION: 55

You have a Microsoft Foundry project that generates product marketing images from text prompts.

After publishing several images, the legal team at your company identifies a competitor ' s logo on a sign in the background of an image.

You need to remove only the logo, while preserving the rest of the image.

What should you do?

- A.** Increase the prompt guidance strength.
- B.** Modify the original prompt to exclude brand names.
- C.** Apply a mask-based inpainting edit to the part of the image that contains the logo.
- D.** Rerun the prompt by using a different random seed.

Answer: C

The correct answer is C because the requirement is a localized image edit: remove only the competitor logo while preserving the rest of the already generated image. Azure OpenAI image editing is designed for modifying existing images based on a text instruction, rather than regenerating the entire image from scratch.

Microsoft's Azure OpenAI image guidance states that the Image Edit API modifies existing images and requires an input image as part of the request. In a mask-based inpainting workflow, the mask identifies the exact region to change, allowing the model to replace

only the logo area while retaining surrounding background, composition, lighting, and product content.

Increasing prompt guidance strength would affect adherence during generation, but it would not safely remove a specific logo from a completed image. Modifying the original prompt and regenerating may create a different image and does not guarantee preservation of the approved visual content. Rerunning with a different random seed also changes the image unpredictably and may introduce new brand or legal issues.

Mask-based inpainting is the minimal-change remediation method for post-generation brand cleanup.

Reference topics: Azure OpenAI image editing, inpainting, mask-guided edits, image generation governance, and computer vision solutions.

QUESTION: 56

You are creating an agent workflow in a Microsoft Foundry project to support natural voice interactions.

The agent must receive continuous audio input, convert the input into text for reasoning, and then return spoken responses to a user. The workflow must meet the following requirements:

- . Support turn-taking dynamics, where the agent begins to generate the speech output before the user finishes speaking.**
- . Operate with low latency to maintain a conversational experience.**

You need to enable both speech to text and text to speech in a real-time agent interaction.

What should you do?

- A.** Use an embeddings model to encode the audio, and then decode the audio into text and speech.
- B.** Use batch transcription to convert the audio input and return text responses from the agent.
- C.** Use speech translation to convert the audio into another language and return the translated text.
- D.** Use real-time speech to text for incoming audio and text to speech for agent responses.

Answer: D

The correct answer is D. Use real-time speech to text for incoming audio and text to speech for agent responses . The workflow requires continuous audio input, low-latency

transcription for reasoning, and spoken output back to the user. Azure Speech in Foundry Tools real-time speech to text is designed for immediate transcription from streaming audio, which satisfies the incoming-audio side of the interaction. Text to speech provides the outbound spoken response path after the agent generates its answer.

This pattern aligns with Microsoft's real-time voice-agent architecture. The Voice Live API overview explains that low-latency speech-to-speech systems integrate speech recognition, generative reasoning, and text-to-speech functionality to create natural voice experiences. It also identifies contact centers as a key scenario and highlights low perceived latency for end users. Embeddings do not decode audio into conversational speech. Batch transcription introduces file-oriented delay and is not suitable for turn-taking.

Speech translation is only appropriate when translating between languages and does not provide the required reasoning-plus-spoken-response loop. Reference topics: Azure Speech in Foundry Tools, real-time speech to text, text to speech, voice agents, low-latency interaction, and conversational turn-taking.

QUESTION: 57

You have a Python application named App1 that integrates with a Microsoft Foundry project named Project1.

You need to ensure that App1 meets the following requirements:

- * Authenticates by using a Microsoft Entra managed identity**
 - * Sends prompts to a deployed model by using the Azure OpenAI Responses API**
- How should you complete the Python code? To answer, select the appropriate options in the answer area.**

NOTE: Each correct selection is worth one point.

```
from azure.identity import DefaultAzureCredential

from azure.ai.projects import AIProjectClient

credential =  ()
    AzureKeyCredential
    ClientSecretCredential
    DefaultAzureCredential

project_client = AIProjectClient(
    endpoint="https://contosoai.services.ai.azure.com/api/projects/project1",
    credential=credential,
)

with project_client.get_openai_client() as openai_client:
    response = openai_client.responses.  (
        compact
        create
        retrieve

    model="trail-guide-chat",
    input="Create a 3-day hiking itinerary near Seattle.",
)

print(response.output_text)
```

Answer:

```

from azure.identity import DefaultAzureCredential
from azure.ai.projects import AIProjectClient

credential =  ()

project_client = AIProjectClient(
    endpoint="https://contosoai.services.ai.azure.com/api/projects/project1",
    credential=credential,
)

with project_client.get_openai_client() as openai_client:
    response = openai_client.responses. (

        model="trail-guide-chat",
        input="Create a 3-day hiking itinerary near Seattle.",
    )

    print(response.output_text)

```

Explanation:

`credential = DefaultAzureCredential`

`response = openai_client.responses.create`

The correct authentication class is `DefaultAzureCredential` because the requirement is Microsoft Entra-based authentication, and this credential supports managed identity in hosted Azure environments. The Microsoft Foundry quickstart for Python shows the project client being created with `AIProjectClient (endpoint=PROJECT_ENDPOINT, credential=DefaultAzureCredential())`, which enables the Foundry SDK to authenticate without using API keys or embedded secrets. The same guidance shows creating an OpenAI-compatible client from the project by calling `project.get_openai_client()`.

The correct Responses API method is `create` because the application must send a new prompt to the deployed model and receive generated output. Microsoft's Foundry quickstart demonstrates the exact pattern: `response = openai.responses.create(model= " gpt-5-mini " , input= " ... " )`, followed by reading `response.output_text`.

The `retrieve` operation is used to fetch an existing response, not submit a new inference request, and `compact` is not the correct method for generating a model response.

`AzureKeyCredential` would violate the Microsoft Entra managed identity requirement, while

ClientSecretCredential uses an application secret rather than managed identity. Reference topics: Microsoft Foundry SDK, AIProjectClient, Microsoft Entra authentication, DefaultAzureCredential, and Azure OpenAI Responses API.

QUESTION: 58

You have a Microsoft Foundry project that contains a model deployment.

You have an application that calls the deployment by using the Azure OpenAI v1 API and DefaultAzureCredential.

The developers at your company receive HTTP 403 errors when they send inference requests, even after running az login.

You need to ensure that the developers can perform model inference. The solution must follow the principle of least privilege.

Which role-based access control (RBAC) role should you assign to the developers?

- A.** Cognitive Services OpenAI User
- B.** Cognitive Services Data Reader
- C.** Cognitive Services User
- D.** Contributor

Answer: A

The correct role is Cognitive Services OpenAI User . The application is using DefaultAzureCredential, so az login only proves the developer's Microsoft Entra identity and enables token acquisition. It does not by itself grant authorization to the model deployment. Azure OpenAI and Microsoft Foundry separate authentication from authorization; Microsoft Entra ID provides token-based authentication, while Azure RBAC controls whether the signed-in principal can perform data-plane actions such as model inference. Microsoft's Foundry guidance states that Microsoft Entra ID supports granular RBAC and that data-plane operations include runtime usage such as chat completions and embedding generation.

For Azure OpenAI resources, the Cognitive Services OpenAI User role specifically allows users to make inference API calls with Microsoft Entra ID against deployed models, while preventing higher-privilege actions such as creating deployments, copying keys, fine-tuning, or managing the resource.

The other roles are not least privilege. Contributor grants broad management-plane permissions. Cognitive Services User is broader and less specific than the Azure OpenAI inference role. Cognitive Services Data Reader is read-oriented and does not provide the

required model inference data action. Reference topics:

Microsoft Foundry authentication and authorization, Azure RBAC, Microsoft Entra ID keyless authentication, Azure OpenAI v1 API, and least-privilege model inference.

QUESTION: 59

You have a Microsoft Foundry project that contains a customer support agent built by using the Foundry Agent Service.

The agent uploads user-provided screenshots to Azure Storage through a ticketing tool and receives a blob URL for additional reasoning.

You need to use image moderation during agent runs and prevent harmful content from being returned during runs. Azure AI Content Safety must access the images by using the blob URL. The solution must follow the principle of least privilege.

What should you configure for Content Safety? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Guardrails:

- Select Tool call and set Action to Block.
- Select User input and Output and set Action to Annotate.
- Select User input and Tool response and set Action to Annotate.
- Select User input, Output, Tool response, and Tool call and set Action to Block.

Storage access:

- Storage account access keys.
- A user-assigned identity that is assigned the Storage Queue Data Contributor role
- A system-assigned managed identity that is assigned the Storage Blob Data Reader role
- A system-assigned managed identity that is assigned the Storage Blob Data Contributor role

Answer:

Guardrails:
 Select User input and Output and set Action to Annotate.
 Select User input and Tool response and set Action to Annotate.
 Select User input, Output, Tool response, and Tool call and set Action to Block.

Storage access:
 A user-assigned identity that is assigned the Storage Queue Data Contributor role
 A system-assigned managed identity that is assigned the Storage Blob Data Reader role
 A system-assigned managed identity that is assigned the Storage Blob Data Contributor role

Explanation:

Guardrails: Select User input, Output, Tool response, and Tool call and set Action to Block.

Storage access: A system-assigned managed identity that is assigned the Storage Blob Data Contributor role. The guardrail must be applied to User input, Output, Tool response, and Tool call with the action set to Block. Microsoft Foundry guardrails support four intervention points: user input, tool call, tool response, and output. This scenario includes user-provided screenshots, a ticketing tool that uploads images and returns blob URLs, and final agent responses. Applying blocking controls at all four points ensures harmful image-related content is inspected throughout the agent run and prevented from continuing or being returned to the user.

Microsoft's guardrails guidance also states that tool call and tool response controls are specifically required when harmful content can pass through agent tools.

For storage, configure the Azure AI Content Safety resource with a system-assigned managed identity and grant it Storage Blob Data Contributor on the storage account or container. The Content Safety image moderation quickstart states that images can be supplied by blob storage URL and that the Content Safety resource must be given storage access by enabling its system-assigned managed identity and assigning Storage Blob Data Contributor or Owner; Contributor is the least-privileged valid option shown. Reference topics: Foundry guardrails, agent intervention points, image moderation, managed identity, and Azure Storage RBAC.

QUESTION: 60

You have a Microsoft Foundry project that generates product marketing images from text prompts.

After publishing several images, the legal team at your company identifies a competitor 's logo on a sign in the background of an image.

You need to remove only the logo, while preserving the rest of the image.

What should you do?

- A. Modify the original prompt to exclude brand names.
- B. Increase the prompt guidance strength.
- C. Rerun the prompt by using a different random seed.
- D. Apply a mask-based inpainting edit to the part of the image that contains the logo.

Answer: D

QUESTION: 61

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses.

You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model.

You need to implement controls to mitigate the risk.

Solution: You configure image moderation to block unsafe content before processing the images.

Does this meet the goal?

A. No

B. Yes

Answer: A

QUESTION: 62

You have a Microsoft Foundry project that contains an agent. The agent uses Azure Speech in Foundry Tools.

You fine-tune a baseline speech to text model for the en-us locale and publish the model.

The agent calls the Speech to text REST API and returns an error message indicating that the project ID is invalid.

You need to set the project property to the correct ID.

To what should you set the project property?

- A.the project ID
- B.the project URL
- C.the custom speech endpoint URL
- D.the custom speech project ID

Answer: D

QUESTION: 63

This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem. After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen. You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses. You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model. You need to implement controls to mitigate the risk. Solution: You configure protected material detection. Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation: The solution does not meet the goal. Protected material detection is intended to identify large language model output that matches known protected text or code, such as copyrighted text, selected web content, song lyrics, articles, recipes, or code. Microsoft describes protected material detection as a control for preventing AI- generated content from reproducing known protected material, not as a control for image safety or prompt injection.

The stated risk has two parts: users can upload unsafe images, and users can embed hidden instructions in images to manipulate the model. Unsafe image uploads require image moderation, because Azure AI Content Safety provides image APIs that detect harmful content across modalities and can support blocking decisions by harm category and severity. Hidden instructions extracted from images are indirect prompt injection or document attacks; Microsoft Prompt Shields are the capability designed to detect user prompt attacks and document attacks, including harmful instructions embedded in third-

party content. Therefore, protected material detection alone does not mitigate either primary risk. Reference topics: Azure AI Content Safety, image moderation, Prompt Shields, document attacks, indirect prompt injection, and protected material detection. Topic 1, Case Study Contoso, Ltd Overview Contoso, Ltd is a multinational retail company that builds, deploys, and manages generative AI and agent- based solutions by using Microsoft Foundry. Identity Environment: Contoso uses Microsoft Entra ID for identity management, authentication, and authorization capabilities that enable agents to access organizational resources and services. Contoso recently formed a new AI engineering team named Agent1Dev Team to optimize and maintain existing AI solutions. The team collaborates with solution architects, DevOps engineers, and security engineers to design, implement, monitor, and secure AI applications. Contoso also has a team named Agent1Test Team that is responsible for validating AI solutions before the solution deployments. Generative Environment:

Contoso has a Microsoft Foundry deployment that contains two projects named Project1 and Project2. Project1 Project1 contains a customer support agent named Agent1 that assists customers with product inquiries and troubleshooting requests. Agent1 has the following configurations: * Agent1 uses a base model deployment. * A safety evaluation pipeline is NOT enabled. * Tool invocation approval workflows are NOT enabled. * Conversation memory constraints are NOT configured. Agent1 interacts with customers by using digital support channels and answers general questions about Contoso products. Project1 is deployed to an Azure region located in the European Union (EU). Agent1Dev Team will use Project1 to optimize and maintain Agent1. Project2 Project2 contains a deployed video generation model. The marketing department at Contoso has access to Project2 and plans to use the model to develop a video creation solution.

Development of the solution is incomplete. Data Environment: Contoso stores product-related information in Azure resources that support AI applications. The Azure environment contains an Azure Blob Storage account named storage1 that stores product detail sheets for all the Contoso products. The product sheets include specifications, feature descriptions, and product support information that Agent1 can use to answer customer questions. The product sheets are stored in the PDF format. Problem Statement: Contoso identifies the following issues: * Agent1 has only general knowledge of the Contoso products. * A recent chat interaction with Agent1 was analyzed for sentiment. The results of the analysis have NOT been processed yet. * Agent1 does NOT use the detailed product information in the product sheets stored in storage1 when responding to customer questions. * The finance department at Contoso reports that vendor invoices must be reviewed manually to ensure that the invoices match the terms defined in the vendor contracts. The invoices contain tables, logos, and varied layouts that make the documents

difficult to process consistently. Requirement: Planned Changes: Contoso plans to implement the following changes: * Implement a solution for Project1 that analyzes the vendor invoices by evaluating both the visual layout and the textual content of the invoices, so that the invoice details can be verified against the vendor contract terms. * Update the base model deployment used by Agent1 and standardize the model version to ensure continuity and consistent responses. * Enable Agent1 to retrieve and use the detailed product information from the product sheets stored in storage1. * Implement an indexing solution for the product sheets that Agent1 can use to answer customer questions.

* Complete the development of the video creation solution. Technical Requirements: Contoso identifies the following technical requirements: * The model deployment used by Agent1 must support scalable, high-throughput generative AI workloads and dynamically scale to handle variable customer support traffic, without requiring reserved throughput capacity. * The product sheets must be processed by using an indexing pipeline that enables semantic and vector search, so that Agent1 can retrieve the relevant product information. * Responses generated by using the product sheet information must be relevant, complete, and accurate. * Agent1 must be able to use the product sheets to answer natural language questions about product details. * The model version used by Agent1 must remain consistent to ensure stable responses. * The data processed by the model must remain within the EU. Safety and Compliance Requirements: Contoso identifies the following security and compliance requirements: * API keys must NOT be used to access Foundry-deployed models. * Access to the Azure resources must follow the principle of least privilege. * The developers at Contoso must authenticate to Microsoft Foundry resources by using Microsoft Entra authentication. * Access to Project1 must be assigned to the members of Agent1Dev Team by using a security group named SC_Agent1_Dev. * Access to Project1 must be assigned to the members of Agent1Test Team by using a security group named SC_Agent1_Test. * Agent1 must never reveal customer information, even if a document that contains customer data is added erroneously to the product sheet repository in storage1. * The product sheets might contain images that include embedded text. Agent1 must be protected from malicious instructions potentially hidden within the images.

Business Information: Contoso identifies the following business requirements: * Users that interact with Agent1 must have a personalized experience in future interactions, including the ability for Agent1 to retain conversation context and recall relevant information from previous interactions. * Agent1 must answer questions only about the products sold by Contoso.

QUESTION: 64

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You increase the value of the temperature parameter.

Does this meet the goal?

A.No

B.Yes

Answer: A

QUESTION: 65

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses.

You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model.

You need to implement controls to mitigate the risk.

Solution: You configure image moderation to block unsafe content before processing the images.

Does this meet the goal?

A.No

B.Yes

Answer: A

QUESTION: 66

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses.

You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model.

You need to implement controls to mitigate the risk.

Solution: You configure a prompt shield for user prompts.

Does this meet the goal?

A.No

B.Yes

Answer: B

QUESTION: 67

HOTSPOT -

You have a Microsoft Foundry project that contains an agent.

You use a GitHub Actions workflow for CI/CD.

You need to configure the workflow to automatically evaluate the agent when a pull request (PR) is created and prevent branches from merging if the evaluation results do NOT meet the defined thresholds.

How should you configure the workflow? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Authentication method:

- A personal access token (PAT)
- A user-assigned managed identity
- An Azure Login action that uses OpenID Connect (OIDC)

If the evaluation results are **NOT** met, configure the workflow to:

- Lock the target branch
- Send an alert
- Fail

Answer:

Answer Area

Authentication method:

- A personal access token (PAT)
- A user-assigned managed identity
- An Azure Login action that uses OpenID Connect (OIDC)

If the evaluation results are **NOT** met, configure the workflow to:

- Lock the target branch
- Send an alert
- Fail

QUESTION: 68

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will **NOT** be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You add a reflection pass that regenerates the response if the required clauses are missing.

Does this meet the goal?

A.No

B.Yes

Answer: B

QUESTION: 69

HOTSPOT -

You have a Microsoft Foundry project that contains an agent.

The agent accepts user-uploaded screenshots and uses a multimodal chat model.

Some screenshots contain potentially malicious embedded text.


You need to prevent a prompt injection attack and ensure that third-party content is treated as lower trust.

How should you configure prompt shields for document attacks? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Prompt shields action:



A dropdown menu with a downward arrow icon. The menu is open, showing three options: "Disable the shield.", "Set action to block.", and "Set action to annotate."

Additional mitigation:



A dropdown menu with a downward arrow icon. The menu is open, showing three options: "Enable Spotlighting.", "Create a custom blocklist.", and "Use optical character recognition (OCR) to extract the text from the images first."

Answer:

Answer Area

Input Shield mode:

▼

Enabled for upload

Disabled for upload

Enabled for content

Document attack protection:

▼

Enabled, but override on

On, and override on individual documents

On, include jailbreak and prompt injection protection from Azure AI Prompt Shields

QUESTION: 70

You are building a text-to-speech solution that uses Azure Speech in Foundry Tools to read instructions from the script in a text file.

You discover that the solution often pronounces technical terms incorrectly.

You need to prevent the incorrect pronunciations. The solution must minimize development effort.

What should you do?

- A.** From Speech Studio, train a custom neural voice
- B.** Use Speech Synthesis Markup Language (SSML) to specify phonemes.
- C.** Use Speech Synthesis Markup Language (SSML) to apply say as rules.
- D.** Use Speech Synthesis Markup Language (SSML) to adjust the prosody of the voice.
- E.** From Azure OpenAI use the Whisper model.

Answer: B

Using Speech Synthesis Markup Language (SSML) with the <phoneme> element is the ideal way to fix mispronunciations for technical terms.

The <phoneme> tag lets you override the default text-to-speech model by explicitly defining the sounds using the International Phonetic Alphabet (IPA).

Reference:

<https://learn.microsoft.com/en-us/answers/questions/5729867/pronunciation-issue-when-generating-audio-from-ssm>

QUESTION: 71

You have a Microsoft Foundry project that contains a Retrieval Augmented Generation (RAG) chat solution used by customer support agents.

You are adding an automated pre-production evaluation step to a CI/CD pipeline named Pipeline1. The evaluation will run against a labeled test dataset that contains support questions and the expected grounding context.

You need to ensure that Pipeline1 fails if unsupported content or a retrieval mismatch exceeds a defined threshold:

- responses include claims not supported by the retrieved source content**

- retrieved source content does not align with the labeled expected context**

Which two built-in evaluators should you use in Pipeline1? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

A. Retrieval

B. Fluency

C. Coherence

D. Groundedness

E. Response Completeness

Answer: A,D

The Groundedness Evaluator evaluator validates that the model's responses include only claims supported by the retrieved source content. It flags ungrounded content or hallucinations. If the average score drops below your defined threshold, it triggers a pipeline failure.

The correct additional built-in evaluator appropriate for the pipeline is Retrieval (specifically, the RetrievalEvaluator or DocumentRetrievalEvaluator).

A standard RAG evaluation pipeline assesses both the generator (the LLM producing the answer) and the retriever (the search system pulling documentation). The CI/CD requirements specify two distinct failure thresholds:

Responses including claims not supported by the retrieved source content: This checks for model hallucinations and is handled by the Groundedness Evaluator.

Retrieved source content not aligning with the labeled expected context: This explicitly measures the performance of your search step against your ground-truth data. The built-in Retrieval evaluator maps to this requirement. It computes metrics like context recall to ensure your system successfully retrieves the exact reference documents specified in your

labeled test dataset Reference:

<https://learn.microsoft.com/en-us/microsoft-365/copilot/extensibility/evaluations-cli-evaluators>

QUESTION: 72

A multimodal agent processes user-uploaded images. An attacker embeds hidden text instructions inside an image, hoping the agent will read and act on them. Which Azure AI Content Safety capability is designed to detect this kind of attack?

- A.** A content filter set to High severity for the Violence category
- B.** Prompt Shields, which detects document and indirect prompt injection
- C.** Raising the image resolution threshold on the analyser
- D.** Disabling the agent's file search tool

Answer: B

Prompt Shields detects prompt injection attacks, including indirect injection where malicious instructions are embedded in content the model processes, such as text hidden inside an image.

QUESTION: 73

You are building a customer support web app named App1 in Microsoft Foundry that uses a GPT realtime model.

App1 must support:

- Live, low-latency voice conversations that use Azure OpenAI**
- Streaming audio input from users and playback audio responses**

You need to configure a connection method that supports real-time audio streaming in client application and targets approximately 100 ms latency.

Which connection method should you use?

- A.** RTMP
- B.** WebRTC
- C.** SIP
- D.** WebSocket

Answer: D

WebSockets is the best connection method for this application because it enables full-duplex, bi-directional streaming over a single TCP connection, meeting the strict ~100 ms latency requirement for real-time audio.

Reference:

<https://www.chat-data.com/blog/implement-openai-realtime-api-for-chatgpt-voice>

QUESTION: 74

You have a Microsoft Foundry project that contains a support-ticket triage agent built by using the Foundry Agent Service.

The agent uses tool to classify the ticket type and set the ticket priority.

Sometimes, the same support case continues across multiple sessions over several days.

You need to persist state by using a durable ID to ensure that the agent can automatically reuse the full interaction history. The solution must preserve previous user messages, tool calls and tool outputs across turns and sessions.

Which runtime component should you use?

- A. output item**
- B. agent**
- C. conversation**
- D. response**

Answer: C

To achieve state persistence and ensure that the agent automatically reuses the full multi-session interaction history (including previous user messages, tool calls, and tool outputs), you must include a conversation component.

References:

<https://learn.microsoft.com/en-us/azure/foundry/agents/concepts/runtime-components>

QUESTION: 75

Drag and Drop Question

You have a web app that uses Azure AI Search.

When reviewing activity you see greater than expected search query volumes. You suspect that the query key is compromised.

You need to prevent unauthorized access to the search endpoint and ensure that users only have read only access to the documents collection. The solution must minimize app downtime.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Regenerate the primary admin key	
Regenerate the secondary admin key	1
Change the app to use the secondary admin key	2
Add a new query key	3
Change the app to use the new key	
Delete the compromised key	

Answer:

Actions**Answer Area**

Regenerate the primary admin key	
Regenerate the secondary admin key	1 Delete the compromised key
Change the app to use the secondary admin key	2 Add a new query key
Add a new query key	3 Change the app to use the new key
Change the app to use the new key	
Delete the compromised key	

Explanation:

Enforces Read-Only Permissions: Query keys are specifically designed to provide read-only access to the documents collection of an index. Admin keys provide full read-write administrative privileges and should never be distributed to consumer-facing applications. Zero Downtime: Azure AI Search lets you generate up to 50 individual query keys. Creating a new one allows the app to stay online throughout the entire key rotation process

Reference:

<https://learn.microsoft.com/en-us/azure/search/search-security-api-keys>

QUESTION: 76

You have a large collection of image files and PDF documents stored in an Azure Storage account. The documents contain tabular data.

You need to extract the tables into a structured format that can be imported into a database. The solution must minimize development effort. What should you use?

- A. Azure Vision in Foundry Tools
- B. Azure AI Video Indexer
- C. Azure Content Understanding in Foundry Tools
- D. Azure Document Intelligence in Foundry Tools

Answer: D

To extract tables from numerous images and PDF documents in Azure Storage into a structured database with low development, you should utilize Azure AI Document Intelligence (formerly Form Recognizer) within the Azure AI Foundry ecosystem. The Layout Model is best for extracting tables, rows, columns, and headers.

Reference:

<https://azure.microsoft.com/en-us/products/ai-foundry/tools/document-intelligence>

QUESTION: 77

You have a Microsoft Foundry project that contains a customer support agent built on a deployed chat model.

The agent responses are validated by using an automated testing system that compares generated answers to stored expected outputs. Identical prompts must return consistent response to prevent automated test failures.

You need to reduce response variability, without modifying the prompt or reducing factual accuracy.

What should you do for the model?

- A.** Increase the max_tokens parameter.
- B.** Remove stop sequences from the requests.
- C.** Decrease the temperature parameter.
- D.** Increase the temperature parameter.

Answer: C

To reduce response variability and ensure identical prompts return consistent answers, you should decrease the temperature parameter.

Temperature controls the randomness of the model's output. Setting the temperature closer to 0 makes the model deterministic. It forces the model to choose the highest-probability words every time, ensuring that identical prompts consistently yield identical or near-identical responses to pass your automated testing.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry/agents/concepts/runtime-components>

QUESTION: 78

You have a custom named entity recognition (NER) project in Azure Language in Foundry Tools for support tickets. The schema for the project contains an entity type

named ContactInfo.

In tagged training files, ContactInfo is used for phone numbers, email addresses, and social media handles.

Model evaluation shows low precision for ContactInfo, including false positives in which nearby text is extracted as ContactInfo.

You need to improve the precision of the project.

What should you do before retraining the model?

- A.** Lower the confidence threshold for ContactInfo.
- B.** Trigger an auto-labeling job.
- C.** Add more support tickets as training data and label more ContactInfo entities.
- D.** Replace ContactInfo by using Phone, Email, and SocialMedia entities. Relabel every matching span.

Answer: D

To improve the precision of the project before retraining the model, you should replace ContactInfo by using Phone, Email, and SocialMedia entities, and then relabel every matching span.

Reduces Model Confusion: Combining distinct types of data (like phone numbers, emails, and usernames) into a single generic entity (ContactInfo) creates conflicting patterns for the model.

Splitting them allows the model to learn the specific, unique syntax rules of each type.

Eliminates Over-extraction: When an entity is too broad, the model struggles to identify where the entity begins and ends, leading to false positives by grabbing adjacent text.

Distinct entities provide sharper, cleaner boundaries.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/custom-named-entity-recognition/overview>

QUESTION: 79

You have an application that uses a prebuilt Azure Language model.

You need to generate a summary that meets the following requirements:

- Identify each presenter in the video and attribute each text file to sentences**
- Preserves the original sentence order**

- Returns exactly three sentences

Which Language service feature should you use?

- A.** key phrase extraction
- B.** sentiment analysis
- C.** extractive summarization
- D.** abstractive summarization

Answer: C

The Azure AI Language service's extractive summarization feature is the appropriate choice to summarize video transcripts while preserving original sentence structure and order.

Using the REST API or client library, you can set sentenceCount to exactly three to generate a concise summary from key phrases.

This approach avoids generating new, abstract text, fulfilling the constraint to keep original text, unlike abstractive summarization.

Key Implementation Details:

Service: Azure AI Language, specifically the Summarization feature.

Method: Extractive summarization selects the most relevant sentences directly from the source text, which is ideal for maintaining the original context, speaker attribution, and sentence structure.

Configuration: The API call allows setting the sentenceCount parameter to three, ensuring the output is limited to that exact number.

Workflow: This model works best on conversational transcripts processed via speech-to-text (e.g., Azure AI Speech), where speakers are identified, allowing the summarization tool to aggregate sentences while preserving the speaker-to-text relationship.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/summarization/how-to/text-summarization>

QUESTION: 80

Hotspot Question

You have a Python application that redacts sensitive information before sending prompt text to a language model. The application has the following code:

```

sample_text = "Contact John Doe at 312-555-1234 or
john.doe@contoso.com. His SSN is 859-98-0987."

def redact_for_model(input_text):
    payload = {
        "kind": "PiiEntityRecognition",
        "parameters": {
            "modelVersion": "latest",
            "piiCategories": ["Person", "PhoneNumber"],
            "redactionPolicies": [
                {
                    "policyKind": "entityMask"
                }
            ]
        },
        "analysisInput": {
            "documents": [
                {
                    "id": "1",
                    "language": "en",
                    "text": input_text
                }
            ]
        }
    }
    ...
    return {
        "text_for_model": doc.get("redactedText", ""),
        "audit": [
            (entity["text"], entity["category"], entity
["confidenceScore"])
            for entity in doc["entities"]
        ]
    }
result = redact_for_model(sample_text)

```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Statements	Yes	No
For sample_text, audit will include entity records for Contact and SSN.	☐	☐
For sample_text, text_for_model will include john.doe@contoso.com and 859-98-0987.	☐	☐
For sample_text, text_for_model will contain entity type masks for John Doe and 312-555-1234.	☐	☐

Answer:

Statements	Yes	No
For sample_text, audit will include entity records for Contact and SSN.	☐	☒
For sample_text, text_for_model will include john.doe@contoso.com and 859-98-0987.	☒	☐
For sample_text, text_for_model will contain entity type masks for John Doe and 312-555-1234.	☒	☐

Explanation:

Box 1: No

The audit list will not include entity records for "Contact" and "SSN" because neither word is a target PII entity type defined in the configuration.

The code limits scanning to Person and PhoneNumber via the piiCategories list.

"Contact" Definition: This is a regular English verb/noun, not a person's name or a telephone number.

"SSN" Definition: While "SSN" is a text label that points to sensitive data, it is a descriptor rather than an actual Social Security Number.

Box 2: Yes

For the given sample_text, text_for_model will include both john.doe@contoso.com and 859-

98-0987 completely unredacted.

The payload explicitly restricts the Personally Identifiable Information (PII) recognition to only two specific categories by setting piiCategories: ["Person", "PhoneNumber"].

Email Address Excluded: Because the Email category is omitted from the configuration list, the service bypasses the email address (john.doe@contoso.com) and leaves it fully intact.

SSN Excluded: Similarly, because the USSSN (or equivalent Social Security Number) category is not specified in the piiCategories array, the SSN (859-98-0987) is completely ignored by the redaction policy and remains visible.

Box 3: Yes

text_for_model will contain entity type masks for "John Doe" and "312-555-1234", but only for those two specific items.

The piiCategories array explicitly requests Person and PhoneNumber.

Applied Redaction: The Azure AI Language service will replace "John Doe" and "312-555-1234" with masks like * or entity labels (depending on exact policy settings).

Skipped Data: The email (john.doe@contoso.com) and the SSN (859-98-0987) will not be redacted because their categories (Email and USSSN) were omitted from your piiCategories list.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/personally-identifiable-information/conversations-entity-categories>

QUESTION: 81

You are building a chatbot.

You need to ensure that the chatbot can classify user input into separate categories. The categories must be dynamic and defined at the time of inference. Which service should you use to classify the input?

- A.** Azure AI Language custom named entity recognition (NER)
- B.** Azure AI Language custom text classification
- C.** Azure OpenAI text classification
- D.** Azure OpenAI text summarization

Answer: C

To classify user input into separate categories that are dynamic and defined at the time of inference, Azure OpenAI text classification is the best option. OpenAI models such as GPT can dynamically classify text into categories based on context and user-defined instructions. Since the categories need to be flexible and determined at inference time, OpenAI's natural language understanding capabilities are well-suited for this task.

QUESTION: 82

A finance team must extract named fields, such as invoice number and total, from invoices, and they need a per-field confidence score so that low-confidence values can be routed to a human reviewer. Which Azure Content Understanding mode should they use?

- A.** Pro mode
- B.** Standard mode
- C.** Either mode, because both return confidence scores
- D.** Neither, because confidence scoring requires Azure AI Document Intelligence

Answer: B

Standard mode in Azure Content Understanding supports extract fields and returns confidence scores and grounding, which is exactly what a human-in-the-loop review workflow needs. Pro mode is built for multi-step reasoning across documents and deliberately omits confidence scoring.

Content Understanding pro mode currently doesn't offer confidence scores or grounding. It supports classify and generate fields, but it doesn't support extract fields.

QUESTION: 83

Case Study 1 - Contoso, Ltd

Overview

Company Information

Contoso, Ltd is a multinational retail company that builds, deploys, and manages generative AI and agent-based solutions by using Microsoft Foundry.

Existing Environment

Identity Environment

Contoso uses Microsoft Entra ID for identity management, authentication, and authorization capabilities that enable agents to access organizational resources and services.

Contoso recently formed a new AI engineering team named Agent1Dev Team to optimize and maintain existing AI solutions.

The team collaborates with solution architects, DevOps engineers, and security engineers to design, implement, monitor, and secure AI applications.

Contoso also has a team named Agent1Test Team that is responsible for validating AI solutions before the solution deployments.

Generative Environment

Contoso has a Microsoft Foundry deployment that contains two projects named Project1 and Project2.

Project1

Project1 contains a customer support agent named Agent1 that assists customers with product inquiries and troubleshooting requests.

Agent1 has the following configurations:

- Agent1 uses a base model deployment.**
- A safety evaluation pipeline is NOT enabled.**
- Tool invocation approval workflows are NOT enabled.**
- Conversation memory constraints are NOT configured.**

Agent1 interacts with customers by using digital support channels and answers general questions about Contoso products.

Project1 is deployed to an Azure region located in the European Union (EU).

Agent1Dev Team will use Project1 to optimize and maintain Agent1.

Project2

Project2 contains a deployed video generation model. The marketing department at Contoso has access to Project2 and plans to use the model to develop a video creation solution.

Development of the solution is incomplete.

Data Environment

Contoso stores product-related information in Azure resources that support AI applications.

The Azure environment contains an Azure Blob Storage account named storage1 that stores product detail sheets for all the Contoso products.

The product sheets include specifications, feature descriptions, and product support information that Agent1 can use to answer customer questions. The product sheets are stored in the PDF format.

Problem Statements

Contoso identifies the following issues:

- Agent1 has only general knowledge of the Contoso products.**
- A recent chat interaction with Agent1 was analyzed for sentiment. The results of the analysis have NOT been processed yet.**
- Agent1 does NOT use the detailed product information in the product sheets stored in storage1 when responding to customer questions.**
- The finance department at Contoso reports that vendor invoices must be reviewed manually to ensure that the invoices match the terms defined in the vendor contracts. The invoices contain tables, logos, and varied layouts that make the documents difficult to process consistently.**

Requirements

Planned Changes

Contoso plans to implement the following changes:

- Implement a solution for Project1 that analyzes the vendor invoices by evaluating both the visual layout and the textual content of the invoices, so that the invoice details can be verified against the vendor contract terms.**
- Update the base model deployment used by Agent1 and standardize the model version to ensure continuity and consistent responses.**
- Enable Agent1 to retrieve and use the detailed product information from the product sheets stored in storage1.**
- Implement an indexing solution for the product sheets that Agent1 can use to answer customer questions.**
- Complete the development of the video creation solution.**

Technical Requirements

Contoso identifies the following technical requirements:

- The model deployment used by Agent1 must support scalable, high-throughput generative AI workloads and dynamically scale to handle variable customer support traffic, without requiring reserved throughput capacity.**
- The product sheets must be processed by using an indexing pipeline that enables semantic and vector search, so that Agent1 can retrieve the relevant product**

information.

- Responses generated by using the product sheet information must be relevant, complete, and accurate.
- Agent1 must be able to use the product sheets to answer natural language questions about product details.
- The model version used by Agent1 must remain consistent to ensure stable responses.
- The data processed by the model must remain within the EU.

Security and Compliance Requirements

Contoso identifies the following security and compliance requirements:

- API keys must NOT be used to access Foundry-deployed models.
- Access to the Azure resources must follow the principle of least privilege.
- The developers at Contoso must authenticate to Microsoft Foundry resources by using Microsoft Entra authentication.
- Access to Project1 must be assigned to the members of Agent1Dev Team by using a security group named SC_Agent1_Dev.
- Access to Project1 must be assigned to the members of Agent1Test Team by using a security group named SC_Agent1_Test.
- Agent1 must never reveal customer information, even if a document that contains customer data is added erroneously to the product sheet repository in storage1.
- The product sheets might contain images that include embedded text. Agent1 must be protected from malicious instructions potentially hidden within the images.

Business Requirements

Contoso identifies the following business requirements:

- Users that interact with Agent1 must have a personalized experience in future interactions, including the ability for Agent1 to retain conversation context and recall relevant information from previous interactions.
- Agent1 must answer questions only about the products sold by Contoso.



Hotspot Question

You need to ensure that the marketing department can generate videos by using the model deployed to Project2.

How should you complete the Python code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
import time
from openai import OpenAI
client = OpenAI(
    base_url= "https://Contoso.openai.azure.com/openai/v1/",
    api_key=...,
)
video = client.videos.  (
    create
    download_content
    list
    retrieve
    model=deployment_name,
    prompt= "A video of our products",
)
while video.status not in ["completed", "failed", "cancelled"]:
    time.sleep(20)
    video = client.videos.  (video.id)
    create
    download_content
    list
    retrieve

    print(video.status)
```

Answer:

Answer Area

```
import time
from openai import OpenAI
client = OpenAI(
    base_url= "https://Contoso.openai.azure.com/openai/v1/",
    api_key=...,
)
video = client.videos. (
    model=deployment_name,
    prompt= "A video of our products",
)
while video.status not in ["completed", "failed", "cancelled"]:
    time.sleep(20)
    video = client.videos. (video.id)

    print(video.status)
```

Explanation:

Box 1: create

`client.videos.create(...)` initializes the generation job. This returns a video object containing a unique id and a status of "queued" or "in_progress".

Box 2: retrieve

`client.videos.retrieve(video.id)` refreshes the video object state from the server by fetching its updated completion progress and status until it finishes Reference:

<https://developers.openai.com/api/reference/python/resources/videos/methods/create>

QUESTION: 84

Drag and Drop Question

You have a Microsoft Foundry project that contains a multi-agent solution. The agents use tool calling to query internal systems.

You need to implement responsible AI auditing to meet the following requirements:

- Capture all the nested operations across the entire agent run.
- Record tool invocation arguments and returned results as metadata.

What should you use for each requirement? To answer, drag the appropriate options to the correct targets Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Options

Answer Area

- ☐ Hierarchical spans
- ☐ A KQL query filter
- ☐ Sampling
- ☐ Tool call attributes
- ☐ Trace sampling policy

Capture all the nested operations across the entire agent run:

Record tool invocation arguments and results:

Answer:

Options

Answer Area

- ☐ Hierarchical spans
- ☐ A KQL query filter
- ☐ Sampling
- ☐ Tool call attributes
- ☐ Trace sampling policy

Capture all the nested operations across the entire agent run:

Record tool invocation arguments and results:

Explanation:

Box 1: Hierarchical spans

Hierarchical spans should be used in this case.

Parent-Child tracking: In OpenTelemetry and Azure Monitor Application Insights (which back Microsoft Foundry environments), hierarchical spans utilize a unique Trace ID for the entire request and distinct Span IDs for individual operations.

Nested capture: When a primary agent calls a sub-agent, or an agent invokes a specific system tool, each subsequent operation is recorded as a child span. This explicitly preserves the parent- child relationship, allowing auditors to reconstruct the exact execution tree of all nested operations across the entire multi-agent run

Box 2: Tool call attributes To record tool invocation arguments and returned results as metadata for

responsible AI auditing in a Microsoft Foundry multi-agent project, you should use Tool call attributes.

In GenAI and multi-agent systems leveraging OpenTelemetry semantic conventions (which Microsoft Foundry utilizes for its observability, tracing, and logging pillars), tool execution details are captured via specific span attributes. Configured tool call attributes explicitly map and record the parameters (tool.call.arguments) and outputs (tool.call.result) as metadata keys inside the tracing context for auditing and validation.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry/observability/concepts/trace-data>

<https://techcommunity.microsoft.com/blog/azure-ai-foundry-blog/ai-observability-starter-kit-for-microsoft-foundry-agents/4522751>

QUESTION: 85

You have an Azure subscription that contains an Azure AI Foundry hub named Hub1, an Azure OpenAI resource named resource1, and a user named User1.

You need to ensure that User1 can create a new Azure AI Content Understanding project in Hub1. The solution must follow the principle of least privilege.

Which role should you assign to User1?

- A.** Cognitive Services OpenAI User
- B.** Cognitive Services OpenAI Contributor
- C.** Azure AI Administrator
- D.** Azure AI Developer

Answer: B

Cognitive Services Contributor

This role is typically granted access at the resource group level for a user in conjunction with additional roles. By itself this role would allow a user to perform the following tasks.

- * Create new Azure OpenAI resources within the assigned resource group.

- * Etc.

To create an Azure AI Content Understanding project on an Azure AI Foundry hub, you need either the Azure AI Account Owner, Contributor, or Cognitive Services Contributor role at the subscription level to create the necessary Azure AI Foundry resource, and then the user must have the Azure AI User role at the project level to create the agent within that project.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-foundry/openai/how-to/role-based-access-control>

QUESTION: 86

You have a Microsoft Foundry project that contains an agent named Agent1. Agent runs successful, but Foundry Control Plane does NOT display values for error rates, runs, and token usage, and the Traces tab is empty.

You need to ensure that Found Control Plane displays the appropriate values for Agent1.

What should you do?

- A.** Update Agent1 to a new version.
- B.** Restart Agent from Foundry Control Plan
- C.** Assign to a Log Analytics workspace to Agent1.
- D.** Enable Application Insights for Agent1.

Answer: D

To resolve this issue, you must connect and configure an Azure Application Insights resource for your Microsoft Foundry project.

The Foundry Control Plane, its Agent Monitoring Dashboard, and the Traces tab rely directly on telemetry data stored within the connected Application Insights instance. If this resource is missing, unlinked, or improperly configured, the dashboard cannot display runs, error rates, token usage, or transaction spans.

Reference:

<https://techcommunity.microsoft.com/blog/azure-ai-foundry-blog/monitoring--observability-in-microsoft-foundry/4517250>

QUESTION: 87

You are developing an application that will use Azure AI Search for internal documents.

You need to implement document-level filtering for Azure AI Search.

Which three actions should you include in the solution? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A.** Add allowed groups to each index entry
- B.** Create one index per group.

- C. Send access tokens from Microsoft Entra ID, with the search request.
- D. Retrieve all the groups.
- E. Retrieve the group memberships of the user
- F. Supply the groups as a filter for the search requests

Answer: A,E,F

[A]

Add allowed groups to each index entry: Each document in your search index must include a field (e.g., group_ids) listing the security groups or principals that have permission to view that document.

[E]

Retrieve the group memberships of the user: When a user submits a query, your application must query Microsoft Entra ID to fetch the list of all security groups the user belongs to.

[F]

Supply the groups as a filter for the search requests: Your application will then translate the user's group memberships into an OData filter (using the search.in() function) and append it to the search API query, restricting the results only to documents the user has access to.

Reference:

<https://learn.microsoft.com/en-us/azure/search/search-security-trimming-for-azure-search>

QUESTION: 88

You build a chatbot that uses the Azure OpenAI GPT-4 model to generate song lyrics.

You need to ensure that responses do NOT contain lyrics from popular songs that might have been ingested during model training.

Which Azure AI Content Safety API should you use?

- A. Custom categories
- B. Analyze Text
- C. Groundedness detection
- D. Protected material text detection

Answer: D

The Protected material text detection feature in Azure AI Content Safety is specifically designed to identify and block responses that may contain copyrighted or protected content, such as lyrics from popular songs. Since GPT-4 may have been trained on publicly available data, this feature helps ensure that the chatbot does not generate copyrighted lyrics by detecting and filtering out protected material.

QUESTION: 89

You are building a voice agent for a pharmacy. It must accept spoken questions, reply with synthesised speech, and accurately recognise specialised medicine names that a general speech model often mishears. Which capability should you configure?

- A.** A custom speech model for speech-to-text within the agent
- B.** Only the default text-to-speech voice
- C.** Azure Translator to convert the audio
- D.** Azure AI Document Intelligence on the audio stream

Answer: A

A custom speech model improves recognition of specialised vocabulary, such as medicine names, by adapting speech-to-text (STT) to your domain, and AI-103 covers integrating speech, including custom speech models, as an agent modality. This directly addresses the accuracy problem in the scenario.

QUESTION: 90

You have a custom agent named Agent1.

You need to control access to and monitor activity for Agent1 by using Microsoft Foundry.

What should you do first?

- A.** Add Agent1 to the Microsoft Foundry playground.
- B.** Add Agent1 to a Microsoft Foundry project.
- C.** Provision an Application Insights resource.
- D.** Create a Microsoft Foundry project.

Answer: D

QUESTION: 91

You have a product support manual.

You need to build a product support chatbot based on the manual. The solution must minimize development effort and costs.

What should you use?

- A.** Azure OpenAI GPT-4 with grounding data that uses Azure AI Search
- B.** Azure AI Document intelligence
- C.** Azure AI Language Custom question answering
- D.** Azure AI Phi-3-medium with fine-tuning

Answer: A

Azure OpenAI On Your Data makes it easier for developers to connect, ingest and ground their enterprise data to create personalized copilots (preview) rapidly. It enhances user comprehension, expedites task completion, improves operational efficiency, and aids decision- making.

Azure OpenAI On Your Data enables you to run advanced AI models such as GPT-35-Turbo and GPT-4 on your own enterprise data without needing to train or fine-tune models. You can chat on top of and analyze your data with greater accuracy. You can specify sources to support the responses based on the latest information available in your designated data sources.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/use-your-data>

QUESTION: 92

You have a Microsoft Foundry project that contains an agent for a customer support chat app.

The agent uses a memory store and a memory search tool.

You need to ensure that the conversation history does NOT persist across separate sessions.

To what should you set the scope of the memory tool?

- A.** session
- B.** {{\$conversationId}}
- C.** {{\$userId}}

D. global

Answer: A

To ensure that conversation history does not persist across separate sessions, you must set the scope to session.

The session scope restricts data access to the current active chat instance. It automatically wipes or ignores previous data when a new session starts.

Incorrect:

[Not C]

user or {{\$userId}} scope: Persists data across multiple sessions for that specific user. This would cause the exact issue you want to avoid by carrying historical context into new conversations.

References:

<https://ai.gopubby.com/agents-with-memory-conceptual-undestanding-part-01-f6caedfcd96d?gi=a9aa95df15c6>

QUESTION: 93

You have an Azure subscription that contains an Azure OpenAI resource.

You need to build an app that will suggest product names from a given product description.

Which model should you use?

A. Whisper

B. DALL-E

C. embeddings

D. GPT-4

Answer: D

To suggest product names from a given product description, the ideal model is GPT-4. GPT-4 is a powerful language model that can understand and generate human-like text based on given inputs, making it well-suited for tasks such as generating creative suggestions like product names from descriptions.

Whisper is used for speech-to-text conversion, not for generating text suggestions.

DALL-E is designed for generating images from text prompts, not for text generation tasks like naming products.

Embeddings are used for tasks involving similarity search and semantic understanding but not for generating new text like product names.

QUESTION: 94

You plan to build an agent that will combine and process multiple files uploaded by users.

You are evaluating whether to use the Azure AI Agent Service to develop the agent. What is the maximum size of each file that can be uploaded to the service?

- A. 10 MB**
- B. 256 MB**
- C. 512 MB**
- D. 1 GB**

Answer: C

<https://learn.microsoft.com/en-us/azure/ai-services/agents/quotas-limits>

QUESTION: 96

You have a Microsoft Foundry agent that grounds responses from an Azure Search index that contains the following:

- Searchable text fields for product names and product codes**
- A vector field that stores embeddings for product descriptions**

You need to ensure that users can query the index by using the following:

- Exact product names or codes**
- Natural language descriptions of the products**

What should you configure?

- A. vector search only**
- B. hybrid search**
- C. keyword search only**
- D. semantic search only**

Answer: B

To meet your requirements, you need to configure a Hybrid Search with Semantic Ranking in Azure AI Search. This setup combines keyword matching for exact identifiers with vector search for natural language queries, delivering the most accurate grounding data to your Microsoft Foundry agent.

Reference:

<https://www.tredence.com/blog/searchsmart-enhancing-rag-with-azure-ai-search-service>

QUESTION: 97

In mid-2026 you are starting a brand-new agentic application on Azure and want to build on the generally available, supported entry point for Foundry agents. Which API should you target?

- A. The Assistants API
- B. The Responses API via Foundry Agent Service
- C. The legacy text Completions API
- D. The LUIS runtime API

Answer: B

The Responses API is the generally available single entry point for Foundry Agent Service and supports both prompt agents and hosted agents. It is the API Microsoft directs new agentic development towards.

QUESTION: 98

A team wants to deploy a model from the Foundry catalogue with the fastest possible setup, pay only for the tokens they consume, and avoid provisioning any virtual machine (VM) quota. Which deployment option fits?

- A. Managed compute (managed endpoint) deployment
- B. Serverless API deployment, also called models as a service (MaaS)
- C. Self-hosted deployment on Azure Kubernetes Service
- D. A local container running on a developer workstation

Answer: B

Serverless API deployment runs on Microsoft-managed infrastructure, bills per input and output token, and needs no VM quota, so it is the fastest path that matches token-based billing. It is the standard choice when you want pay-per-use economics without managing compute.

QUESTION: 99

Hotspot Question

You need to create a new resource that will be used to perform sentiment analysis and optical character recognition (OCR). The solution must meet the following requirements:

- Use a single key and endpoint to access multiple services.
- Consolidate billing for future services that you might use.
- Support the use of Azure Vision in Foundry Tools in the future.

How should you complete the HTTP request to create the new resource? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

	▼	https://management.azure.com/subscriptions/
PATCH		xxxxxxxx-xxxx-xxxx-xxxx-
POST		xxxxxxxxxxxx/resourceGroups/RG1
PUT		

/providers/Microsoft.CognitiveService/accounts/CS1?api-version=2021-04-30

```
{  
  "location": "West US",  
  "kind": "  
    CognitiveServices  
    ComputerVision  
    TextAnalytics  
  ",  
  "sku": {  
    "name": "S0"  
  },  
  "properties": {},  
  "identity": {  
    "type": "SystemAssigned"  
  }  
}
```


Answer:

Answer Area

	▼	https://management.azure.com/subscriptions/ xxxxxxxx-xxxx-xxxx-xxxx- xxxxxxxxxxxx/resourceGroups/RG1
PATCH		
POST		
PUT		

/providers/Microsoft.CognitiveService/accounts/CS1?api-version=2021-04-30

```
{  
  "location": "West US",  
  "kind": "
```

Foundry Tools to analyze marketing videos.

Video segmentation is enabled.

You need to configure an analyzer to output a generated JSON field that describes the color scheme of each video segment.

How should you configure the analyzer? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values	Answer Area
classify	
generate	Field value type:
group	Field method:
string	
table	

Answer:

Values	Answer Area
classify	
generate	Field value type: string
group	Field method: generate
string	
table	

Explanation:

Box 1: string

The Type (string): While you want the output schema to contain a description of the color palette or style, Azure Content Understanding schemas map the underlying data types

using fundamental types like string, number, or array. The field will output the final result in the JSON response under a key like "valueString" Box 2: generate The Method (generate): The extract method is only used for pulling literal text exactly as it appears in content (such as speech-to-text transcripts or OCR). Because analyzing visual aesthetics and synthesizing descriptive text about a "color scheme" requires multimodal AI reasoning, you must use the generate method. This instructs the underlying model to analyze the segment and synthesize a completely new descriptive insight.

References:

<https://learn.microsoft.com/en-us/azure/ai-services/content-understanding/video/elements>

QUESTION: 101

A customer support agent must answer from an internal knowledge base that is updated weekly with new product documentation, and its answers must always reflect the latest content. Which approach should you implement?

- A.** Fine-tune the base model on the new documents every week
- B.** Implement retrieval-augmented generation (RAG) over an index of the documents
- C.** Raise the model temperature to improve recall of recent facts
- D.** Apply Direct Preference Optimization (DPO) to the documents

Answer: B

RAG injects fresh, frequently changing knowledge at inference time by retrieving relevant content from an index and grounding the model's answer in it, which is exactly the "chat with your data" pattern this scenario describes. Microsoft's guidance is direct on the boundary between the two techniques.

Use fine-tuning when you need to change model behavior, style, or task performance, rather than add fresh knowledge.

QUESTION: 102

You have a web app named App1 that processes user prompts by integrating with a Microsoft Foundry project named Project1. App1 performs the following actions:

- Sends prompts directly to a model by using the Azure OpenAI Responses API**
- Invokes the Azure AI Content Safety tool by using a Foundry connection within the same request**

You need to configure end-to-end visibility into each step of the request workflow. What should you do?

- A. Enable logging by using the client SDK for Content Safety.
- B. Enable logging by using Foundry Local.
- C. Enable application tracing in Project1.
- D. Route requests through the Azure OpenAI endpoint.

Answer: C

To configure end-to-end visibility into each step of the request workflow, you should Enable application tracing in the project.

Project-Level Observability: Enabling application tracing allows Microsoft Foundry to natively track, trace, and monitor every component of the workflow-including calls to the Azure OpenAI Responses API and associated tools like Azure AI Content Safety-within the same request without requiring complex, separate SDK configurations.

Centralized Tracking: Tracing captures the exact execution path across the project endpoints, allowing you to see metrics, latency, and inputs/outputs sequentially for both the model invocation and the content safety check.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry-classic/how-to/develop/trace-production-sdk>

QUESTION: 103

Hotspot Question



You have a Python application collects customer comments before posting them to a public forum.

You need to send a text comment to Azure AI Content Safety and return the self-harm severity from the response.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
def get_self_harm_severity (comment: str) -> int:
    key = os.environ["CONTENT_SAFETY_KEY"]
    endpoint = os.environ
["CONTENT_SAFETY_ENDPOINT"]
    client = ContentSafetyClient(endpoint,
AzureKeyCredential(key))
    request = 
    response = 
    result = next(
        item for item in response.categories_analysis
        if item.category == TextCategory.SELF_HARM
    )
    return result.severity
```

The first dropdown menu for 'request' contains the following options:



- AnalyzeTextOptions(categories=comment)
- AnalyzeTextOptions(text=[comment])
- AnalyzeTextOptions(text=comment)
- TextCategory.SELF_HARM(comment)

The second dropdown menu for 'response' contains the following options:

- client.analyze_image(request)
- client.analyze_text(request)
- client.moderate_text(request)
- client.path("/text:analyze").post(request)

Answer:

Answer Area

```
def get_self_harm_severity (comment: str) -> int:
    key = os.environ["CONTENT_SAFETY_KEY"]
    endpoint = os.environ
["CONTENT_SAFETY_ENDPOINT"]
    client = ContentSafetyClient(endpoint,
AzureKeyCredential(key))
    request = 
    response = 
    result = next(
        item for item in response.categories_analysis
        if item.category == TextCategory.SELF_HARM
    )
    return result.severity
```

Explanation:

Box 1: AnalyzeTextOptions(text=comment)

To set up the request in your Python script, you must use

AnalyzeTextOptions(text=comment), passing the text as a single string rather than a list.

Example:

Build the request options focusing on Text analysis

You can optionally restrict analysis to specific categories

```
request_options = AnalyzeTextOptions(
    text=user_comment,
    categories=[TextCategory.SELF_HARM]
```

Box 2: client.analyze_text(request)

You must use client.analyze_text(request)

1. Send the request: Call client.analyze_text(request) to retrieve the multi-severity results.
2. Extract the self-harm result: Locate the self-harm categories from the response list.
3. Return the severity: Access the .severity attribute

Reference:

<https://learn.microsoft.com/en-us/python/api/overview/azure/ai-contentssafety-readme>

QUESTION: 104

You have an Azure subscription.

You plan to build an app that will use the Azure OpenAI DALL-E model.

You need to deploy the model.

What should you use?

- A.** the Azure SDK for Python and PowerShell cmdlets
- B.** the Azure portal and Microsoft Graph API
- C.** the Azure SDK for JavaScript and Azure Machine Learning Studio
- D.** Azure AI Studio and Azure Command-Line Interface (CLI)

Answer: D

To deploy the Azure OpenAI DALL-E model, the most appropriate tools to use are Azure AI Studio for model management and interaction, and the Azure Command-Line Interface (CLI) for deployment and configuration tasks. Azure AI Studio provides an interactive interface for working with Azure AI models, and the Azure CLI allows you to manage resources, deployments, and other Azure services from the command line.

QUESTION: 105

You have a Microsoft Foundry project that generates short promotional product videos.

After several clips are approved, reviewers notice a small watermark in the top-right corner of some videos.

You need to remove the watermark without regenerating the videos.

What should you do?

- A.** Modify the original prompt to exclude watermarks.
- B.** Crop the video by using the size parameter.
- C.** Increase the guidance scale.
- D.** Apply a mask-based inpainting edit to the affected part of the video.

Answer: D

The correct action to take is to apply a mask-based inpainting edit to the affected part of the video.

This is the only option that directly modifies the existing, approved video files. Inpainting allows you to isolate the specific top-right corner using a mask and seamlessly blend it with the surrounding pixels to erase the watermark without altering or regenerating the rest of the video footage.

QUESTION: 106

You have a Microsoft Foundry project.

You need to deploy a model from the model catalog to support real-time inference.

The solution must meet the following requirements:

- Use key-based authentication.**
- Support real-time REST API access.**
- NOT consume the vCPU quota of the virtual machines in the Azure subscription.**

Which type of deployment should you use?

- A. serverless API**
- B. batch**
- C. self-hosted container**
- D. standard**

Answer: D

The most appropriate type of deployment is a serverless API deployment (also referred to as a standard deployment or Models-as-a-Service / MaaS).

Key-Based Authentication: Serverless API deployments natively provision an endpoint URL alongside primary and secondary API keys to secure your client applications.

Real-Time REST API Access: When the model is successfully deployed, it exposes a scalable, real-time HTTP/REST endpoint matching the standardized Azure AI Model Inference API.

No Virtual Machine vCPU Quota Consumption: Unlike managed compute deployments - which provision dedicated virtual machines in your subscription and require VM vCPU quota - serverless API deployments are completely hosted and managed by Microsoft. They operate on a pay-as-you-go, token-based billing structure and do not consume any VM vCPU quota from your Azure subscription.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry/concepts/deployments-overview>

QUESTION: 107

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

You need to improve response completeness. The solution must be implemented in the logic of the application code before responses are returned.

What should you do?

- A.** Add a retry evaluation before the responses are returned.
- B.** Decrease the value of the max_tokens parameter.
- C.** Switch to Retrieval Augmented Generation (RAG).
- D.** Replace the model with a smaller deployment.

Answer: A

To enhance response completeness in your Microsoft Foundry agent, you must intercept the retrieved documents and the generated summary within your backend application logic before returning the payload to the user.

1. Implement Completeness Verification Logic

Add a verification step in your orchestration code (e.g., in your Python/Semantic Kernel or LangChain pipeline) that compares the generated summary against the retrieved chunks.'

Map Key Assertions: Extract main policy rules from retrieved text.

Cross-Reference Entities: Verify all key entities are in the summary.

Check Scope Coverage: Ensure every retrieved document is represented.

Scan for Gaps: Identify critical missing constraints or exceptions.

2. Apply Application-Level Mitigation Strategies

If the verification step detects that the summary is incomplete, use your code to correct it before the final response leaves your system.

Reference:

<https://dev.to/moonrunnerkc/how-i-built-a-verification-layer-for-copilot-clis-multi-agent-output-4b7h>

QUESTION: 108**Drag and Drop Question**

You are developing an application that will detect faulty components produced on a factory production line. The components are specific to your business.

You need to use the Azure Custom Vision API to help detect common faults.

Which three actions should you perform in sequence? To answer, move the

appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Initialize the training dataset.	
Train the classifier model.	1
Create a project.	2
Upload and tag images.	3
Train the object detection model.	

Answer:

Actions	Answer Area
Initialize the training dataset.	
Train the classifier model.	1 Create a project.
Create a project.	2 Upload and tag images.
Upload and tag images.	3 Train the object detection model.
Train the object detection model.	

Explanation:

Step 1: Create a project

You must first set up a new project in the Custom Vision portal or via the SDK to store your data and configuration.

Step 2: Upload and tag images

Next, you need to upload photos of your specific factory components and draw bounding boxes around the faults to tag them.

Step 3: Train the object detection model

Because you need to identify and locate specific, localizable faults on a component, you must train an object detection model rather than a general classifier model.

Reference:

<https://thegroundtruth.blog/tag/computervision/>

QUESTION: 109

Hotspot Question

You have a Microsoft Foundry project that contains two agents named PolicyWriter and RskReviewer.

PolicyWriter generates draft updates for customer policies, and RiskReviewer reviews the drafts.

In the visual builder, you need to create a workflow that meets the following requirements:

- Finalizes low-risk updates without manual intervention
- Ensures predictable execution across the agents
- Requires user approval for high updates

What should you configure? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Orchestration pattern:


	▼
The sequential template that passes outputs node by-node	
The group chat template to dynamically route control between the agents	
The human-in-the-loop template that pauses execution of the workflow for input	


Approval checkpoints:

	▼
Add a Basic chat node.	
Add a Condition statement.	
Add an Ask a question node.	

Answer:

Answer Area

Orchestration pattern: 

Approval checkpoints: 

Explanation:

Box 1: The human-in-the-loop template that pauses execution of the workflow for input Human-in-the-loop template is correct. This pattern allows the system to run automatically for predictable, node-by-node agent execution while explicitly pausing the workflow when a high-risk condition is met to wait for human intervention.

Box 2: Add a Condition statement

Condition statement is correct. You must evaluate the risk level (low vs. high) before deciding whether to finalize the policy. A condition node checks the risk score output from RiskReviewer. If it is low, it routes to auto-finalization. If it is high, it routes to an approval step.

Reference:

<https://techcommunity.microsoft.com/blog/azure-ai-foundry-blog/building-human-in-the-loop-ai-workflows-with-microsoft-agent-framework/4460342>

QUESTION: 110

You are defining an agent in Microsoft Foundry Agent Service. The agent uses a catalogue model for reasoning, a system instruction that sets its goals, and a file search capability that reads a knowledge store. Which three components does this combination represent?

- A. Endpoint, API key, deployment name
- B. Model, instructions, tools
- C. Trigger, action, condition
- D. Embedding, retriever, generator

Answer: B

Microsoft Foundry Agent Service defines every agent as the combination of three core components: a model that provides reasoning, instructions that set goals and behaviour, and tools that give access to data or actions. The scenario maps exactly, with the catalogue model, the system instruction, and the file search tool.

QUESTION: 112

You have an Azure AI agent solution.

You plan to create an agent-based app named App1 that will analyze and summarize data for users and generate data-driven recommendations. App1 will be used by non-technical business users and must adapt to new and unforeseen business challenges and improve its performance over time.

You need to identify which type of agent to use in App1. The solution must meet the following requirements:

- Adapt and improve the agents' performance over time based on user feedback.**
- Provide tailored recommendations to help users make informed decisions.**
- Provide the best possible performance of the app.**

Which agent type should you identify?

- A. cognitive**
- B. prompt-and-response**
- C. autonomous**
- D. task automation**

Answer: A

A cognitive agent is designed to mimic human thought processes, making it ideal for the complex, data-driven tasks described in the query.

Analyze and summarize data and generate data-driven recommendations:

Cognitive agents use machine learning and natural language processing to analyze historical data and act on their learnings, which includes generating insights and recommendations.

Adapt and improve over time based on user feedback:

They are designed to learn and adapt to user behavior and preferences over time, improving their responses and performance through feedback.

Used by non-technical business users:

These agents can be accessed through user-friendly interfaces, often natural language-based (like virtual assistants), making them suitable for non-technical users.

Provide tailored recommendations to help users make informed decisions:

By learning from data and interactions, they can deliver personalized recommendations that support better decision-making.

QUESTION: 113

You have a factory that contains a production line.

You need to develop a solution that will analyze a video feed to detect and track the movement of employees as they work on the production line.

Which service should you include in the solution?

- A.** Azure AI Video Indexer
- B.** Azure Vision in Foundry Tools Spatial Analysis
- C.** Azure Vision in Foundry Tools Image Analysis
- D.** Azure AI Face

Answer: B

Azure Vision in Foundry Tools (Spatial Analysis) can be used to analyze a video feed to detect and track the movement of employees as they work on the production line. This technology processes live or recorded video to understand spatial relationships and human activity.

Key capabilities for this application include:

Person Tracking & Movement: Spatial Analysis models identify human presence by detecting a body bounding box and tracking it in real time, enabling monitoring of movement patterns throughout a facility.

Zone & Line Monitoring: You can configure the system to detect when a person enters or exits a designated zone (e.g., a specific workstation) or crosses a line.

Dwell Time Calculation: The system can measure how long an employee spends in a specific zone, which can be useful for identifying bottlenecks.

Safety and Compliance: It can monitor for safety infractions, such as workers entering prohibited areas or failing to wear protective equipment (PPE).

Efficiency Analysis: By integrating these insights into Palantir Foundry's ontology, manufacturers can track assembly processes, optimize workflows, and enhance productivity.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/whats-new>

QUESTION: 114

You have an Azure subscription that contains an Azure OpenAI resource.

You deploy the GPT-4 model to the resource.

You need to ensure that you can upload files that will be used as grounding data for the model.

Which two types of resources should you create? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

A. Azure AI Document Intelligence

B. Azure AI Search

C. Azure SQL

D. Azure Blob Storage

E. Azure AI Bot Service

Answer: B,D

Azure OpenAI On Your Data enables you to run advanced AI models such as GPT-35-Turbo and GPT-4 on your own enterprise data without needing to train or fine-tune models.

For some data sources such as uploading files from your local machine (preview) or data contained in a blob storage account (preview), Azure AI Search is used.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/openai/concepts/use-your-data>

QUESTION: 115

Case Study 1 - Contoso, Ltd

Overview

Company Information

Contoso, Ltd is a multinational retail company that builds, deploys, and manages generative AI and agent-based solutions by using Microsoft Foundry.

Existing Environment

Identity Environment

Contoso uses Microsoft Entra ID for identity management, authentication, and

authorization capabilities that enable agents to access organizational resources and services.

Contoso recently formed a new AI engineering team named Agent1Dev Team to optimize and maintain existing AI solutions.

The team collaborates with solution architects, DevOps engineers, and security engineers to design, implement, monitor, and secure AI applications.

Contoso also has a team named Agent1Test Team that is responsible for validating AI solutions before the solution deployments.

Generative Environment

Contoso has a Microsoft Foundry deployment that contains two projects named Project1 and Project2.

Project1

Project1 contains a customer support agent named Agent1 that assists customers with product inquiries and troubleshooting requests.

Agent1 has the following configurations:

- Agent1 uses a base model deployment.
- A safety evaluation pipeline is NOT enabled.
- Tool invocation approval workflows are NOT enabled.
- Conversation memory constraints are NOT configured.

Agent1 interacts with customers by using digital support channels and answers general questions about Contoso products.

Project1 is deployed to an Azure region located in the European Union (EU).

Agent1Dev Team will use Project1 to optimize and maintain Agent1.

Project2

Project2 contains a deployed video generation model. The marketing department at Contoso has access to Project2 and plans to use the model to develop a video creation solution.

Development of the solution is incomplete.

Data Environment

Contoso stores product-related information in Azure resources that support AI applications.

The Azure environment contains an Azure Blob Storage account named storage1 that stores product detail sheets for all the Contoso products.

The product sheets include specifications, feature descriptions, and product support information that Agent1 can use to answer customer questions. The product sheets are stored in the PDF format.

Problem Statements

Contoso identifies the following issues:

- Agent1 has only general knowledge of the Contoso products.
- A recent chat interaction with Agent1 was analyzed for sentiment. The results of the analysis have NOT been processed yet.
- Agent1 does NOT use the detailed product information in the product sheets stored in storage1 when responding to customer questions.
- The finance department at Contoso reports that vendor invoices must be reviewed manually to ensure that the invoices match the terms defined in the vendor contracts. The invoices contain tables, logos, and varied layouts that make the documents difficult to process consistently.

Requirements

Planned Changes

Contoso plans to implement the following changes:

- Implement a solution for Project1 that analyzes the vendor invoices by evaluating both the visual layout and the textual content of the invoices, so that the invoice details can be verified against the vendor contract terms.
- Update the base model deployment used by Agent1 and standardize the model version to ensure continuity and consistent responses.
- Enable Agent1 to retrieve and use the detailed product information from the product sheets stored in storage1.
- Implement an indexing solution for the product sheets that Agent1 can use to answer customer questions.
- Complete the development of the video creation solution.

Technical Requirements

Contoso identifies the following technical requirements:

- The model deployment used by Agent1 must support scalable, high-throughput generative AI workloads and dynamically scale to handle variable customer support traffic, without requiring reserved throughput capacity.
- The product sheets must be processed by using an indexing pipeline that enables semantic and vector search, so that Agent1 can retrieve the relevant product information.
- Responses generated by using the product sheet information must be relevant, complete, and accurate.
- Agent1 must be able to use the product sheets to answer natural language questions about product details.
- The model version used by Agent1 must remain consistent to ensure stable responses.
- The data processed by the model must remain within the EU.

Security and Compliance Requirements

Contoso identifies the following security and compliance requirements:

- API keys must NOT be used to access Foundry-deployed models.
- Access to the Azure resources must follow the principle of least privilege.
- The developers at Contoso must authenticate to Microsoft Foundry resources by using Microsoft Entra authentication.
- Access to Project1 must be assigned to the members of Agent1Dev Team by using a security group named SC_Agent1_Dev.
- Access to Project1 must be assigned to the members of Agent1Test Team by using a security group named SC_Agent1_Test.
- Agent1 must never reveal customer information, even if a document that contains customer data is added erroneously to the product sheet repository in storage1.
- The product sheets might contain images that include embedded text. Agent1 must be protected from malicious instructions potentially hidden within the images.

Business Requirements

Contoso identifies the following business requirements:

- Users that interact with Agent1 must have a personalized experience in future interactions, including the ability for Agent1 to retain conversation context and recall relevant information from previous interactions.
- Agent1 must answer questions only about the products sold by Contoso.

You need to configure personalized user interactions for Agent1. The solution must meet the business requirements.

What should you include in the solution?

- A.** knowledge
- B.** memory
- C.** guardrails
- D.** tools

Answer: B

Scenario, Business Requirements:

*-> Users that interact with Agent1 must have a personalized experience in future interactions, including the ability for Agent1 to retain conversation context and recall relevant information from previous interactions.

Agent1 must answer questions only about the products sold by Contoso.

To configure personalized user interactions that retain context and recall previous information, you should use the Memory in Foundry Agent Service.

Managed Long-Term Memory: This allows the agent to extract meaningful information from user conversations and store it as durable knowledge (e.g., user preferences or historical details) to be recalled across future interactions.

Scoped Memory Stores: By using the scope parameter when creating memory stores, you can segment memories to ensure a secure, personalized experience for distinct users.

Automatic Injection: The agent will automatically inject these learned facts and user profile details into the prompt before new interactions, eliminating the need to rehydrate context manually.

References:

<https://learn.microsoft.com/en-us/azure/foundry/agents/concepts/what-is-memory>

QUESTION: 116

You need to add an automated grounding check to a RAG application's continuous evaluation.

The check must return a simple pass or fail result and must not require you to deploy a separate judge model. Which evaluator should you use?

- A.** Groundedness, which returns a 1 to 5 score
- B.** Groundedness Pro
- C.** Fluency
- D.** Relevance

Answer: B

Groundedness Pro returns a binary pass or fail result and runs on the Azure AI Content Safety service, so it does not require you to deploy a model to act as a judge. That matches both requirements in the scenario.

QUESTION: 117

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

You need to improve response completeness. The solution must be implemented in the logic of the application code before responses are returned.

What should you do?

- A.** Add a retry evaluation before the responses are returned.
- B.** Decrease the value of the temperature parameter.

- C. Increase the value of the presence_penalty parameter
- D. Replace the model with a smaller deployment.

Answer: A

You should implement an evaluation and retry loop in your application code.

You must wrap the agent call inside a conditional code loop that evaluates the output against required criteria (such as a checklist of regulatory clauses), and programmatically forces a retry if information is omitted.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry/agents/how-to/tools/ai-search>

QUESTION: 118

Hotspot Question

You have a Microsoft Foundry project that contains a customer support application.

You create an evaluation named Run1 that has the following configurations:

- Includes risk and safety metrics
 - Includes the protected material evaluation
 - Includes harmful content metrics that use a medium severity threshold
- You create an evaluation named Run2 that has the following configurations:**

- Includes risk and safety metrics
- Includes the protected material evaluation
- Includes harmful content metrics that use a high severity threshold

You run both evaluations against a dataset named DB1 and receive the following results:

- Content harm defect rate of Run1: 12%
- Content harm defect rate of Run2: 4%
- Protected material evaluation of Run1: 6%
- Protected material evaluation of Run2: 6%

You start a fine-tuning job by using DB1. The job fails during automatic RAI checks for multiple content harm types.

You discover that the content filtering configuration is set to high severity.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
Changing the content filtering configuration to low severity will resolve the fine-tuning job issues.	☐	☐
The difference between the 12% and 4% content harm defect rate is consistent with the different severity thresholds used in Run1 and Run2.	☐	☐
The identical 6% protected material evaluation values across Run1 and Run2 indicate that this metric is unaffected by the change in the severity threshold.	☐	☐

Answer:

Answer Area

Statements	Yes	No
Changing the content filtering configuration to low severity will resolve the fine-tuning job issues.	☐	☒
The difference between the 12% and 4% content harm defect rate is consistent with the different severity thresholds used in Run1 and Run2.	☒	☐
The identical 6% protected material evaluation values across Run1 and Run2 indicate that this metric is unaffected by the change in the severity threshold.	☒	☐

Explanation:

Box 1: No

Changing the content filtering configuration to low severity will not resolve the fine-tuning job issues.

Box 2: Yes

The difference between a 12% defect rate in Run1 and a 4% defect rate in Run2 is entirely consistent with the different severity thresholds configured for the two evaluations.

Box 3: Yes

The protected material evaluation metric is entirely unaffected by changes to the content harm severity threshold.

Reference:

<https://learn.microsoft.com/en-us/python/api/azure-ai-evaluation/azure.ai.evaluation>

QUESTION: 119

You are building an app by using the Semantic Kernel.

You need to include complex objects in the prompt templates of the app. The solution must support objects that contain subproperties.

Which two prompt templates can you use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. YAML**
- B. Liquid**
- C. Handlebars**
- D. Semantic Kernel**
- E. JSONL**

Answer: B,C

Semantic Kernel provides support for the following template formats:

semantic-kernel - Built-in Semantic Kernel format.

handlebars - Handlebars template format.

liquid - Liquid template format

The Semantic Kernel prompt template language is a simple way to define and compose AI functions using plain text. You can use it to create natural language prompts, generate responses, extract information, invoke other prompts or perform any other task that can be expressed with text.

Reference:

<https://learn.microsoft.com/en-us/semantic-kernel/concepts/prompts/prompt-template-syntax>

QUESTION: 120

You have a Semantic Kernel agent that provides answers to product questions.

You plan to extend the capabilities of the agent to answer questions about the status of orders.

You need to ensure that the agent answers order status questions as quickly as possible. The solution must minimize development effort.

What should you include in the solution?

- A. a kernel
- B. a plugin
- C. a function
- D. an agent

Answer: B

To enable the agent to answer questions about order status, you need a plugin.

In Semantic Kernel, a plugin acts as a container for the specific capabilities (functions) that allow an agent to interact with external data or systems, such as an order database or API.

To help me give you a more specific example or guide, could you tell me:

The data source where order info is kept (e.g., a SQL database, a REST API, or Shopify) The programming language you're using for the Kernel (C# or Python) Reference:

<https://learn.microsoft.com/en-us/semantic-kernel/frameworks/agent/agent-functions>

QUESTION: 121

You have a web app named App1 that sends requests to a multimodal chat model deployment in a Microsoft Foundry project.

User messages can contain both text and images.

Currently, App1 includes image URL: as plain text inside the message content so the model cannot recognize them as images.

Traces show that the requests contain a single text message instead of a multimodal content array.

You need to send the message as a structured array that includes both the text portion and the image reference to ensure that the model can process the image correctly.

What should you do?

- A. Set the user message content array to include items that have type: text and type: image_url.
- B. Encode the image to base64 and include the encoded data inside the content string of the user message.
- C. Add the image URL to the request metadata section, so the model can resolve the processing issue automatically.
- D. Place the image URL inside the System Message and set type to image_url so the model loads the image at initialization.

Answer: A

To fix this, you must change your request payload structure from a single string to a structured content array. Multimodal models in Azure AI Foundry expect an array of distinct objects for text and images.

Reference:

<https://towardsdatascience.com/multimodal-ai-search-for-business-applications-65356d011009/>

QUESTION: 122


Hotspot Question


You are creating an enrichment pipeline that will use Azure AI Search. The knowledge store contains unstructured JSON data and the text from scanned PDF documents.

Which projection type should you use for each data type? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.


Answer Area


JSON data: 

Extracted text data: 

Answer:

Answer Area

JSON data: 

Extracted text data: 

Explanation:

Box 1: Object projection

For the unstructured JSON data, you must use object projection, and for the extracted text data from the scanned PDF documents, you must use file projection.

Object Projection: Projects your data as a full JSON representation. It is ideal for maintaining the original structure alongside any applied AI enrichments within a single JSON document inside Azure Blob Storage.

Box 2: File projection

File Projection: Captures the binary or image extraction layer required during Optical Character Recognition (OCR) processing. It isolates and outputs raw text and graphic details directly from unstructured physical binary files into a designated container

Reference:

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-concept-image-scenarios>

QUESTION: 123

Case Study 1 - Contoso, Ltd

Overview

Company Information

Contoso, Ltd is a multinational retail company that builds, deploys, and manages generative AI and agent-based solutions by using Microsoft Foundry.

Existing Environment

Identity Environment

Contoso uses Microsoft Entra ID for identity management, authentication, and

authorization capabilities that enable agents to access organizational resources and services.

Contoso recently formed a new AI engineering team named Agent1Dev Team to optimize and maintain existing AI solutions.

The team collaborates with solution architects, DevOps engineers, and security engineers to design, implement, monitor, and secure AI applications.

Contoso also has a team named Agent1Test Team that is responsible for validating AI solutions before the solution deployments.

Generative Environment

Contoso has a Microsoft Foundry deployment that contains two projects named Project1 and Project2.

Project1

Project1 contains a customer support agent named Agent1 that assists customers with product inquiries and troubleshooting requests.

Agent1 has the following configurations:

- Agent1 uses a base model deployment.
- A safety evaluation pipeline is NOT enabled.
- Tool invocation approval workflows are NOT enabled.
- Conversation memory constraints are NOT configured.

Agent1 interacts with customers by using digital support channels and answers general questions about Contoso products.

Project1 is deployed to an Azure region located in the European Union (EU).

Agent1Dev Team will use Project1 to optimize and maintain Agent1.

Project2

Project2 contains a deployed video generation model. The marketing department at Contoso has access to Project2 and plans to use the model to develop a video creation solution.

Development of the solution is incomplete.

Data Environment

Contoso stores product-related information in Azure resources that support AI applications.

The Azure environment contains an Azure Blob Storage account named storage1 that stores product detail sheets for all the Contoso products.

The product sheets include specifications, feature descriptions, and product support information that Agent1 can use to answer customer questions. The product sheets are stored in the PDF format.

Problem Statements

Contoso identifies the following issues:

- Agent1 has only general knowledge of the Contoso products.
- A recent chat interaction with Agent1 was analyzed for sentiment. The results of the analysis have NOT been processed yet.
- Agent1 does NOT use the detailed product information in the product sheets stored in storage1 when responding to customer questions.
- The finance department at Contoso reports that vendor invoices must be reviewed manually to ensure that the invoices match the terms defined in the vendor contracts. The invoices contain tables, logos, and varied layouts that make the documents difficult to process consistently.

Requirements

Planned Changes

Contoso plans to implement the following changes:

- Implement a solution for Project1 that analyzes the vendor invoices by evaluating both the visual layout and the textual content of the invoices, so that the invoice details can be verified against the vendor contract terms.
- Update the base model deployment used by Agent1 and standardize the model version to ensure continuity and consistent responses.
- Enable Agent1 to retrieve and use the detailed product information from the product sheets stored in storage1.
- Implement an indexing solution for the product sheets that Agent1 can use to answer customer questions.
- Complete the development of the video creation solution.

Technical Requirements

Contoso identifies the following technical requirements:

- The model deployment used by Agent1 must support scalable, high-throughput generative AI workloads and dynamically scale to handle variable customer support traffic, without requiring reserved throughput capacity.
- The product sheets must be processed by using an indexing pipeline that enables semantic and vector search, so that Agent1 can retrieve the relevant product information.
- Responses generated by using the product sheet information must be relevant, complete, and accurate.
- Agent1 must be able to use the product sheets to answer natural language questions about product details.
- The model version used by Agent1 must remain consistent to ensure stable responses.
- The data processed by the model must remain within the EU.

Security and Compliance Requirements

Contoso identifies the following security and compliance requirements:

- API keys must NOT be used to access Foundry-deployed models.
- Access to the Azure resources must follow the principle of least privilege.
- The developers at Contoso must authenticate to Microsoft Foundry resources by using Microsoft Entra authentication.
- Access to Project1 must be assigned to the members of Agent1Dev Team by using a security group named SC_Agent1_Dev.
- Access to Project1 must be assigned to the members of Agent1Test Team by using a security group named SC_Agent1_Test.
- Agent1 must never reveal customer information, even if a document that contains customer data is added erroneously to the product sheet repository in storage1.
- The product sheets might contain images that include embedded text. Agent1 must be protected from malicious instructions potentially hidden within the images.

Business Requirements

Contoso identifies the following business requirements:

- Users that interact with Agent1 must have a personalized experience in future interactions, including the ability for Agent1 to retain conversation context and recall relevant information from previous interactions.
- Agent1 must answer questions only about the products sold by Contoso.

You need to recommend a solution to assess the responses generated by Agent1 when the agent uses the product information stored in storage1. The solution must meet the technical requirements. What should you include in the recommendation?

- A.** a Retrieval Augmented Generation (RAG) evaluator
- B.** a custom guardrail
- C.** model fine-tuning
- D.** a groundedness evaluator

Answer: A

The best way for the Foundry Agent to access the product information stored in the Azure Blob Storage account is by connecting an Azure AI Search index as an indexed knowledge source via the Agent's file search tool.

This implementation creates a robust Retrieval-Augmented Generation (RAG) framework, ensuring the agent retrieves grounded, high-relevance product information.

Scenario:

Technical Requirements;

*-> Responses generated by using the product sheet information must be relevant, complete, and accurate.

Data environment: The Azure environment contains an Azure Blob Storage account named storage1 that stores product detail sheets for all the Contoso products.

Planned changes: Enable Agent1 to retrieve and use the detailed product information from the product sheets stored in storage1.

Reference:

<https://learn.microsoft.com/en-nz/answers/questions/5868427/uploading-data-in-microsoft-foundry-agent>

QUESTION: 124

You have an Azure AI Search indexer that ingests PDF policy manuals.

Client applications must display page-level citations that have bounding polygons for both text and images.

You need to add a single built-in multimodal content extraction skill to the Azure AI Search skillset. The solution must meet the following requirements:

- Provide text and image location metadata.**
- Extract tables that span multiple pages.**

What should you add?

- A. Document Extraction**
- B. Azure Content Understanding in Foundry Tools**
- C. GenAI Prompt**
- D. Document Layout**

Answer: B

To meet all requirements for this project, you need the Azure Content Understanding skill (integrated via Azure AI Foundry / Microsoft Foundry Tools).

1. Extracting Cross-Page Tables: A critical constraint is the ability to recognize and extract tables that span multiple pages as a single unit. The Azure Content Understanding skill natively supports this capability, whereas the older Document Layout skill outputs layout content as flattened text/markdown boundaries that break across pages, leading to information loss.

2. Page-Level Bounding Polygons: It functions as a layout-aware, built-in multimodal content extraction skill that maps out text and image location metadata (location_metadata), satisfying your client application requirement for strict page-level

citation boundaries.

Reference:

<https://learn.microsoft.com/en-us/azure/search/cognitive-search-skill-content-understanding>

QUESTION: 125

You are developing a new sales system that will process user-generated video and text from a public-facing website.

You plan to notify users that their data has been processed by the sales system.

Which responsible AI principle does this help meet?

- A.** fairness
- B.** transparency
- C.** inclusiveness
- D.** reliability and safety

Answer: B

Notifying users that their data has been processed by your sales system fulfills the Transparency responsible AI principle. Transparency ensures that users are informed about how their data is collected and processed, which helps foster understanding and build trust between users and technology providers.

Reference:

<https://learn.microsoft.com/en-us/microsoft-copilot-studio/guidance/responsible-ai>

QUESTION: 126

An agent must call your organisation's internal inventory system, which is exposed as a set of tools. You want the tool definitions maintained centrally and reusable across several agents rather than redefined in each agent. Which Foundry capability fits?

- A.** Add a Model Context Protocol (MCP) server as an agent tool
- B.** Increase the model's context window
- C.** Fine-tune the model on the inventory system's API schema
- D.** Enable the code interpreter tool

Answer: A

A Model Context Protocol (MCP) server exposes a reusable, centrally maintained set of tools that any agent can connect to, and Foundry Agent Service supports adding MCP servers from the tools catalogue. This is the standard pattern for connecting agents to external systems through shared tool definitions.

QUESTION: 127

Hotspot Question

You have a Microsoft Foundry project.

- You need to create a customer support agent that meets the following requirements:

- Grounds responses only in company policy documents stored in curated repositories

- Retains customer preferences across separate chat sessions

How should you configure the agent? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Knowledge grounding:

	▼
Configure revival from approved data sources.	
Upload the policy documents direct to the agent.	
Embed the policy documents directly into the agent instructions.	

Memory:

	▼
Use orchestration-managed session context.	
Enable agent memory that uses persistent storage.	
Retain user preferences in the state of the client application	

Answer:

Answer Area

Knowledge grounding: 

- Configure revival from approved data sources.
- Upload the policy documents direct to the agent.
- Embed the policy documents directly into the agent instructions.

Memory: 

- Use orchestration-managed session context.
- Enable agent memory that uses persistent storage.
- Retain user preferences in the state of the client application

Explanation:

Box 1: Configure retrieval from approved data sources

You should configure retrieval from approved data sources to meet this requirement. In Microsoft Foundry (and broader Azure AI Foundry / Azure OpenAI architectures), grounding an agent exclusively in company policy documents requires implementing a Retrieval-Augmented Generation (RAG) pattern. This ensures the AI model only answers using the provided context and does not hallucinate or rely on its public training data.

Box 2: Enable agent memory that uses persistent storage

To retain user preferences across completely separate chat sessions or conversations, you must use persistent agent memory. Traditional chat history only tracks messages within a single active thread or session. By contrast, the Memory Store feature provides a managed, long-term memory system that structurally indexes user preferences and facts across different devices and separate sessions.

Reference:

<https://techcommunity.microsoft.com/blog/azure-ai-foundry-blog/announcing-new-capabilities-for-azure-openai-on-your-data/4144636>

QUESTION: 128

You have a Microsoft Azure AI Foundry project named Project1.

You plan to create an app named App1 that will connect to Project1 and chat by using a generative AI model.

You need to connect App1 to Project1 by using the Azure AI Foundry SDK. The solution must minimize development effort.

What should you configure in App1?

- A.** a connection string
- B.** a SASCredentials object

- C. an AIProjectClient object
- D. a project scope key

Answer: A

To minimize development effort when using the Azure AI Foundry SDK, you should configure the AIProjectClient object using the project connection string.

This approach is highly efficient because the connection string is a single string that encapsulates multiple required parameters--such as the subscription ID, resource group name, and project name--allowing the client to initialize and authenticate with minimal code.

How to Configure the AIProjectClient

1. Retrieve Connection String: In the Azure AI Foundry portal, go to the Overview page of your project. Under Project details, copy the Project connection string.
2. Initialize the Client: Use the from_connection_string method (or equivalent initialization depending on SDK version) to create the client.

Minimal Code Example (Python)

```
import os
from azure.ai.projects import AIProjectClient
from azure.identity import DefaultAzureCredential
# Load the connection string from an environment variable for security
conn_str = os.environ["PROJECT_CONNECTION_STRING"]
# Load the connection string from an environment variable for security
conn_str = os.environ["PROJECT_CONNECTION_STRING"]
```

Reference:

https://workshop.aifoundry.app/1-introduction/3-quick_start

QUESTION: 129

You have an Azure subscription that contains an Azure OpenAI resource.

You plan to build an agent by using the Azure AI Agent Service. The agent will perform the following actions:

- Interpret written and spoken questions from users.
- Generate answers to the questions.
- Output the answers as speech.

You need to create the project for the agent.

What should you use?

- A. the Azure portal

- B. Language Studio
- C. Azure AI Foundry
- D. Speech Studio

Answer: C

Azure AI Foundry is a platform for designing, customizing, managing, and supporting AI applications and agents. It acts as an AI app factory, providing a unified environment with tools, models, and deployment pipelines for various AI tasks. It enables teams to build and operate AI solutions, including those powered by generative AI, while ensuring security, governance, and cost-efficiency.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-foundry/what-is-azure-ai-foundry>

QUESTION: 130

Hotspot Question

You have a Microsoft Foundry project that contains an agent.

The agent uses a stored access key to retrieve secrets from an Azure key vault, which violates a keyless-credentials requirement.

You need to ensure that the agent can retrieve the secrets. The solution must follow the principle of least privilege.

What should you configure? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Managed identity scope:	▼ Enable a system-assigned managed identity at the Foundry level. Enable a system-assigned managed identity at the project level. Create a service principal and store the principal's client secret.
Key Vault authorization method:	▼ Add an API key to application settings. Add a Key Vault access policy for the secrets. Assign the Key Vault Secrets User role to the managed identity.

Answer:

Answer Area

Managed identity scope:

▼
Enable a system-assigned managed identity at the Foundry level.
Enable a system-assigned managed identity at the project level.
Create a service principal and store the principal's client secret.

Key Vault authorization method:

▼
Add an API key to application settings.
Add a Key Vault access policy for the secrets.
Assign the Key Vault Secrets User role to the managed identity.

Explanation:

Box 1: Enable a system-assigned managed identity at the project level

The correct configuration is to enable a system-assigned managed identity at the project level.

This is the correct approach. Azure AI Foundry (formerly Azure Machine Learning) allows you to enable a system-assigned managed identity directly at the project level. This identity acts as a keyless credential native to Azure. To follow the principle of least privilege, you then assign this project-level identity a specific role (such as Key Vault Secrets User) restricted only to the required Azure Key Vault. This ensures that the agent can retrieve secrets securely without exposing credentials to other projects or resources. Box 2: Assign the Key Vault Secrets User role to the managed identity To comply with the keyless-credentials requirement and follow the principle of least privilege, you should configure the Key Vault authorization method to assign the Key Vault Secrets User role to the managed identity using Azure Role-Based Access Control (Azure RBAC).

Reference:

<https://learn.microsoft.com/en-us/azure/foundry/agents/how-to/ai-gateway>

<https://learn.microsoft.com/en-us/azure/app-service/app-service-key-vault-references>

QUESTION: 131

A copilot must answer multi-part questions that depend on earlier turns in the conversation. You want the system to break each complex question into focused subqueries, run them in parallel, and semantically rerank the results before composing an answer. Which Azure AI Search capability provides this?

A. Classic keyword search with a scoring profile

B. Agentic retrieval

- C. A single vector query with a high similarity threshold
- D. Increasing the top-k value on a single keyword query

Answer: B

Agentic retrieval in Azure AI Search uses a large language model to decompose a complex query into focused subqueries, runs them in parallel, semantically reranks each set of results, and merges them into a unified response, while taking conversation history into account. That is precisely the behaviour the scenario describes.

QUESTION: 132

You have an Azure subscription that contains an Azure OpenAI resource named AI1. You build a chatbot that uses AI1 to provide generative answers to specific questions. You need to ensure that questions intended to circumvent built-in safety features are blocked.

Which Azure AI Content Safety feature should you implement?

- A. Moderate text content
- B. Protected material text detection
- C. Prompt Shields
- D. Monitor online activity

Answer: C

Prompt Shields in Azure AI Content Safety are specifically designed to detect and block prompts or questions that attempt to bypass safety features in generative AI models, such as those used in chatbots. This feature ensures that the chatbot adheres to safety and ethical guidelines by preventing inappropriate or harmful content generation.

QUESTION: 133

A legal team must extract clauses and inferred fields from unstructured contracts. One required field is the contract end date, which is not stated explicitly and must be derived from the start date plus the term. The team has no labelled training data. Which tool best fits?

- A. An Azure AI Document Intelligence custom model trained on labelled samples

- B.** Azure Content Understanding using a prebuilt-contract or zero-shot custom analyser
- C.** The Azure AI Search semantic ranker
- D.** Azure Translator in Foundry Tools

Answer: B

Azure Content Understanding handles unstructured documents, works zero-shot without labelled data, and can infer fields that are not explicitly present, such as deriving a contract end date from the start date and term. That matches every requirement in the scenario.

QUESTION: 134

You have a Microsoft Foundry project that contains an agent and uses a GitHub repository. The repository contains a YAM file named File1 that defines the evaluation settings of the agent. You need to create a GitHub Actions workflow that runs the evaluation defined in File1 when a pull request (PR) is opened. How should you configure the workflow?

- A.** Set project-endpoint to the endpoint of the project.
- B.** Set evaluation-config to the path of the YAML file.
- C.** Set model-deployment-name to the deployed model.
- D.** Set tenant-id to the Microsoft Entra tenant ID

Answer: A

The correct configuration choice is to set fproject-endpoint to the endpoint of the project.
azure-ai-project-endpoint: This parameter is required and must be set to the URL endpoint of your Microsoft Foundry project. This aligns with setting project-endpoint.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry/how-to/evaluation-github-action>

QUESTION: 135

You have a Microsoft Foundry project that contains an agent.

The agent uses Azure AI Search for Retrieval Augmented Generation (RAG).

You plan to ingest and index PDF product manuals.

You need to build a solution that supports semantic similarity matching. The solution

must ensure that the agent retrieves relevant data when user questions use different wording than the product manuals.

Which indexing approach should you use?

- A.** vector search
- B.** semantic ranking
- C.** suggesters
- D.** analyzers

Answer: A

Vector search is the correct indexing and retrieval type to use for this solution.

Vector search versus semantic ranking

While both features help bridge the gap between different wordings, vector search is specifically designed at the indexing and retrieval layer to handle semantic similarity matching by converting text into mathematical vectors (embeddings) based on conceptual meaning. Semantic ranking is a secondary re-ranking layer applied after initial retrieval to improve precision, but it cannot function as the primary indexing method on its own.

Reference:

<https://www.scribd.com/document/866453317/Agentforce-Specialist>

QUESTION: 136

Case Study 1 - Contoso, Ltd

Overview

Company Information

Contoso, Ltd is a multinational retail company that builds, deploys, and manages generative AI and agent-based solutions by using Microsoft Foundry.

Existing Environment

Identity Environment

Contoso uses Microsoft Entra ID for identity management, authentication, and authorization capabilities that enable agents to access organizational resources and services.

Contoso recently formed a new AI engineering team named Agent1Dev Team to optimize and maintain existing AI solutions.

The team collaborates with solution architects, DevOps engineers, and security engineers to design, implement, monitor, and secure AI applications.

Contoso also has a team named Agent1Test Team that is responsible for validating AI solutions before the solution deployments.

Generative Environment

Contoso has a Microsoft Foundry deployment that contains two projects named Project1 and Project2.

Project1

Project1 contains a customer support agent named Agent1 that assists customers with product inquiries and troubleshooting requests.

Agent1 has the following configurations:

- Agent1 uses a base model deployment.**
- A safety evaluation pipeline is NOT enabled.**
- Tool invocation approval workflows are NOT enabled.**
- Conversation memory constraints are NOT configured.**

Agent1 interacts with customers by using digital support channels and answers general questions about Contoso products.

Project1 is deployed to an Azure region located in the European Union (EU).

Agent1Dev Team will use Project1 to optimize and maintain Agent1.

Project2

Project2 contains a deployed video generation model. The marketing department at Contoso has access to Project2 and plans to use the model to develop a video creation solution.

Development of the solution is incomplete.

Data Environment

Contoso stores product-related information in Azure resources that support AI applications.

The Azure environment contains an Azure Blob Storage account named storage1 that stores product detail sheets for all the Contoso products.

The product sheets include specifications, feature descriptions, and product support information that Agent1 can use to answer customer questions. The product sheets are stored in the PDF format.

Problem Statements

Contoso identifies the following issues:

- Agent1 has only general knowledge of the Contoso products.**
- A recent chat interaction with Agent1 was analyzed for sentiment. The results of the analysis have NOT been processed yet.**
- Agent1 does NOT use the detailed product information in the product sheets stored in storage1 when responding to customer questions.**
- The finance department at Contoso reports that vendor invoices must be reviewed manually to ensure that the invoices match the terms defined in the vendor contracts. The invoices contain tables, logos, and varied layouts that make the documents**

difficult to process consistently.

Requirements

Planned Changes

Contoso plans to implement the following changes:

- Implement a solution for Project1 that analyzes the vendor invoices by evaluating both the visual layout and the textual content of the invoices, so that the invoice details can be verified against the vendor contract terms.
- Update the base model deployment used by Agent1 and standardize the model version to ensure continuity and consistent responses.
- Enable Agent1 to retrieve and use the detailed product information from the product sheets stored in storage1.
- Implement an indexing solution for the product sheets that Agent1 can use to answer customer questions.
- Complete the development of the video creation solution.

Technical Requirements

Contoso identifies the following technical requirements:

- The model deployment used by Agent1 must support scalable, high-throughput generative AI workloads and dynamically scale to handle variable customer support traffic, without requiring reserved throughput capacity.
- The product sheets must be processed by using an indexing pipeline that enables semantic and vector search, so that Agent1 can retrieve the relevant product information.
- Responses generated by using the product sheet information must be relevant, complete, and accurate.
- Agent1 must be able to use the product sheets to answer natural language questions about product details.
- The model version used by Agent1 must remain consistent to ensure stable responses.
- The data processed by the model must remain within the EU.

Security and Compliance Requirements

Contoso identifies the following security and compliance requirements:

- API keys must NOT be used to access Foundry-deployed models.
- Access to the Azure resources must follow the principle of least privilege.
- The developers at Contoso must authenticate to Microsoft Foundry resources by using Microsoft Entra authentication.
- Access to Project1 must be assigned to the members of Agent1Dev Team by using a security group named SC_Agent1_Dev.
- Access to Project1 must be assigned to the members of Agent1Test Team by using a

security group named SC_Agent1_Test.

- Agent1 must never reveal customer information, even if a document that contains customer data is added erroneously to the product sheet repository in storage1.
- The product sheets might contain images that include embedded text. Agent1 must be protected from malicious instructions potentially hidden within the images.

Business Requirements

Contoso identifies the following business requirements:

- Users that interact with Agent1 must have a personalized experience in future interactions, including the ability for Agent1 to retain conversation context and recall relevant information from previous interactions.
- Agent1 must answer questions only about the products sold by Contoso.

You need to configure an indexing pipeline for Agent1 to retrieve the relevant product information in storage1. The solution must meet the technical requirement.

Which two built-in skills should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Azure OpenAI Embedding
- B. Entity Recognition
- C. Text Split
- D. Merge
- E. Language Detection
- F. key phrase extraction

Answer: A,C

The most essential skills for this scenario are Azure OpenAI Embedding and Text Split.

Azure OpenAI Embedding: This skill is critical for generating the vector representations (embeddings) of your text, which directly enables the required vector search capability.

Text Split: This skill is essential because LLMs and embedding models have strict token limits.

Breaking large product detail sheets into smaller chunks ensures the text fits into the embedding model and improves the accuracy of semantic search.

Incorrect:

[Not B]

Entity Recognition: This extracts specific entities like names, dates, or locations. While helpful for advanced filtering, it is not a foundational requirement to enable basic semantic

or vector search.

[Not D]

Merge: This skill combines text from multiple fields into a single string. Since product sheets are already unified documents, splitting and chunking them is the priority rather than merging separate fields.

[Not E]

Language Detection: This identifies the language of the input text. Unless your product sheets are completely multilingual and require conditional routing to different language models, this skill is secondary.

[Not F]

Key Phrase Extraction: This pulls out main talking points or keywords. This is primarily used for traditional keyword tagging or basic search indexing, whereas your requirement specifically dictates vector and semantic-based retrieval.

Scenario:

Technical Requirements;

*-> The product sheets must be processed by using an indexing pipeline that enables semantic and vector search, so that Agent1 can retrieve the relevant product information.

Data environment: The Azure environment contains an Azure Blob Storage account named storage1 that stores product detail sheets for all the Contoso products.

Planned changes: Enable Agent1 to retrieve and use the detailed product information from the product sheets stored in storage1.

Reference:

<https://www.rheininsights.com/blog/en/Retrieval+Augmented+Generation+with+Azure+AI+Search+and+Atlassian+Confluence.php>

QUESTION: 137

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You increase the value of the max_tokens parameter.

Does this meet the goal?

A. Yes

B. No

Answer: B

Correct:

* You add a reflection pass that regenerates the response if the required clauses are missing.

This is Self-Correction Strategy: A reflection pass allows an agent to evaluate its own initial output against specified constraints (e.g., checking for the presence of mandatory regulatory clauses). If the required text is missing, the agent triggers a programmatic self-correction or regeneration loop to include them before final delivery.

Incorrect:

* You increase the value of the max_tokens parameter.

Increasing the max_tokens parameter prevents the response from being cut off mid-sentence due to length constraints. However, it does not force the model's logic to explicitly include missing information that it chose to leave out earlier in the text.

* You increase the value of the temperature parameter.

Raising the temperature parameter increases randomness and creativity. For rigid compliance tasks like summarizing regulatory documents, higher temperature actually increases the risk of hallucination and omission.

* You run an evaluation flow that scores responses for completeness and blocks responses that fall below a defined threshold.

Evaluation Flow Block: Running an evaluation flow to score completeness and blocking bad responses identifies and stops low-quality outputs, but it does not fix or actively improve the response completeness. It simply filters failures out of the system.

Reference:

<https://pub.towardsai.net/reflection-with-llm-how-to-make-ai-review-its-own-work-2db122fca1d8>

QUESTION: 138

Hotspot Question

You plan to create a Microsoft Foundry project named Project1 that will contain an agent and use an Azure key vault named KV1.

You need to configure a connection from Project1 to KV1.

How should you complete the Bicep code? To answer, select the appropriate options in the answer area?


NOTE: Each correct selection is worth one point.

Answer Area

```
resource existingKeyVault 'Microsoft.KeyVault/vaults@2024-11-01'
existing = {
  name: 'KV1'
  scope: resourceGroup()
}
resource connection
'Microsoft.CognitiveSevices/accounts/connections@2025-04-01-
preview' = {
  name: '${aiFoundryName}-keyvault'
  parent: aiFoundry
  properties: {
    category: 
    target: existingKeyVault.id
    authType: 
    isSharedToAll:
      metadata: {
        ApiType: 'Azure'
        ResourceId: existingKeyVault.id
        location: existingKeyVault.location
      }
  }
}
```

Answer:

Answer Area

```
resource existingKeyVault 'Microsoft.KeyVault/vaults@2024-11-01'
existing = {
  name: 'KV1'
  scope: resourceGroup()
}
resource connection
'Microsoft.CognitiveServices/accounts/connections@2025-04-01-
preview' = {
  name: '${aiFoundryName}-keyvault'
  parent: aiFoundry
  properties: {
    category: 
    target: existingKeyVault.id
    authType: 
    isSharedToAll:
      metadata: {
        ApiType: 'Azure'
        ResourceId: existingKeyVault.id
        location: existingKeyVault.location
      }
  }
}
```

Explanation:

Box 1: AzureKeyVault

To configure a connection from a Microsoft Foundry project to an Azure Key Vault, you must use the AzureKeyVault connection category.

Box 2: AccountManagedIdentity

To connect a Microsoft Foundry project to an Azure Key Vault, you must use AccountManagedIdentity as the authType.

Microsoft Foundry secures its outbound connections to core Azure storage and security infrastructure using Microsoft Entra ID role-based access control (RBAC). Using AccountManagedIdentity completely eliminates the need to handle or rotate explicit credentials or api keys within your configuration.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry/how-to/set-up-key-vault-connection>

QUESTION: 139

You are developing an app that will perform a sentiment analysis of social media posts by using the Azure AI Language service.

You perform a test on a sample post.

You need to quantify the results of the test.

Which JSON property should you review?

A. confidenceThreshold

B. sentiment

C. confidenceScores

D. relations

Answer: C

Sentiment analysis

The sentiment analysis feature assigns sentiment labels, such as "negative," "neutral," and "positive." The service determines these labels using the highest confidence score.

Sentiment is evaluated at both the sentence level and the document level. This feature also returns confidence scores between 0 and 1 for each document & sentences within it for positive, neutral, and negative sentiment.

In Azure AI Language's sentiment analysis, confidenceScores are numerical values between 0 and 1 that represent the probability that the text belongs to a specific sentiment (positive, neutral, or negative). A score closer to 1 indicates a higher confidence from the service that the text exhibits that sentiment, while a lower score signifies less confidence in that particular label. The service calculates these scores for both individual sentences and the entire document, providing a granular understanding of sentiment.

How to interpret confidenceScores:

High Score (close to 1): The model is very sure about the assigned sentiment. For example, a positive score of 0.95 means the model is 95% confident the text is positive.

Low Score (close to 0): The model is not very sure about the assigned sentiment.

Scores for each sentiment: For any given piece of text, the service returns a score for positive, neutral, and negative sentiment. The sentiment label that receives the highest score is assigned as the overall sentiment for that text.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/language-service/sentiment-opinion-mining/overview>

QUESTION: 140

Hotspot Question

You develop a test method to verify the results retrieved from a call to the Azure Vision in Foundry Tools API. The call is used to analyze the existence of company logos in images. The call returns a collection of brands named brands.

You have the following code segment:

```
for brand in image_analysis.brands:
    if brand.confidence >= 0.75:
        print(f"\nLogo of {brand.name} between
{brand.rectangle.x}, {brand.rectangle.y} and {brand.rectangle.w},
{brand.rectangle.h}")
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The code will display the name of each detected brand with a confidence equal to or higher than 75 percent.	☐	☐
The code will display coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.	☐	☐
The code will display coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.	☐	☐

Answer:

Answer Area

Statements	Yes	No
The code will display the name of each detected brand with a confidence equal to or higher than 75 percent.	☒	☐
The code will display coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.	☒	☐
The code will display coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.	☐	☒

Explanation:

Box 1: Yes

The code segment correctly filters for and displays the name (and coordinates) of each detected brand only if the model's confidence score is 75 percent or higher. The expression `if brand.confidence >= 0.75` guarantees that only brands meeting or exceeding this threshold are printed.

Box 2: Yes

The code segment will display the coordinates. Specifically, it prints the x and y values of the rectangle's top-left corner alongside its width (w) and height (h) for any detected brand with a confidence score equal to or greater than 0.75 (75%).

The provided code uses the properties directly to extract the bounding box:

`brand.rectangle.x` and `brand.rectangle.y`: The coordinates of the top-left corner of the bounding box.

`brand.rectangle.w` and `brand.rectangle.h`: The width and height of the bounding box.

Box 3: No

See Box 2 above.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/computer-vision/concept-brand-detection>

QUESTION: 141

You are designing a solution that will answer questions about human resources (HR) policies stored in the PDF format.

You need to ensure that the identical answer to a specific question is returned every time. The solution must minimize development effort.

Which service should you include in the solution?

- A.** Azure OpenAI
- B.** Azure AI Document Intelligence
- C.** Azure Machine Learning
- D.** Azure AI Language

Answer: A

You can create a Custom FAQs PDF Solution Powered by Azure OpenAI with Citations from Grounded Data.

Business Problem

Imagine an organization that frequently deals with FAQs, legal documents, or reports. They want an automated system that can:

Generate answers from questions using the information found in existing PDF documents. Ensure citations from those documents are included in the responses. Output a structured PDF with the Q&A and citations formatted in a professional and readable way.

Solution Overview

By leveraging Azure OpenAI on your data and Azure Cognitive Search, we can ground the OpenAI GPT model with custom documents like PDFs. When queried, the model retrieves the most relevant responses from the documents and includes citations such as document titles or page numbers and create a question-and-answer format JSON file. The system then generates a structured, professional-looking PDF with bold questions, clear answers, and citations.

Reference:

<https://www.linkedin.com/pulse/creating-custom-faqs-pdf-solution-powered-azure-openai-saqlain-tahir-x9gse>

QUESTION: 143

Drag and Drop Question

You have a Microsoft Foundry project that contains an agent. The agent uses threads and file uploads and calls an Azure OpenAI model deployment.

During load testing, calls intermittently fail and return an HTTP 429 rate limit exceeded error.

Some user uploads fail and generate an HTTP 400 file size exceeded error.

You need to mitigate the errors and reduce call failures. The solution must remain within the service and model limits.

What should you do to resolve each error? To answer, drag the appropriate actions to the correct errors. Each action may be used once, more than once or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each comet selection is worth one point.

Actions

Answer Area

Increase tenant-wide quotas.

Move large content to files
and use file search.

Use additional agent tools
to reduce the message size.

Implement exponential backoff
and jitter in the retry logic.

Split content into smaller files
before uploading the files.

HTTP 429:

HTTP 400:

Answer:

Actions

Answer Area

Increase tenant-wide quotas.

Move large content to files
and use file search.

Use additional agent tools
to reduce the message size.

Implement exponential backoff
and jitter in the retry logic.

Split content into smaller files
before uploading the files.

HTTP 429:

HTTP 400:

Implement exponential backoff
and jitter in the retry logic.

Split content into smaller files
before uploading the files.

Explanation:

Box 1: Implement exponential backoff and jitter in the retry logic

To remedy intermittent HTTP 429 rate limit errors while staying within service limits, you must implement a retry policy that uses exponential backoff and jitter.

Because the agent uses threads and file uploads, load testing triggers short-window concurrency spikes (bursting). Azure OpenAI evaluates rate limits in small slices (like 1- to 10-second windows), meaning overlapping file processing or concurrent thread steps will trigger a 429 even if your total Tokens-Per-Minute (TPM) or Requests-Per-Minute (RPM) look safe overall.

Box 2: Split content into smaller files before uploading the files.

Instead of uploading a large document in a single standard files.create request, utilize the Chunked Uploads REST API path. This splits the file payload on your client application layer before streaming it to Azure.

Note: To resolve the HTTP 400 file size exceeded error in your Microsoft Foundry agent, you need to bypass the strict payload constraints of standard single-request uploads. In Azure OpenAI and Microsoft Foundry Agent architectures, a multipart standard upload has a rigid request body limit (typically 30 MB or 50 MB).

Reference:

<https://learn.microsoft.com/en-us/answers/questions/2265002/getting-error-after-deployed-a-model-in-azure-ai-f>

[https://learn.microsoft.com/en-us/answers/questions/5521436/getting-400-error\(request-body-too-large\)-for-batc](https://learn.microsoft.com/en-us/answers/questions/5521436/getting-400-error(request-body-too-large)-for-batc)

QUESTION: 144

You have an invoice-processing application named App1 that uses Azure Constant Understanding in Foundry Tools.

You are building a new Content Understanding pipeline named Pipeline1 that must meet the following requirements:

- Compare an invoice to its related purchase order**
- Validate the voice against static vendor contact documents**
- Return a single structured output that includes discrepancy findings**

You need to configure Pipeline1 and expose the pipeline as a single analyzer endpoint.

What should you configure?

- A.** a single-file task in standard mode that uses the vendor contract provided as an additional document during analysis.
- B.** a single-file task in standard mode that uses confidence scores enabled for the extracted fields.
- C.** a multiple-file task in pro mode that uses the vendor contract files as reference data
- D.** a multi-file task in standard mode that uses the invoice and purchase order as input to the analyzer

Answer: C

Multiple-file task: Required over a single-file task. Your pipeline needs to reconcile data across two separate active transactional documents (the invoice and the purchase order) within a single analyzer request.

Pro mode is required instead of standard mode. Pro mode is specifically designed for advanced scenarios requiring multi-step reasoning, cross-file analysis, and validation against a knowledge base. Note that Pro mode currently supports classify and generate fields but does not support confidence scores for specific extracted fields.

Vendor contract files as reference data: Required. Because the vendor contracts are static compliance documents, they should be uploaded and treated as the analyzer's background knowledge base (reference data) to guide the validation logic.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/content-understanding/concepts/standard-pro-modes>

QUESTION: 145

You have an Azure subscription that contains an Azure App Service app named App1.

You provision a Microsoft Foundry Service resource named CSAccount1.

You need to configure App1 to access CSAccount1. The solution must minimize administrative effort.

What should you use to configure App1?

- A.** the endpoint URI and subscription key
- B.** the endpoint URI and an OAuth token
- C.** the endpoint URI and a shared access signature (SAS) token
- D.** a system assigned managed identity and an X.509 certificate

Answer: A

To configure the App Service app to access the Microsoft Foundry Service (part of Azure AI Services) with low administration, you need to configure the app with the endpoint URI and a subscription key.

These credentials can be found in the Azure Portal under the Keys & Endpoint section of your provisioned Foundry resource, allowing the app to seamlessly authenticate API calls.

Reference:

<https://learn.microsoft.com/en-us/azure/ai-services/authentication>

QUESTION: 146

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses.

You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model.

You need to implement controls to mitigate the risk.

Solution: You configure a prompt shield for documents.

Does this meet the goal?

A. Yes

B. No

Answer: A

Correct:

* You configure a prompt shield for documents.

Prompt Shield for Documents: Highly Effective (Critical Defense)

How it helps: This shield specifically scans untrusted, third-party data inputs (like external documents or text extracted from uploaded images).

Mechanism: It evaluates the extracted image text before it is sent to the LLM to identify hidden jail

* You configure a prompt shield for user prompts.

Prompt Shield for User Prompts: Partially Effective (Defense in Depth)

How it helps: This shield targets direct jailbreak attempts written manually by the user in the text prompt field accompanying the upload.

Mechanism: It prevents the user from typing supporting instructions that prime the model to execute the hidden instructions found within the image.

* You configure image moderation to block unsafe content before processing the images.

Implementing rigorous image moderation is one of the most effective ways to secure multimodal AI systems against these threats. Moderation acts as a necessary gatekeeper, preventing malicious inputs from ever reaching the generative model.

Incorrect:

* You configure protected material detection.

Protected Material Detection: Ineffective for this Threat

Why it does not help: This feature is designed to scan model outputs to prevent the generation of copyrighted text, proprietary source code, or licensed imagery.

Limitation: It does not scan inputs for adversarial instructions and will not prevent a user from manipulating the model's logic.

Reference:

<https://www.upgrad.com/blog/what-is-multimodal-ai/>

<https://learn.microsoft.com/en-us/azure/ai-services/content-safety/concepts/jailbreak-detection>


QUESTION: 147


Implement generative AI and agentic solutions

You have an Azure subscription. You need to create a new resource that will generate fictional stories in response to user prompts. The solution must ensure that the resource uses a customer-managed key to protect data.

How should you complete the script?

Answer Area

```
az cognitiveservices account create -n myresource -g
myResourceGroup --kind  --sku S -l WestEurope

 '{
  "keySource": "Microsoft.KeyVault",
  "keyVaultProperties": {
    "keyName": "KeyName",
    "keyVersion": "secretVersion",
    "keyVaultUri": "https://issue23056kv.vault.azure.net/"
  }
}'
```

Answer:

Answer Area

```
az cognitiveservices account create -n myresource -g  
myResourceGroup --kind
```

	▼
AIService	
LanguageAuthoring	
OpenAI	

	▼
--api-properties	
--assign-identity	
--encryption	

```
'{  
  "keySource": "Microsoft.KeyVault",  
  "keyVaultProperties": {  
    "keyName": "KeyName",  
    "keyVersion": "secretVersion",  
    "keyVaultUri": "https://issue23056kv.vault.azure.net/"  
  }  
}'
```

QUESTION: 148

Implement text analysis solutions

You are designing a content management system.

You need to ensure that the reading experience is optimized for users who have reduced comprehension and learning differences, such as dyslexia. The solution must minimize development effort.

Which Azure service should you include in the solution?

AAzure Document Intelligence in Foundry Tools

BAzure Language in Foundry Tools

CAzure AI Immersive Reader

DAzure Translator in Foundry Tools

Answer: C


Explanation

Azure AI Immersive Reader is specifically designed to improve reading comprehension for people with learning differences such as dyslexia, as well as for new readers and language learners. It provides built-in, proven accessibility features—text-to-speech read-aloud, syllabification, parts-of-speech highlighting, line focus, adjustable spacing/fonts, and a picture dictionary—through a ready-to-use client library and UI, which can be embedded into an application with minimal custom development. This distinguishes it from Document Intelligence (document data extraction), Language (general NLP tasks like sentiment/key phrase analysis), and Translator (language translation), none of which target reading comprehension accessibility.

QUESTION: 149

Implement computer vision solutions

You are building a model to detect objects in images. The performance of the model based on training data is shown in the following exhibit.



Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

Answer Area

The percentage of false positives is:

	▼
0	
25	
30	
50	
100	

The value for the number of true positives divided by the total number of true positives and false negatives is:

	▼
0	
25	
30	
50	
100	

Answer:

Answer Area

The percentage of false positives is:

	▼
0	
25	
30	
50	
100	

The value for the number of true positives divided by the total number of true positives and false negatives is:

	▼
0	
25	
30	
50	
100	

Explanation:

Box 1: 0

The percentage of false positives is 0%.

Because the model made zero incorrect positive predictions, the count and percentage of false positives must be exactly zero.

Box 2: 25

The value for the number of true positives divided by the total number of true positives and false negatives is 25% (or 0.25).

High Precision (100%): Every single object your model detects is correct; it generates zero false alarms.

Low Recall (25%): Your model misses 75% of the actual objects it was supposed to find.

Reference:

<https://medium.com/grabngoinfo/how-to-evaluate-the-performance-of-a-binary-classification-model-6e7193dcbbf9>

QUESTION: 150

You configure Azure AI Content Safety to moderate user-uploaded images in a generative application. During testing, a colleague claims an image returned a Violence severity score of 3.

Why is this result impossible?

- A.** Content Safety does not moderate images, only text
- B.** The image content model returns only severity values 0, 2, 4, and 6
- C.** Severity 3 is reserved for the Self-Harm category only
- D.** Image moderation always returns a binary blocked or allowed flag, never a severity

Answer: B

The Azure AI Content Safety image model uses a trimmed severity scale that returns only 0, 2, 4, and 6, so an odd value such as 3 cannot occur for image content. The text model uses the full 0 to 7 scale, which is where odd-numbered severities appear.

QUESTION: 151

A production application authenticates to Microsoft Foundry using an API key stored in an environment variable. Your security team requires that you remove hardcoded secrets and enable per-principal auditing of every call. Which approach meets both requirements?

- A.** Move the API key into Azure Key Vault and rotate it monthly
- B.** Authenticate with Microsoft Entra ID using a managed identity (keyless)
- C.** Share a single API key across services and restrict it with IP allow-listing

D. Embed the API key in the deployment pipeline as a masked variable

Answer: B

Keyless authentication with Microsoft Entra ID issues short-lived OAuth bearer tokens scoped per principal, which removes hardcoded secrets and produces per-principal audit trails. A managed identity extends this to service-to-service calls without storing any credential in code or configuration.

QUESTION: 152

An application must translate a high volume of standard product descriptions into 20 languages quickly, consistently, and at the lowest cost. The text is straightforward, with no domain-specific nuance. Which approach best fits?

A. A large language model translation flow with a custom system prompt

B. Azure Translator in Foundry Tools

C. A separately fine-tuned model for each of the 20 target languages

D. Azure AI Document Intelligence

Answer: B

Azure Translator in Foundry Tools provides deterministic, broad-language translation at scale and at lower cost, which suits high-volume, straightforward content.

QUESTION: 153

You are building an image sharing app that will use Azure AI to prevent users from sharing sexually explicit images.

You need to ensure that inappropriate images are identified correctly. The solution must minimize development effort.

What should you use?

A. Azure AI Studio

B. Visual Studio

C. Azure AI Content Safety Studio

D. Vision Studio in Azure AI Vision

Answer: C

To prevent users from sharing sexually explicit images, you need a tool that can handle content moderation specifically designed for inappropriate content such as adult or sexually explicit material. Azure AI Content Safety Studio is a service designed for this purpose, providing pre-built AI models that can detect inappropriate content like adult images, violence, and other sensitive material. It minimizes development effort by offering an easy-to-use, pre-configured service for content safety without the need to build custom models.

QUESTION: 154

You plan to configure an evaluation in Microsoft Foundry for a Retrieval Augmented Generation (RAG) chat app.

You need to provide scores for groundedness, relevance, and harmful content categories.

Which two evaluation categories can you use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A.** risk and safety metrics
- B.** fluency evaluator
- C.** similarity evaluators
- D.** AI quality (NLP) metrics
- E.** AI quality (AI assisted) metrics

Answer: A,E

The two evaluation categories that would be of use are AI quality (AI assisted) metrics and risk and safety metrics. Risk and safety metrics: Scoring harmful content categories (such as violence, sexual content, hate speech, and self-harm) is handled directly by the specialized safety evaluators contained within this category.

AI quality (AI assisted) metrics: In Microsoft Foundry, evaluating RAG-specific criteria like groundedness and relevance requires large language models (LLMs) acting as a judge to assess semantic context. These are explicitly categorized under AI-assisted quality metrics rather than standard natural language processing (NLP) rules.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry/concepts/evaluation-evaluators/rag-evaluators>

<https://learn.microsoft.com/en-us/azure/foundry-classic/how-to/develop/evaluate-sdk>

QUESTION: 155

A developer needs a managed identity to call a Microsoft Foundry project endpoint to run inference and use agents. Security policy requires the least-privilege role that grants data-plane access and nothing more. Which built-in role should you assign to the managed identity?

- A.** Cognitive Services Contributor
- B.** Foundry User
- C.** Azure AI Developer
- D.** Foundry Account Owner

Answer: B

Foundry User is the least-privilege built-in role for Microsoft Foundry projects and grants the data- plane actions needed to call a project endpoint, run inference, and use agents. Microsoft recommends assigning the project's managed identity the Foundry User role on the Foundry resource.

QUESTION: 156

You have a Microsoft Foundry project.

You need to deploy a model from the model catalog to support a search solution for internal policy documents. The model must generate vector representations of the text in the documents and of user queries.

Which type of model should you use?

- A.** an embedding model
- B.** an image generation model
- C.** a large language model (LLM)
- D.** a small language model (SLM)

Answer: A

To fulfill this requirement, you need an embeddings model from the Microsoft Foundry AI Model Catalog that captures semantic meaning and outputs numerical vectors for text chunks.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry-classic/openai/concepts/understand-embeddings>

QUESTION: 157

You have an Azure subscription.

You plan to build an app that will use the Azure AI DALL-E model.

You need to deploy the model.

What should you use?

- A.** the Azure SDK for Python and PowerShell cmdlets.
- B.** the Azure SDK for JavaScript and Azure Machine Learning Studio.
- C.** Microsoft Foundry and the Azure Command Line Interface (CLI)
- D.** the Azure portal and Microsoft Graph API

Answer: C

Explanation:

To deploy the DALL-E model, you need Microsoft Foundry and the Azure Command Line Interface (CLI).

Deployment vs. Application Logic: The Azure CLI along with Microsoft Foundry handles the administrative infrastructure layer required to provision, configure, and manage model deployments within your subscription.

Incorrect: [Not B] The Role of the SDK: The Azure SDK for JavaScript is a client-side development library. It is used within your code to consume or call the model's API after it has already been deployed. It cannot deploy the DALL-E model asset itself.

Studio Clarification: While Azure Machine Learning Studio manages custom ML pipelines, foundational generative AI models like DALL-E are provisioned and hosted through the Microsoft Foundry portal and its associated ecosystem.

Reference:

<https://learn.microsoft.com/en-us/azure/foundry-classic/openai/how-to/create-resource>